

摘要：

GitHub Copilot按Token计费新政让部分用户月账单从44美元骤升至847美元，一场AI"补贴换增长"模式的崩塌就此引爆。OpenAI利润率接近负122%、Uber四个月耗尽全年AI预算.....当资本耐心触及极限，Coinbase、Hugging Face掌门人已给出答案：成本低99%的廉价模型与开源小模型，或将接管八成AI工作负载。

GitHub Copilot的定价转型正在引发AI行业的连锁反应，一场关于AI商业模式可持续性的深层争论由此浮出水面。随着按量计费取代固定订阅，用户账单骤然飙升，而Coinbase、Hugging Face等科技巨头的掌舵者正在给出截然不同的应对路径，廉价模型的崛起，或将从根本上重塑AI算力的成本格局。

6月1日，微软旗下GitHub Copilot正式将计费模式从按请求次数收费切换为按Token用量计费，部分重度用户的月度账单预计从此前的数十美元骤升至数百美元。这一变化迅速在社交媒体引发强烈反弹，有用户晒出内部成本估算截图，**显示其月度费用将从44.68美元跳升至754.29美元，另有用户预计账单将高达847美元。**

这场定价风波的背后，是AI行业长期以来以补贴换增长模式的集中爆发。Coinbase首席执行官Brian Armstrong对此作出回应，**预判80%的AI工作负载将在12至18个月内迁移至成本低99%的模型，能源与算力将成为真正的瓶颈所在。**

Hugging Face首席执行官Clement Delangue则援引斯坦福大学研究数据，为本地化、开源小模型的大规模替代提供了实证支撑。

GitHub Copilot定价转型：补贴时代终结

GitHub Copilot的定价调整并非突然之举。今年4月，GitHub首席产品官Mario Rodriguez已公开表示，随着智能体AI的兴起，现行定价模式"已不再可持续"——一次简短的对话提问与一次耗时数小时的自主编程任务，此前对用户收取相同费用，而GitHub一直在背后吸收不断攀升的推理成本。

新政策于6月1日正式生效。**在新计费体系下，用量成本按所使用的AI模型及消耗的Token数量折算为AI积分，每个积分等值0.01美元。订阅用户获得固定基础积分配额，并根据订阅层级获得额外弹性积分。**由于前沿AI模型通常消耗更多Token，不同模型之间的实际成本差异悬殊。

用户的反应迅速而激烈。在GitHub的Reddit社区，一名自称从首日起便订阅Copilot Pro+的用户写道："每月39美元感觉已经很贵了，但还算值。现在换成这套AI积分机制，我算了一下，下个月预计账单：847美元。"多名用户将这一变化比作Uber的商业路径——以超低价格培养用户依赖，待用户形成习惯后再大幅提价。

←  r/github · 2d ago
Skauzor

Time to move on?

GitHub Copilot is moving to usage-based billing

Starting June 1, 2026, Copilot usage will be measured in AI Credits (AICs) instead of Premium Requests (PRUs). **1 AIC = \$0.01**. This is a preview estimate based on your uploaded report. Actual bills under usage-based billing may differ based on model mix and final pricing.

Note: This is a preview estimate based on your uploaded report (premiumRequestUsageReport_1_870c459e998e49f6dc5d7ba5017f33b.csv). Actual bills under usage-based billing may differ based on model mix and final pricing.

[Learn more about usage-based billing →](#)

Your **April 2026** usage would cost **\$709.61 more** under usage-based billing

Upgrading to **Max** would reduce your total monthly bill by **\$69.00**. [Learn more about Max plan.](#)

CURRENT BILLING (PRUs)		USAGE-BASED BILLING (AICs)	
\$44.68	1,642 PRUs 1 PRU = \$0.04	\$754.29	78,529.105 AICs 1 AIC = \$0.01
Consumed PRUs	\$65.68	Consumed AICs	\$785.29
Included PRUs	-\$60.00	Included AICs	-\$70.00
Overages	\$5.68	Additional usage	\$715.29
License cost	\$39.00	License cost	\$39.00
Total (license + overages)	\$44.68	Total (license + additional usage)	\$754.29

Flex allotment is used in this simulation. [Learn more about flex allotment.](#)

📌 Estimated projection for illustrative purposes only. Actual usage may differ.
Charges are calculated from actual usage emissions processed by the billing platform, separate from the preview data pipeline. Possible gaps are a reporting issue, not a billing issue.

Been using copilot quite a bit for side-projects because I don't have that much free-time anymore with family, work etc. I knew changes were coming and a lot of people complained, but hadn't actually read much about what was going to change. This seals the deal, lol. I guess there's nothing to do except finding another service?

👍 58 🔄 69

Gartner分析师Arun Chandrasekaran在接受Business Insider采访时表示，Copilot的案例“可能只是一个早期样本”，预计随着高级推理模型和智能体 workflow 推动推理端算力消耗大幅上升，更多企业将转向按Token或按用量计费模式。

补贴模式的系统性风险

这场定价风波折射出AI行业更深层的结构性矛盾。投资人Tommy Shaughnessy在社交媒体上发文，系统梳理了他所认为的“AI最显而易见的崩溃路径”。

他指出，按席位订阅的固定费用长期以来被大幅补贴，远低于重度使用的实际成本。一旦企业出于数据保护、合规审批等需求转向API调用，便会直面按量计费的真实价格，而实际消耗速度往往远超此前预期。他援引多个案例佐证这一趋势，包括Uber在2026年将全年AI预算在四个月内耗尽。



Tommy  
@Shaughnessy119



The most basic way AI could blow up imo. I'm not saying it does but this is the most obvious way I can see it happening

- Per seat subscriptions are massively subsidized. The flat fee was priced way below what heavy usage actually costs
- For real business use you have to move to the API anyway. Data protections, work integrations and compliance officer approval
- On the API you pay metered rates, and businesses are burning credits way faster than the per seat pricing ever led them to expect
- This is everywhere right now. Internally for us, Codex users, Uber torching its entire 2026 AI budget in 4 months, the Microsoft comments. Just go try an API



Shaughnessy进一步指出，当前AI大厂的利润率已深度为负——据报道OpenAI的利润率接近负122%——这意味着其完全依赖外部资本来购买GPU、训练模型并持续补贴用量。他认为，一旦投资者对回报预期失去信心，整个资本流动将面临逆转风险。

不过，他也提示了这一逻辑的边界：若AI真正催生新药研发或全新商业形态，用户对高价AI服务的支付意愿将大幅提升，届时上述压力或将缓解。

Coinbase CEO：廉价模型将主导未来

面对算力成本的持续攀升，Coinbase首席执行官Brian Armstrong给出了他的判断框架。他认为，对智能的需求近乎无限，但市场将快速分化：**80%的工作负载将在12至18个月内迁移至成本低99%的模型，剩余20%对智能上限有极致要求的任务**——如科学突破、高层级编排智能体——仍将运行在最新一代前沿模型上。



Brian Armstrong  
@brian_armstrong



Good take

My guess is

- demand for intelligence is near infinite
- but 80% of workloads will be running on 99% cheaper models within 12-18 months
- 20% of workloads will still run on latest gen models where IQ maxing is important (scientific breakthroughs, higher level orchestrator agents?)
- rough analogy might be what % of macbooks or gaming PCs sold have the maxed out specs for CPU/GPU, prices are falling much faster than Moore's law here though
- this leads me to think the limiting factor will be energy and compute, not better models

At Coinbase we're working hard on routing prompts to cheaper models where appropriate, and in some cases have been able to keep costs roughly flat, while token usage continues to grow exponentially.

Armstrong将这一趋势类比为消费电子市场：购买顶配MacBook或游戏PC的用户始终是少数，而AI领域的价格下降速度甚至比摩尔定律更快。他由此得出结论，未来真正的制约因素将是能源与算力，而非模型能力本身。

Armstrong还披露了Coinbase的内部实践：公司正积极推进提示词路由策略，将请求分配至成本更低的模型，在部分场景下已实现总成本基本持平，而Token用量仍在指数级增长。

开源小模型：多模型未来的实证支撑

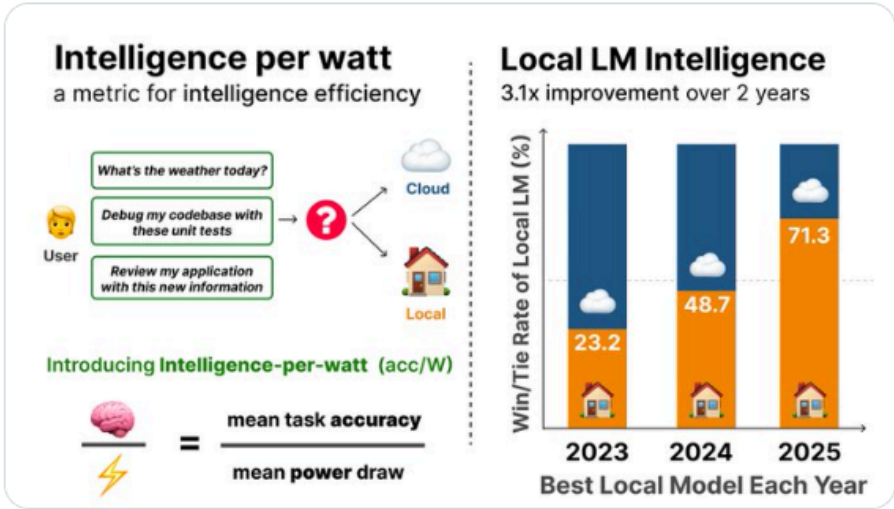
Hugging Face首席执行官Clement Delangue援引斯坦福大学研究数据，为廉价模型的替代潜力提供了量化依据：本地模型在真实世界的对话与推理查询中的准确率已从2023年的23.2%跃升至71.3%，且成本与能耗均仅为前沿API的一小部分。

Delangue据此提出“多模型未来”的判断：**对于大多数工作负载，本地化、开源、小型且廉价的模型将成为主流选择；只有在别无选择时，才需要调用前沿API。**



Narrative violation: according to [@Stanford](#) research, local models can answer 71.3% of real-world chat and reasoning queries accurately, up from 23.2% in 2023. Obviously at a fraction of the cost and energy consumption of frontier APIs.

The obvious conclusion: you don't need a frontier model for most tasks. The future is multi-model: local, open-source, smaller and cheaper for the majority of workloads, frontier APIs when no other choices!



Shaughnessy的分析与此形成呼应。他指出，DeepSeek V4在SWE-bench编程基准测试中的表现与Anthropic Claude Opus相近，但价格约为后者的三十分之一；最廉价的开源模型价格甚至低至约百分之一。他认为，中国实验室持续开源前沿级模型，使推理服务提供商得以免费获取最核心的模型成本，这正在从根本上压制闭源AI大厂的定价能力与利润空间。

风险提示及免责声明

市场有风险，投资需谨慎。本文不构成个人投资建议，也未考虑到个别用户特殊的投资目标、财务状况或需要。用户应考虑本文中的任何意见、观点或结论是否符合其特定状况。据此投资，责任自负。

 259  收藏

分享到:  

写评论

请发表您的评论

 表情  图片

发表评论