AUGUST 19, 2026 ACCELERATOR & HBM

Broadcom TPU Lower but MTIA and OAI Higher, MediaTek TPU 8t Increased, Qualcomm ByteDance Back, TikTok ASIC, Nvidia Updates // Accelerator Model Update – 19 August 2026

6 MINUTES

By Myron Xie, Ivan Chiam and Louis Lu

Download Model

Executive Summary

- We implement changes to Broadcom TPU output we flagged in our note last week, lowering TPU 7 and 8i production across this year and next. For next year, Broadcom is allocating more capacity than expected to Meta, OpenAI, and a new project codenamed “Lyra” whose end customer is yet to be determined.
- We see more capacity for TPU 8t with MediaTek for this year and next
- We make several updates to Nvidia, primarily lifting share of Rubin Ultra units next year. We also incorporate more HBM de-spec for Feynman and beyond
- According to industry chatter, the Qualcomm and ByteDance ASIC project is now back on track in August, after being suspended in June. In addition, ByteDance is exploring an ASIC for its US TikTok operations with Qualcomm and Marvell. This ASIC would be deployed at Oracle Cloud datacenters.
- On HBM despecing more broadly, while we are yet to implement this in our model outside of Nvidia, we plan to review our HBM content assumptions, especially in the light of customer demand for 4-hi that we wrote about last week. As HBM is the primary constraint on shipments, reducing stack heights to allow more cube production could unlock more total shipments.

Broadcom Update

In our note last week, we provided on update on Broadcom's output for 2026 and 2027. We implement these changes in the model and share additional details.

TPU

We are seeing lower TPU output in 2H26 than previous expectations. 2026 CoWoS output for TPU v7x (Ironwood) is reduced by ~32k, reducing 2026-unit production by ~500k from 3.2mn to 2.7mn. TPU v8i (Sunfish) is also reduced by ~5k wafers, reducing 2026-unit production by ~80k. Ironwood upstream output now peaks in 2Q26 with a slight decline in run-rate in 2H, versus our prior forecast of a continued ramp from 2Q26 onwards. We believe this reduction is supply related and due to challenges ramping CoWoS-S. While CoWoS-S is a mature technology, TPU is the only major platform using CoWoS-S, with other customers having transitioned to either CoWoS-L or CoWoS-R.   This brings total Broadcom 2026 CoWoS output to 215k wafers from 250k previously. We therefore reduce Broadcom's 2026 TPU shipments by 400k units across TPU 7x and 8i. Net of these reductions, our estimate for Broadcom AI ASIC revenue for FY2026 comes to $34bn revenue (previously $36bn), and with the additional AI networking revenue should allow Broadcom to meet its $56bn guidance.

For 2027, we lower TPU v7x and v8i shipments from our previous expectation of 6mn units to 3.8mn units. We see total Broadcom CoWoS at ~460k wafers vs 483k previously, so this TPU reduction is driven by greater share of allocation towards other customers rather than a significant total output drop for Broadcom in 2027. We also don't think this as a reduction of Google TPU demand, rather we see Broadcom wanting to support other key customers such as Meta and OpenAI to help progress their roadmaps. In addition, we see it as a fair possibility that these other projects do not progress as planned and TPU can absorb freed up capacity as long as they are N3 based.

Meta

For Meta, there are 55k wafers allocated primarily to MTIA 400 Iris+ in 2027. There is very little allocated towards MITA 450 Arke, which is the 2nm refresh of Iris+. However, this remains in flux as Meta is known for abandoning orders, as we saw in May, where Meta agreed to cut its own Broadcom production allocation to support greater TPU supply for Google in exchange for access to TPU compute. This could play out again and, if Meta abandons its allocation, this could be taken up by TPU. We note that the timing is heavily weighted to fourth quarter output, which provides more lead time for these volumes to be switched to another design if necessary. We therefore add 220k and 170k shipments of Iris in 2027 and 2028 respectively. We also reduce MTIA 300 shipments by 36k in 2026 to align with lower output than previously estimated.

OpenAI

The first gen OAI ASIC that will ship to production is Jalapeno which is the smaller chip of the 2 floorplans we have previously shown. We therefore consolidate units to this SKU. Overall, we see this program shipping total lifecycle units of just under 500k through 2027 with 33k CoWoS allocated, with this volume heavily 4Q 2027 weighted. Overall, this is a decently sized program considering it is a first-generation silicon effort.

New Project: Lyra, Customer TBC

We now see a new project codenamed “Lyra” which has significant wafer allocation of 44k CoWoS wafers in 2027, with which is 1.35 million units of package output. It is so far unknown what the project is and for which end customer, it could even be merchant silicon. We know that it’s 32 chips per CoWoS-S wafer so this is close to Hopper-class size wise. This project has appeared very abruptly and was able to tape out very quickly, suggesting it is derived from another design. “Lyra” has some relation to Greek mythology which is some indication that this could be for Meta. While details are still being figured out for Lyra, we incorporate this SKU as it accounts for nearly 10% of Broadcom's CoWoS allocation and should also be revenue generating (we assume ASP of $7-8k). Like with the discussion of Meta, this allocation could also potentially be diverted to support other chips such as TPU.

BROADCOM DEMAND	1Q26	2Q26	3Q26	4Q26	1Q27	2Q27	3Q27	4Q27	1Q28	2Q28	3Q28	4Q28
19 August 2026 update												
TPU v5p (Viperfish) Units	2,430	2,652	-	-	-	-	-	-	-	-	-	-
TPU 7x (Ironwood) - Internal Units	450,000	700,000	530,000	525,000	150,000	75,000	75,200	39,723	-	-	-	-
TPU 7x (Ironwood) - Anthropic Units	-	-	321,000	300,000	-	-	-	-	-	-	-	-
TPU 7x (Ironwood) Units	450,000	700,000	851,000	825,000	150,000	75,000	75,200	39,723	-	-	-	-
TPU 8i (Sunfish) - Internal Units	-	35,000	45,000	49,000	345,000	275,000	448,000	400,000	1,212,096	299,164	-	-
TPU 8i (Sunfish) - Anthropic Units	-	-	-	-	250,000	400,000	650,000	700,000	-	-	-	-
TPU 8i (Sunfish) Units	-	35,000	45,000	49,000	595,000	675,000	1,098,000	1,100,000	1,212,096	299,164	-	-
MTIA 300 (Athena) Units	27,000	5,000	5,000	2,116	-	-	-	-	-	-	-	-
MTIA 400 (Iris+) Units	5,000	1,860	-	2,000	13,000	13,000	45,000	150,000	150,000	17,740	-	-
MTIA 450 (Arke) Units	-	-	30	70	100	100	150	150	100	70	-	-
MTIA 500 (Astrid) Units	-	-	-	-	-	-	-	140	1,400	31,500	43,400	44,800
Lyra Units	-	-	-	40,000	120,000	250,000	350,000	450,000	237,390	-	-	-
OAI Jalapeno (Titan 1/Nexus 1) Units	-	1,000	5,000	5,000	18,000	22,000	65,000	250,000	80,000	43,750	-	-
OAI Titan 2 Units	-	-	-	-	-	-	-	-	1,400	25,750	35,750	42,500
OAI Titan 3 Units	-	-	-	-	-	-	-	-	-	-	-	-
TPU v5p (Viperfish) Revenue	13	15	-	-	-	-	-	-	-	-	-	-
TPU 7x (Ironwood) - Internal Revenue	5,974	9,293	7,036	6,969	2,340	1,170	-	-	-	-	-	-
TPU 7x (Ironwood) - Anthropic Revenue	-	-	4,261	3,983	-	-	-	-	-	-	-	-
TPU 7x (Ironwood) Revenue	5,974	9,293	11,297	10,952	2,340	1,170	-	-	-	-	-	-
TPU 8i (Sunfish) - Internal Revenue	-	525	675	735	6,693	5,335	8,691	7,760	21,090	5,205	-	-
TPU 8i (Sunfish) - Anthropic Revenue	-	-	-	-	4,850	7,760	12,610	13,580	-	-	-	-
TPU 8i (Sunfish) Revenue	-	525	675	735	11,543	13,095	21,301	21,340	21,090	5,205	-	-
TPU v9i (Whalefish) Revenue	-	-	-	-	-	-	-	-	5,205	1,840	-	-
TPU v9i (Pumafish) Revenue	-	-	-	-	-	-	-	-	2,013	24,156	37,576	38,583
TPU v10i Revenue	-	-	-	-	-	-	-	-	-	-	-	-
TPU v11i Revenue	-	-	-	-	-	-	-	-	-	-	-	-
MTIA 200 (Artemis) Revenue	40	40	40	40	0	0	0	0	0	0	0	0
MTIA 300 (Athena) Revenue	243	45	45	19	0	0	0	0	0	0	0	0
MTIA 400 (Iris+) Revenue	85	32	-	34	342	342	-	-	-	-	-	-
MTIA 450 (Arke) Revenue	-	-	1	2	3	3	5	5	3	2	-	-
MTIA 500 (Astrid) Revenue	-	-	-	-	-	-	-	6	56	1,263	1,740	1,796
MTIA 600 (Apollo) Revenue	-	-	-	-	-	-	-	-	-	-	-	-
Bytedance Custom ASIC Gen 1 Revenue	-	-	-	-	-	-	-	-	-	-	-	-
Bytedance Custom ASIC Gen 2 Revenue	-	-	-	-	-	-	-	-	0	0	0	0
Lyra Revenue	-	-	-	280	960	2,000	2,800	3,600	1,899	-	-	-
OAI Jalapeno (Titan 1/Nexus 1) Revenue	-	-	75	75	409	499	1,476	5,675	0	0	0	0
OAI Titan 2 Revenue	-	-	-	-	-	-	-	-	87	1,604	2,227	2,648
OAI Titan 3 Revenue	-	-	-	-	-	-	-	-	0	0	0	0
Broadcom AI ASIC Revenue	**6,359**	**9,956**	**12,152**	**12,206**	**15,953**	**17,465**	**25,937**	**30,982**	**31,518**	**35,057**	**42,382**	**43,865**
% change vs previous												
TPU v5p (Viperfish) Units	-	-	(100%)	(100%)	-	-	-	-	-	-	-	-
TPU 7x (Ironwood) - Internal Units	-	(13%)	(5%)	(28%)	(79%)	(75%)	-	-	-	-	-	-
TPU 7x (Ironwood) - Anthropic Units	-	-	-	-	-	-	-	-	-	-	-	-
TPU 7x (Ironwood) Units	-	(13%)	(3%)	(20%)	(79%)	(75%)	-	-	-	-	-	-
TPU 8i (Sunfish) - Internal Units	-	-	(36%)	(51%)	25%	(62%)	(47%)	(56%)	(7%)	-	-	-
TPU 8i (Sunfish) - Anthropic Units	-	-	-	-	-	-	-	-	-	-	-	-
TPU 8i (Sunfish) Units	-	-	(36%)	(51%)	13%	(40%)	(27%)	(31%)	(7%)	-	-	-
MTIA 300 (Athena) Units	(33%)	(78%)	(38%)	(53%)	-	-	-	-	-	-	-	-
MTIA 400 (Iris+) Units	-	(61%)	-	-	-	-	-	-	-	-	-	-
MTIA 450 (Arke) Units	-	-	(79%)	-	(95%)	(100%)	(99%)	(99%)	(98%)	-	-	-
OAI Jalapeno (Titan 1/Nexus 1) Units	-	-	-	-	(10%)	(27%)	117%	733%	611%	-	-	-
TPU v5p (Viperfish) Revenue	-	-	(100%)	(100%)	-	-	-	-	-	-	-	-
TPU 7x (Ironwood) - Internal Revenue	-	(13%)	(5%)	(28%)	(79%)	(75%)	-	-	-	-	-	-
TPU 7x (Ironwood) - Anthropic Revenue	-	-	-	-	-	-	-	-	-	-	-	-
TPU 7x (Ironwood) Revenue	-	(13%)	(3%)	(20%)	(79%)	(75%)	-	-	-	-	-	-
TPU 8i (Sunfish) - Internal Revenue	-	-	(36%)	(51%)	25%	(62%)	(47%)	(56%)	(7%)	-	-	-
TPU 8i (Sunfish) - Anthropic Revenue	-	-	-	-	-	-	-	-	-	-	-	-
TPU 8i (Sunfish) Revenue	-	-	(36%)	(51%)	13%	(40%)	(27%)	(31%)	(7%)	-	-	-
MTIA 300 (Athena) Revenue	(33%)	(78%)	(38%)	(53%)	-	-	-	-	-	-	-	-
MTIA 400 (Iris+) Revenue	-	(61%)	-	-	-	-	-	-	-	-	-	-
MTIA 450 (Arke) Revenue	-	-	(79%)	-	(95%)	(100%)	(99%)	(99%)	(98%)	-	-	-
OAI Jalapeno (Titan 1/Nexus 1) Revenue	-	-	(55%)	(55%)	(59%)	(67%)	(2%)	275%	(100%)	-	-	-
Broadcom AI ASIC Revenue	**(2%)**	**(13%)**	**(7%)**	**(21%)**	**(29%)**	**(40%)**	**(18%)**	**(8%)**	**(1%)**	**0%**	**-**	**-**

Source: SemiAnalysis

MediaTek TPU 8t Zebrafish Increases

We see more supply for MediaTek's TPU 8t Zebrafish with 10k additional CoWoS wafers in 2026 and 35k increase in 2027. Therefore, we increase our 2027 shipment forecast by 630k units for Zebrafish.

We keep total TPU v9 units unchanged across 2028 and 2029 but push back timing further, with Triggerfish starting 2 quarters later than Humufish which we anticipate begins shipping in 1Q28.

MEDIATEK DEMAND	**1Q26**	**2Q26**	**3Q26**	**4Q26**	**1Q27**	**2Q27**	**3Q27**	**4Q27**	**1Q28**	**2Q28**	**3Q28**	**4Q28**	**1Q29**	**2Q29**	**3Q29**	**4Q29**
19 August 2026 update																
TPU 8t (Zebrafish) Units	-	5,000	90,000	560,000	711,205	812,139	893,073	911,323	360,000	90,000	6,680	-	-	-	-	-
TPU v9t (Humufish) Units	-	-	-	-	-	-	-	-	110,000	500,000	725,000	725,000	512,500	460,000	191,594	-
TPU v9i (Triggerfish) Units	-	-	-	-	-	-	-	-	-	-	150,000	450,000	812,500	460,000	51,588	-
TPU v10t (Icefish) Units	-	-	-	-	-	-	-	-	-	-	-	-	720	6,120	12,960	16,920
TPU v11t Units	-	-	-	-	-	-	-	-	-	-	-	-	-	-	-	-
TPU Merope Units	-	-	-	-	-	-	-	-	-	10,000	70,000	153,000	170,000	180,000	120,000	120,000
TPU 8t (Zebrafish) ASP	-	4,000	4,000	4,000	4,500	4,500	4,500	4,500	4,500	4,500	4,500	4,500	4,500	-	-	-
TPU v9t (Humufish) ASP	-	-	-	-	-	-	-	12,000	12,000	12,000	12,000	12,000	12,000	12,000	12,000	12,000
TPU v9i (Triggerfish) ASP	-	-	-	-	-	-	-	12,000	12,000	12,000	12,000	12,000	12,000	12,000	12,000	12,000
TPU 8t (Zebrafish) Revenue	-	20	360	2,240	3,200	3,655	4,019	4,101	1,620	405	30	-	-	-	-	-
TPU v9t (Humufish) Revenue	-	-	-	-	-	-	-	-	1,320	6,000	8,700	8,700	6,150	5,520	2,299	-
TPU v9i (Triggerfish) Revenue	-	-	-	-	-	-	-	-	-	-	1,800	5,400	9,750	5,520	619	-
MediaTek Total Revenue	-	20	360	2,240	3,200	3,655	4,019	4,101	2,940	6,435	10,740	14,559	17,130	17,700	16,238	17,280
% change vs previous																
TPU 8t (Zebrafish) Units	0%	0%	0%	0%	16%	23%	25%	28%	100%	0%	0%	0%	0%	0%	0%	0%
TPU v9t (Humufish) Units	0%	0%	0%	0%	0%	0%	0%	0%	-27%	5%	49%	35%	-9%	48%	271%	0%
TPU v9i (Triggerfish) Units	0%	0%	0%	0%	0%	0%	0%	0%	-100%	-100%	-69%	-16%	44%	48%	0%	0%
TPU 8t (Zebrafish) Revenue	0%	0%	0%	0%	16%	23%	25%	28%	100%	0%	0%	0%	0%	0%	0%	0%
TPU v9t (Humufish) Revenue	0%	0%	0%	0%	0%	0%	0%	0%	-27%	5%	49%	35%	-9%	48%	271%	0%
TPU v9i (Triggerfish) Revenue	0%	0%	0%	0%	0%	0%	0%	0%	-100%	-100%	-69%	-16%	44%	48%	0%	0%
MediaTek Total Revenue	**0%**	**0%**	**0%**	**0%**	**16%**	**23%**	**25%**	**28%**	**-33%**	**-46%**	**-10%**	**9%**	**16%**	**26%**	**12%**	**0%**

Source: SemiAnalysis

Nvidia Changes

We make a host of changes for Nvidia:

Rubin and Rubin Ultra:

We see 2027 CoWoS for Nvidia at 1.2mn wafers, higher than our previous at 1mn wafers. The majority of the incremental is from more Rubin, most of which consists of Rubin Ultra. There is a very rapid transition from Rubin to Rubin Ultra (GR150) from 2Q onwards. Given very few differences at the package level except for different HBM (8-hi over 12-hi), we see this as very achievable at the upstream manufacturing stage. For Rubin Ultra, the big changes will be at the system level with the introduction of the NVL576 architecture, and that may prove to be a challenge that could push back the Rubin Ultra ramp.

Therefore, we increase our mix of Rubin Ultra by 1.2mn units, however, our total Rubin and Rubin Ultra shipments for CY27 stays at 7mn units.

For CY2026, we also pare back our expectations for the beginning of the Rubin shipments for Q3 by 50k units.

Vera:

For CY27, our Vera upstream forecast is reduced to 4.8mn units from 5.5mn units previously and we therefore reduce our standalone Vera shipments from just over 2mn units to 1.5mn units. For 2026 we also push back the start of Vera standalone shipments to 4Q.

H200:

We push out some H200 shipments, reducing 3Q/4Q CY26 by 100k/50k units, respectively, and instead place them into 1Q/2Q CY27. While some H200 shipments to China have been formally approved (around 200k units) by the Chinese government, we understand that there are shipments being shipped to Chinese customers but outside of China.

Groq LP30/LP35

We reduce the early LPU shipment ramp with L6 board assembly now expected to start in 4Q CY26, rather than starting from 3Q CY26 as previously assumed.

Feynman and Beyond

Following on from the HBM downgrade with Rubin Ultra, we now assume Feynman will use 8-hi HBM4E instead of 16-hi previously assumed. We therefore reduce our Feynman ASP assumption to $110k per chip from $143k. We remind subscribers that this is still a very preliminary estimate given the uncertainty around the content. Similarly for the generation after Feynman Ultra we assume it's 8-hi (ASP readjusted to $175k per unit from $190k) and Feynman Next is a downgrade to 12-hi from 16-hi. We also increase Feynman units by 460k and 430k in 2029 and 2030, and Feynman Ultra by 100k units in 2030.

Please see the detailed changed below:

Source: SemiAnalysis

Qualcomm-ByteDance ASIC Back / Potential TikTok ASIC

We previously wrote about the suspension of ByteDance's ASIC program with Qualcomm. Industry chatter suggests that the program may be back on track as of August. In addition, there is a possible TikTok-specific ASIC that is independent of this ByteDance chip. It is said that TikTok is in talks with Marvell and Qualcomm, with TikTok completing import/export compliance and forming a joint entity with Oracle for AI infrastructure. Earlier this year, Oracle took a 15% stake in TikTok's U.S. operations under the TikTok USDS Joint Venture. This could be a way for them to access TSMC's advanced logic. Marvell's involvement could see ByteDance/TikTok emerge as a potential customer for NVLink Fusion as well.

AMD Update

Post Earnings AMD Adjustment

We remove 30k MI308 units from 2Q26 and 3Q26, as Chinese government restrictions have blocked sales into China. Offsetting this, we pull in 10k MI350 units from 4Q26 into 3Q26, add 10k MI325 units to 3Q26, and pull in 20k MI400 units from 1H27 into 3Q26.

Source: SemiAnalysis

Huawei Shipment Update

In our previous China Accelerator update, we expected Huawei to capture majority of the China market with ~1.2mn unit shipments in 2026 of which ~440k were Ascend 910 series and ~760k were Ascend 950 series (~80k Ascend 950DT and ~680k Ascend 950PR). We are now revising the number of Ascend 950 series up to ~910k units (~150k Ascend 950DT and ~760k Ascend

950PR). This will put Huawei's 2026 Ascend shipments at ~1.35mn units and revenue at ~Rmb 120Bn (up from ~Rmb 110Bn).

Industry chatter suggests that ByteDance will be Huawei's single largest external customer representing ~25% of external shipment volume, with Tencent and Alibaba representing ~25% together. RedNote (小红书)，Meituan, Kuaishou and Chinese AI Labs will represent ~25%, with the final ~25% split amongst SOEs and Telecom companies. It remains unclear how much is being consumed internally, but we understand most should be sold externally.

Source: SemiAnalysis

HBM Mix Change

- Rubin Ultra 8-Hi (192GB): moved off a flat 40% SK Hynix / 40% Samsung / 20% Micron split to roughly 30% SK Hynix / 37% Samsung / 34% Micron
- B300, Rubin, and TPU v8t: adjusted the supplier mix on these SKUs to better match HBM revenues

Install Base Changes

- Added 400k MI400/450 units starting from 2Q 2027 to Anthropic
- Added Kimi AI
- Raised DeepSeek's 2026 Ascend 950 installed base to ~56k (from ~16k) of which ~40k is Ascend 950PR and ~16k is Ascend 950DT.

Download Model

Disclaimers

Analyst Certifications and Independence of Research.

Each of the analysts whose names appear in this report hereby certify that all the views expressed in this Report accurately reflect our personal views about any and all of the subject securities or issuers and that no part of our compensation was, is, or will be, directly or indirectly, related to the specific recommendations or views in this Report. SemiAnalysis LLC (the "Company") is an independent equity research provider. The Company is not a member of the Financial Industry Regulatory Authority or the Securities Investor Protection Corporation, and it is not a registered broker dealer or investment adviser.

Private Investment Disclosure.

Each analyst and/or author contributing to this Report hereby discloses that they may hold, directly or indirectly, private investments in companies, sectors, or asset classes discussed or referenced herein. Such holdings, if any, are disclosed externally with a third-party vendor. The Company requires that no author or analyst allow their personal investment positions to influence the views, opinions, analyses, or conclusions expressed in this Report. Readers should be aware that the existence of such holdings, even where disclosed, may present a potential conflict of interest, and are encouraged to weigh this accordingly when evaluating the contents of this Report.

Limitation of Research and Information.

This Report has been prepared for distribution to only qualified institutional or professional clients of the Company. The contents of this Report represent the views, opinions, and analyses of its authors. The information contained herein does not constitute financial, legal, tax or any other advice. All third-party data presented herein were obtained from publicly available sources which are believed to be reliable; however, the Company makes no warranty, express or implied, concerning the accuracy or completeness of such information. In no event shall the Company be responsible or liable for the correctness of, or update to, any such material, or for any direct, indirect, incidental or consequential loss, damage or lost opportunities (including loss of profits, revenue or goodwill) resulting from use of this data or from any investment decisions based on information or research obtained from the Company. Nothing contained in this Report or any distribution by the Company should be construed as any offer to sell, or any solicitation of an offer to buy, any security or investment. Any research or other material received should not be construed as individualized investment advice. Investment decisions should be made as part of an overall portfolio strategy and you should consult with a professional financial advisor, legal and tax advisor prior to making any investment decision.

Reproduction and Distribution Strictly Prohibited.
No user of this Report may reproduce, modify, copy, distribute, sell, resell, transmit, transfer, license, assign or publish the Report itself or any information contained therein. Notwithstanding the foregoing, clients with access to working models are permitted to alter or modify the information contained therein, provided that it is solely for such client's own use. This Report is not intended to be available or distributed for any purpose that would be deemed unlawful or otherwise prohibited by any local, state, national or international laws or regulations or would otherwise subject the Company to registration or regulation of any kind within such jurisdiction.

Copyrights, Trademarks, Intellectual Property.
The Company, and any logos or marks included in this Report are proprietary materials. The use of such terms and logos and marks without the express written consent of the Company is strictly prohibited. The copyright in the pages or in the screens of the Report, and in the information and material therein, is proprietary material owned by the Company unless otherwise indicated. The unauthorized use of any material on this Report may violate numerous statutes, regulations and laws, including, but not limited to,

copyright, trademark, trade secret or patent laws.

ETF and Regulatory Status Disclaimer.
From time to time, The Company may collaborate with asset managers, issuers, sponsors, index providers or other financial institutions in connection with the development, launch or support of exchange-traded funds ("ETFs") and other investment products, including by providing research, market analysis, index-related research or other research related services. The Company does not sponsor, endorse, sell or promote any ETF, and makes no representation or warranty, express or implied, regarding the advisability of investing in any such ETF. In connection with the services described herein, the Company does not act as a fiduciary to any investor, fund, sponsor, issuer or other party. The Company is not registered with the U.S. Securities and Exchange Commission or any state securities authority as an investment adviser or exempt reporting adviser. The Company's role is limited to providing research and related services, and it does not provide individualized investment advice, portfolio management, brokerage services or other services requiring registration as an investment adviser.

// RELATED RESEARCH

Keep reading.

August 14, 2026 ACCELERATOR & HBM

2026 TPU Output Reduction, 2027 Broadcom CoWoS 460k But Lower TPU Allocation than Expected, AMD TPU Involvement

// Accelerator Model Update – 14 August 2026

By Myron Xie and Ivan Chiam

August 14, 2026 ACCELERATOR & HBM MEMORY MODEL

Leading AI Labs and Hyperscalers Want 4-hi HBM

// SemiAnalysis Multi-Vertical Note – 14 August 2026

By Myron Xie and Dylan Patel

August 10, 2026 ACCELERATOR & HBM AI NETWORKING MODEL CORE RESEARCH

Rubin Ultra NVL576 Architecture: Flash Overview

By Tony Chien, Wega Chu, and Myron Xie