

摘要：

Claude被曝“降智+涨价”引发重度用户不满：Anthropic模型推理能力断崖下滑，第三方测试排名从前2跌至第10，复杂任务表现明显退化。官方承认为控成本默认降至“中等努力”模式，同时隐藏思考过程、缩短缓存，加剧体验缩水。另一边企业订阅从包月转向按量计费，费用最高翻三倍。在算力压力下限用高耗工具之际，Opus 4.7与AI设计工具将上线，冲击Adobe、Figma等玩家。

Claude性能降智，已经成了重度用户们绕不开的槽点。

而且发现这一点的，还是AMD高级总监Stella Laurenzo。

在翻烂了近7000份会话日志后直言，Stella发现现在的Claude在处理复杂任务时已经变傻了。



Featured

Is Anthropic 'nerfing' Claude? Users increasingly report performance degradation as leaders push back

Carl Franzen

Published 10:11 am, PT, April 13, 2026

Updated 1:33 pm, PT, April 13, 2026

公众号 · 量子位

第三方跑分数据也印证了这一点，原本稳坐前2名的尖子生，短短几周内排名直接跌到了榜单第10位。

眼看着交了同样的订阅费却只能用上这种减配版模型，不少用户都感觉自己被当成了冤大头。

Claude降智被抓包

前面提到的Stella，对Claude Code的运行数据进行了详细审计。

调查范围覆盖了6852份会话文件、17871个思考块以及超过23万次工具调用。

[MODEL] Claude Code is unusable for complex engineering tasks with the Feb updates #42796

New issue

Closed



stellaraccident

opened 2 weeks ago · edited by stellaraccident

Edits ...

Assignees

No one assigned

公众号 · 量子位

分析结果显示，从2月份开始，模型的推理深度出现了断崖式下滑。

在实际操作中，它在处理复杂工程任务时表现得不再可靠，经常出现逻辑断层或无法理解复杂指令的情况。



原本那个愿意深钻难题的助手现在变得非常敷衍，开始倾向于用最简单的修补方案来应付差事，甚至频繁在任务进行到一半时莫名其妙地停下来。

这种退步在第三方基准测试BridgeBench中表现得更加直观，Opus 4.6的准确率在短时间内从83.3%跌到了68.3%。

由于准确率的大幅缩水，它在测试榜单上的排名也从第二名直接掉到了第十名。



显示翻译

CLAUDE OPUS 4.6 IS NERFED.

BridgeBench just proved it.

Last week Claude Opus 4.6 ranked #2 on the Hallucination benchmark with an accuracy of 83.3%.

Today Claude Opus 4.6 was retested and it fell to #10 on the leaderboard with an accuracy of only 68.3%.

A 98% increase in hallucination.

bridgebench.ai just confirmed that Claude Opus 4.6 has reduced reasoning levels and is nerfed.

Hallucination
How often do AI models fabricate false claims when analyzing code? 30 tasks across 6 clusters with 175 questions — verified by code execution and ground truth. Updated: 2024-04-22

RANK	MODEL	SCORE	ACCURACY	FABRICATION
1	Grok 4.20 Reasoning ai/grok-4.20-reasoning	91.8	98.0%	18.0% →
2	Claude Opus 4.6 anthropic/claude-opus-4.6	87.6	83.3%	16.7% →
3	GPT-5.4 openai/gpt-5.4	86.1	83.3%	16.7% →
4	Qwen 3.6 Plus qwen/qwen3.6-plus	79.7	75.2%	26.5% →
5	Gemini 3.1 Pro google/gemini-3.1-pro-preview	79.1	74.7%	26.7% →
6	Qwen3.5 Plus 2024-02-15 qwen/qwen3.5-plus-02-15	77.3	72.2%	29.8% →
7	Claude Sonnet 4.6 anthropic/claude-sonnet-4.6	76.6	72.4%	28.9% →
8	Grok 4.20 (Non-Reasoning) ai/grok-4.20	76.1	71.9%	29.7% →
9	Gemini 3 Pro google/gemini-3-pro-preview	75.9	71.9%	30.8% →
10	Claude Opus 4.6 (April 12) anthropic/claude-opus-4.6-apr12	73.3	68.3%	

对于这些质疑，官方负责人Boris Cherny承认，目前模型默认开启的是努力程度为85分的“中等努力”模式。

官方给出的解释是，这种调整是为了平衡处理速度和算力成本，给大多数用户提供更快的响应。





bcherry last week · Pinned by bcherry

Collaborator

Hi, thanks for the detailed analysis. Before I keep going, I wanted to say I appreciate the depth of thinking & care that went into this.

There's a lot here, I will try to break it down a bit. These are the two core things happening:

redact-thinking-2026-02-12

This beta header hides thinking from the UI, since most people don't look at it. It **does not** impact thinking itself, nor does it impact thinking budgets or the way extended reasoning works under the hood. It is a UI-only change.

Under the hood, by setting this header we avoid needing thinking summaries, which reduces latency. You can opt out of it with `showThinkingSummaries: true` in your settings.json (see [docs](#)).

If you are analyzing locally stored transcripts, you wouldn't see raw thinking stored when this header is set, which is likely influencing the analysis. When Claude sees lack of thinking in transcripts for this analysis, it may not realize that the thinking is still there, and is simply not user-facing.

Thinking depth had already dropped ~67% by late February

We landed two changes in Feb that would have impacted this. We evaluated both carefully:

1/ Opus 4.6 launch → adaptive thinking default (Feb 9)

Opus 4.6 supports adaptive thinking, which is different from thinking budgets that we used to support. In this mode, the model decides how long to think for, which tends to work better than fixed thinking budgets across the board.

`CLAUDE_CODE_DISABLE_ADAPTIVE_THINKING` to opt out.

2/ Medium effort (85) default on Opus 4.6 (Mar 3)

We found that effort=85 was a sweet spot on the intelligence-latency/cost curve for most users, improving token efficiency while reducing latency. One of our product principles is to avoid changing settings on users' behalf, and ideally we would have set effort=85 from the start. We felt this was an important setting to change, so our approach was to:

1. Roll it out with a dialog so users are aware of the change and have a chance to opt out
2. Show the effort the first few times you opened Claude Code, so it wasn't surprising.

Some people want the model to think for longer, even if it takes more time and tokens. To improve intelligence more, set `effort=high` via `/effort` or in your settings.json. This setting is sticky across sessions, and can be shared among users. You can also use the `ULTRATHINK` keyword to use high effort for a single turn, or set `/effort max` to use even higher effort for the rest of the conversation.

Going forward, we will test defaulting Teams and Enterprise users to high effort, to benefit from extended thinking even if it comes at the cost of additional tokens & latency. This default is configurable in exactly the same way, via `/effort` and settings.json.

79 449 12 1 27 5 25

bcherry closed this as **completed** last week

公众号 · 量子位

但在很多重度用户看来，这种为了省钱而限制模型智能水平的做法，实质上就是一种变相的性能削减。

为了在视觉上掩盖这种变化，程序界面还在2月份悄悄隐藏了具体的思考过程，导致很多习惯观察模型逻辑的用户觉得它根本没在认真动脑子。

此外，后台的技术细节也暴露了一些缩减成本的痕迹，有开发者发现提示词缓存的时长从1小时被骤降到了5分钟。

这种改动意味着在长会话中，只要用户稍微停顿思考，之前生成的缓存就会过期，导致额度消耗速度大幅提升。

这种性能表现和计费规则的双重缩水，让那些需要长时间处理复杂代码的程序员感觉非常受挫。

不过，虽然智商降下来了，可是价格在涨啊。

企业包月套餐涨价

Anthropic对计费规则的大幅度调整让很多企业用户措手不及，原本那种一口价包月畅用的模式已经被彻底拆分。

现在的Claude Enterprise计费方案，更像是一种流量计费——

用户除了要交20美元的固定底薪，剩下的每一份算力开销都要根据实际用量额外结账。

这种计费转变对重度用户来说无异于一场金钱噩梦。



曾经每月只需支付200美元固定费用的公司，在新的规则下极有可能发现账单翻了倍，甚至有部分团队的支出直接暴涨到了原来的三倍。

A社之所以选择在此时跟客户算细账，主要还是因为扛不住成本的压力。



Boris Cherny
@bcherny



显示翻译

We've been working hard to meet the increase in demand for Claude, and our subscriptions weren't built for the usage patterns of these third-party tools. Capacity is a resource we manage thoughtfully and we are prioritizing our customers using our products and API.

上午7:14 · 2026年4月4日 · 42.6万 查看

公众号 · 量子位

由于模型推理成本在过去一年翻了三倍，管理层为了保住毛利率，只能选择从这些高频用户身上寻找盈利空间。

甚至因为算力资源实在捉襟见肘，A社开始限制订阅用户去调用OpenClaw这种极度消耗资源的Agent工具。

看着A社的涨价降智风波，隔壁OpenAI立刻推出了100美元档位的Codex订阅服务。

这种定价策略，或许正是想趁机收割那些对涨价感到不满的优质客户。

Opus 4.7要来了

根据爆料，Anthropic最快在本周就会把Claude Opus 4.7摆上货架。

这款新旗舰模型目前已经准备就绪，预计将作为产品线的最新性能代表直接上线。

同步曝光的还有一款AI设计工具，它能让用户通过自然语言指令直接生成网页、演示文稿、落地页或者产品原型。

这款工具的门槛设得很低，无论是敲代码的还是完全不懂技术的小白，都能通过提要求的方式完成设计。

受此消息影响，Adobe、Wix和Figma这些设计软件巨头的股价在报告发布后的几小时内应声下跌了2%以上。

这种还没见着真家伙就先让友商市值蒸发的表现，确实给市场留足了想象空间。

本文来源：[量子位](#)

风险提示及免责声明

市场有风险，投资需谨慎。本文不构成个人投资建议，也未考虑到个别用户特殊的投资目标、财务状况或需要。用户应考虑本文中的任何意见、观点或结论是否符合其特定状况。据此投资，责任自负。



限时权益申领

* 体验资格每人限申领一次

扫码
免费领取

VIP会员·7天畅读卡 | 大师课·入门串讲



限免

87 收藏

分享到:

写评论

请发表您的评论

表情 图片

发表评论

1 条评论

李慕白
12小时前 中国河南省

一直赔本做买卖吗？用户越多亏损越严重？

0 回复