

继2025年12月小米在人车家全生态合作伙伴大会上发布自研AI大模型Xiaomi MiMo-V2-Flash后，小米再次提速。

3月19日，小米发布旗下首个全模态基座模型 Xiaomi MiMo-V2-Omni。

MiMo-V2-Omni 被设计为一个具备跨模态感知与 GUI（图形用户界面）操作能力的“执行者”，可无缝接入各种 Agent 框架。

此前，该模型以“Healer Alpha”的代号在 OpenRouter 平台上进行盲测，并在各项基准测试中显示出对标甚至部分赶超头部闭源模型的表现。

对于该模型高效的“上新速度”，雷军表示：“我们在AI领域上相对比较低调，实际进展可能比大家看到的要快很多。在AI领域，我们今年的研发和资本投入就将超过160亿元。我相信，只要我们坚持持续投入，小米在AI时代一定会交出一份靓丽的答卷。”

作为该模型核心负责人，罗福莉亦是在海外社交平台上直言：“在明天之前，MiMo 团队里对话测试不足100次的人，可以直接离职。这招奏效了。一旦团队的想象力被智能体系统的能力点燃，这份想象力就直接转化成了研发速度。”

目前小米给出了输入\$0.4 /百万 tokens、输出\$2 / 百万 tokens（支持256K上下文）的API报价。

小米的野心显然不止于向开发者售卖API。

该模型目前已与金山办公（WPS）达成合作，探索文本生成与结构化数据处理的场景。

但从战略纵深来看，MiMo-V2-Omni的商业终局指向小米的“人车家全生态”。

在关于MiMo-V2-Omni 的未来畅想中，小米亦表示“将持续推进长周期智能体规划、实时流式感知、多智能体协同，以及与物理世界更深层的整合。”

如果该模型能作为底层“大脑”深度整合至小米澎湃 OS（HyperOS）中，真正打造一个能够跨端深度理解语音指令、自主调用手机 App 甚至控制小米车机界面的AI基座，将极大程度地提升小米硬件的溢价能力和用户留存率。

尽管技术演示与生态愿景极具吸引力，但目前小米正面临严峻的资源分配与成本控制考验。

当前的小米正处于高压的“多线作战”状态：

一方面，作为现金牛的智能手机业务正遭遇上游存储芯片价格暴涨的逆风，硬件综合毛利率备受挤压；另一方面，汽车业务又处于产能爬坡与全国销售网络扩张的关键期，亟需持续投入。

不仅如此，相比于手握丰厚利润率且拥有庞大云计算底座的纯互联网巨头，小米在AI军备竞赛中的资金筹码并不占优。

从战略愿景来看，MiMo-V2-Omni 无疑是小米补齐“人车家全生态”智能化闭环的最关键一块拼图。

在内存涨价的逆风局中，如何平衡好手机、汽车与大模型基座的“多线投入”，考验着小米管理层的智慧。



Fuli Luo
@_LuoFuli



MiMo-V2-Pro & Omni & TTS is out. Our first full-stack model family built truly for the Agent era.

I call this a quiet ambush — not because we planned it, but because the shift from Chat to Agent paradigm happened so fast, even we barely believed it. Somewhere in between was a process that was thrilling, painful, and fascinating all at once.

The 1T base model started training months ago. The original goal was long-context reasoning efficiency. Hybrid Attention carries real innovation, without overreaching — and it turns out to be exactly the right foundation for the Agent era. 1M context window. MTP inference for ultra-low latency and cost. These architectural decisions weren't trendy. They were a structural advantage we built before we needed it.

What changed everything was experiencing a complex agentic scaffold — what I'd call orchestrated Context — for the first time. I was shocked on day one. I tried to convince the team to use it. That didn't work. So I gave a hard mandate: anyone on MiMo Team with fewer than 100 conversations tomorrow can quit. It worked. Once the team's imagination was ignited by what agentic systems could do, that imagination converted directly into research velocity.

People ask why we move so fast. I saw it firsthand building DeepSeek R1. My honest summary:

- Backbone and Infra research has long cycles. You need strategic conviction a year before it pays off.
- Posttrain agility is a different muscle: product intuition driving evaluation, iteration cycles compressed, paradigm shifts caught early.
- And the constant: curiosity, sharp technical instinct, decisive execution, full commitment — and something that's easy to underestimate: a genuine love for the world you're building for.

We will open-source — when the models are stable enough to deserve it.

From Beijing, very late, not quite awake.

风险提示及免责声明

市场有风险，投资需谨慎。本文不构成个人投资建议，也未考虑到个别用户特殊的投资目标、财务状况或需要。用户应考虑本文中的任何意见、观点或结论是否符合其特定状况。据此投资，责任自负。

限时 0 元领

* 同一用户免费领取一次

扫码
免费领取



限免

VIP会员·7天畅读卡 | 大师课·入门串讲



写评论

请发表您的评论

 表情

 图片

发表评论