

摘要：

Xiaomi MiMo大模型负责人罗福莉指出，OpenClaw产生的请求数是Claude原生框架的数倍；换成API计费，成本将是订阅价格的数十倍。Anthropic封堵第三方订阅接入，让低效的代价暴露，将倒逼开发者精进上下文管理，削减无效token消耗。

Anthropic封堵第三方工具滥用订阅权限的举动，正在揭开AI Agent时代一场被长期忽视的成本危机。

两天前，Anthropic宣布切断第三方调用框架通过Claude订阅接入的通道。Xiaomi MiMo大模型负责人罗福莉（Fuli Luo）随即发文，结合三天前MiMo刚刚推出的Token计划，对这一事件作出深度解读。

她认为，Anthropic此举并非单纯的商业防御，而是在全球算力供给无法追上Agent需求增速的背景下，整个生态系统走向成熟的一个必要节点。

此次受直接冲击的，是OpenClaw、OpenCode等依赖Claude订阅权限运行的第三方调用框架用户。这些用户面临的成本骤升，短期内可能达到此前的数十倍。

但罗福莉认为，这种压力恰恰是倒逼工程质量改善的催化剂——只有让效率低下的成本真实可见，才能推动开发者认真对待上下文管理与缓存优化。

订阅制度背后的成本黑洞

罗福莉指出，Claude Code的订阅体系在算力分配设计上颇为精巧，但她坦言，这套系统大概率无法盈利，甚至可能处于亏损状态。

问题的根源在于第三方框架的调用方式。以OpenClaw为例，其上下文管理存在明显缺陷：在处理单个用户请求时，系统会拆分出多轮低价值工具调用，以独立API请求的形式逐一发出，每次请求携带的上下文窗口往往超过10万个token。

即便存在缓存命中，这种模式也极为浪费；在极端情况下，还会拉高其他请求的缓存未命中率。

罗福莉估算，这类框架每个查询实际产生的请求数，往往是Claude Code原生框架的数倍。换算成API计费，实际成本可能是订阅价格的数十倍。她将这一差距称为“不是缺口，而是天坑”。

Two days ago, Anthropic cut off third-party harnesses from using Claude subscriptions — not surprising. Three days ago, MiMo launched its Token Plan — a design I spent real time on, and what I believe is a serious attempt at getting compute allocation and agent harness development right. Putting these two things together, some thoughts:

1. Claude Code's subscription is a beautifully designed system for balanced compute allocation. My guess — it doesn't make money, possibly bleeds it, unless their API margins are 10-20x, which I doubt. I can't rigorously calculate the losses from third-party harnesses plugging in, but I've looked at OpenClaw's context management up close — it's bad. Within a single user query, it fires off rounds of low-value tool calls as separate API requests, each carrying a long context window (often >100K tokens) — wasteful even with cache hits, and in extreme cases driving up cache miss rates for other queries. The actual request count per query ends up several times higher than Claude Code's own framework. Translated to API pricing, the real cost is probably tens of times the subscription price. That's not a gap — that's a crater.

AI工作坊组织者@newlinedotco评论称：订阅的“all you can eat”套餐从一开始就是定时炸弹——第三方harness（如OpenClaw）24/7循环调用，API成本可能高达5000美元，而订阅仅200美元。官方工具（如Claude Code）因优化prompt cache才可持续。

the all you can eat arbitrage on these subscriptions was always a ticking time bomb. when you have agents running 24/7 loops that would cost \$5k on the api but only \$200 on a max plan, the margins just evaporate.

anthropic isn't banning openclaw, they're just moving it to the api ledger where the compute is actually paid for. first-party tools like claude code and cowork are built to maximize prompt cache hits, which is the only reason those flat-rate tiers are even sustainable for them. third-party harnesses that bypass those efficiencies are essentially backing a truck up to a buffet.

the branding definitely should've been clearer from day one. if they'd just messaged it as "subscriptions cover official apps only," half of this blowback wouldn't exist.

由 Google 翻译自 英语

这些订阅服务的无限量套利一直是个定时炸弹。当代理商 24/7 全天候运行循环任务时，如果使用 API 开发，成本高达 5000 美元，而最高套餐只需 200 美元，利润空间就会瞬间消失殆尽。

Anthropic 并没有禁用 OpenClaw，他们只是将其转移到 API 账本，在那里计算费用才是真正支付的。像 Claude Code 和 Cowork 这样的第一方工具旨在最大化缓存命中率，这也是它们能够维持固定费率模式的唯一原因。绕过这些效率提升的第三方工具，就好比是开着卡车去自助餐摊位。

品牌定位从一开始就应该更清晰。如果他们当初直接说明“订阅仅涵盖官方应用”，那么现在一半的负面评价就不会出现了。

Anthropic此次调整并未完全关闭第三方接入的大门。OpenClaw、OpenCode等工具仍可通过API调用Claude，只是失去了搭乘订阅计划的通道。

这一区别至关重要。对于习惯以订阅价格使用这些工具的用户而言，**成本冲击是即时且显著的**。

但罗福莉认为，这种痛感具有修正意义——**它会迫使框架开发者认真提升上下文管理能力，最大化提示缓存命中率以复用已处理的上下文，并削减无效token消耗。她将这一过程描述为“痛苦最终转化为工程规律”。**

2. Third-party harnesses like OpenClaw/OpenCode can still call Claude via API — they just can't ride on subscriptions anymore. Short term, these agent users will feel the pain, costs jumping easily tens of times. But that pressure is exactly what pushes these harnesses to improve context management, maximize prompt cache hit rates to reuse processed context, cut wasteful token burn. Pain eventually converts to engineering discipline.

她同时提醒各大语言模型公司，在尚未厘清编程计划成本结构之前，不应盲目卷入定价竞底。廉价出售token同时对第三方框架敞开大门，表面上对用户友好，实则是一个陷阱——而Anthropic刚刚走出了这个陷阱。

她进一步指出，若用户在低质量框架、不稳定推理服务和降级模型上消耗大量精力后仍无所获，对用户体验和留存都将造成实质损害。

3. I'd urge LLM companies not to blindly race to the bottom on pricing before figuring out how to price a coding plan without hemorrhaging money. Selling tokens dirt cheap while leaving the door wide open to third-party harnesses looks nice to users, but it's a trap — the same trap Anthropic just walked out of. The deeper problem: if users burn their attention on low-quality agent harnesses, highly unstable and slow inference services, and models downgraded to cut costs, only to find they still can't get anything done — that's not a healthy cycle for user experience or retention.

对此，AI 工程师 @karpathy指出：

“伟大的软件往往诞生于限制之中。如果 Token 是免费的，没人会去写精简的 Prompt 或研究 Context 压缩；当成本变成瓶颈，开发者才会真正思考如何构建‘有脑子’的 Agent。”

MiMo Token计划的不同路径

在解读Anthropic举措的同时，罗福莉也阐明了MiMo Token计划的设计逻辑。

该计划支持第三方调用框架，采用按token配额计费的方式，逻辑上与Claude新近推出的超额用量包一致。

罗福莉强调，MiMo的目标是“**长期稳定地交付高质量模型和服务，而非让用户冲动付费后流失**”。

4. On MiMo Token Plan — it supports third-party harnesses, billed by token quota, same logic as Claude's newly launched extra usage packages. Because what we're going for is long-term stable delivery of high-quality models and services — not getting you to impulse-pay and then abandon ship.

这一表述折射出一种与订阅制度不同的算力分配哲学：通过真实的用量成本约束用户和框架的行为，而非以封闭方式管理滥用风险。

效率竞争，而非算力消耗

罗福莉在文章结尾给出了一个更宏观的判断：全球算力供给已无法跟上Agent产生的token需求增速。

在她看来，出路不在于进一步压低token价格，而在于“更高token效率的Agent框架”与“更大、更高效的模型”之间的协同演进。

Anthropic的这次举措，无论是否出于主动设计，都在将整个生态系统——无论开源还是闭源——推向这一方向。

"Agent时代不属于烧算力最多的人，而属于用得最聪明的人。"罗福莉写道。

The bigger picture: global compute capacity can't keep up with the token demand agents are creating. The real way forward isn't cheaper tokens — it's co-evolution. "More token-efficient agent harnesses" × "more powerful and efficient models." Anthropic's move, whether they intended it or not, is pushing the entire ecosystem — open source and closed source alike — in that direction. That's probably a good thing. The Agent era doesn't belong to whoever burns the most compute. It belongs to whoever uses it wisely.

风险提示及免责条款

市场有风险，投资需谨慎。本文不构成个人投资建议，也未考虑到个别用户特殊的投资目标、财务状况或需要。用户应考虑本文中的任何意见、观点或结论是否符合其特定状况。据此投资，责任自负。

限时 0 元 领

* 同一用户免费领取一次

扫码 免费领取



限免

152 收藏

分享到:

写评论

请发表您的评论

表情

图片

发表评论