

Allianz Research | 12 May 2026

Behind the gate: The promises and perils of private markets democratization

Ludovic Subran
Chief Investment Officer and
Chief Economist
ludovic.subran@allianz.com

Yao Lu
Investment Strategist
yao.lu@allianz.com

Jordi Basco Carrera
Head of Private Markets
jordi.basco_carrera@allianz.com

Nils Bradtke
Senior Investment Strategist
nils.bradtke@allianz.com

America Hernandez Ortiz
Senior Investment Strategist
america.hernandez@allianz.com


In Summary

- **Private markets are opening up to retail, and the product has to be rebuilt to fit.** Global private markets have grown more than twentyfold since 2000 to over USD17trn, propelled by institutional adoption of the Yale endowment model that tilted long-duration capital aggressively into illiquid assets. The next leg of growth runs through wealth channels and mass markets, where investors lack the governance, illiquidity tolerance and manager-selection capability. Semi-liquid evergreens have become the workhorse for wealth, while the mass market is reached primarily via DC plans, regulated retail funds, public-private hybrids, and insurance wrappers, structures in which a fiduciary, insurer, regulatory template or the product design itself makes the call rather than the end investor.
- **Three forces have converged to push democratization: savers need enhanced returns and diversification, governments need to mobilize private savings for public priorities and asset managers need new pools.** Diversified private-market exposure can improve portfolio efficiency, lifting expected returns from 6.2% to 7.9% while keeping stressed downside risk (CVaR) around -3% versus -5% for traditional liquid portfolios. Governments see a policy lever to address a USD400trn retirement gap by 2050, growing inequality as value creation stays private and infrastructure needs that public budgets cannot fund. For asset managers, the move is increasingly necessary: institutional fundraising is slowing as the traditional capital pool matures and large LPs reach their target allocations.
- **The pitch is real, but a closer look reveals a more nuanced picture.** First, the return story is compelling at face value: buyout funds have delivered 12-14% annually against 7-9% for public equities. But strip out leverage and most of the private-equity advantage disappears, leaving manager selection as the real driver of outperformance. Second, retail investors face a compounding triple disadvantage compared to institutional products: additional fee layers that erode the illiquidity premium, potentially weaker underlying assets, and no ability to select managers in an asset class where dispersion runs wide. Third, the semi-liquid promise is not illusory, but it comes with conditions; it works in normal market conditions but exposes investors to gates precisely when they want their capital back. Finally, retail money brings reactive behavior into an asset class increasingly tied to banks and life insurers, reintroducing the maturity transformation that closed-end structures had engineered away.
- **The early-2026 slowdown is a stress test, but structural drivers remain intact.** Expect product recalibration toward larger liquid buffers and tighter gating, consolidation toward mega-GPs with the scale and distribution to run retail wrappers and a new round of regulatory scrutiny on disclosure, leverage and liquidity buffers, adding cost but reinforcing product quality.
- **For investors, the dispersion between well- and poorly-built retail vehicles will widen.** Four markers separate institutional-grade retail product from repackaged versions: liquidity sized to the underlying asset, clean separation of retail and institutional pools, independent valuation governance and real manager-selection capability behind the wrapper. For institutional investors, the illiquidity premium will compress and deal-allocation conflicts will emerge as GPs spread a finite pipeline across vehicles with different liquidity profiles, making vigilance on cross-pool transfers essential.

A new investor base, a new product shelf

Private markets were built around institutions and family offices. The 1980s leveraged-buyout boom, fueled by the advent of junk bond markets, lit the fuse; institutional money turned the resulting industry into a colossus. Global private-capital assets under management (AUM) have grown from roughly USD800bn in the early 2000s to over USD17trn by 2024 (Figure 1), more than twentyfold in two decades.

Figure 1: Global private capital AUM has grown more than twentyfold since 2000




Sources: Preqin, Allianz Research

Note: Secondaries and funds of funds are excluded to avoid double-counting.

The pivotal driver behind the transformation of private markets into a mainstream institutional allocation is Yale's endowment model. When David Swensen took over the university's investment office in 1985, endowments ran the standard 60/40 mix. Over the next three decades he tilted the portfolio aggressively into private assets, arguing that long-duration capital should earn an illiquidity premium that liquid investors could not. By 2021, Yale held roughly 70% in alternatives, nearly ten times the pre-Swensen level, and the portfolio had returned 13.7% a year over Swensen's tenure against 10.5% for 60/40 and 10.3% for the average endowment (Figure 2). Yale's success rippled outward. By the 2010s, the biggest US university endowments held 40-60% in private markets on average. Sovereign wealth funds, public pension funds and insurers followed. Today, global institutional investors allocate around 10% to private assets (Figure 2), with endowments and public pensions allocating more than a third.

Figure 2: The Yale endowment has outperformed traditional benchmarks under Swensen's management (left) and global institutional investors are now allocating around 10% in private assets (right)




Sources: Bloomberg, The Economist, Preqin, Allianz Research

The fit is natural. Most institutional capital sits in drawdown funds run by general partners (GPs) who source, structure and manage the underlying deals. The 10-year-plus horizon suits pension funds and life insurers with liabilities measured in decades; in-house teams, investment committees and boards do the diligence that opaque assets demand and scale buys both fee leverage and access to the best managers. Beyond drawdown funds, institutions increasingly mix in co-investments (alongside the GP, often at low or zero fees), separately managed accounts (SMAs, typically USD100mn-plus, fully customized) and open-ended evergreens (operational simplicity and continuous liquidity, no J-curve). Together these tools let institutions calibrate exposure with a precision the conventional fund template cannot match. **Family offices, the private treasuries of the ultra-wealthy (a USD 30mn threshold, though in practice more is needed to access large GPs), follow the same model in miniature.** They vary in sophistication but use the same vehicles to reach the same managers, sometimes on slightly worse terms.

Now, the action has moved to the wealth and mass-market channels, where investors' needs are entirely different. "Democratization" means giving private-market access to those locked out of the institutional product suite and that requires more than lower minimums. It calls for fresh product design, with new trade-offs between liquidity, fees, transparency and asset quality (Figure 3).


Figure 3: Investor segmentation and product structures across the wealth pyramid



Source: Allianz Research

The wealth channel covers the high- and very-high-net-worth segment, growing fast in both numbers and assets. Entry starts at USD1mn of investable assets; the more active VHNW tier kicks in at USD5mn. US households with USD5mn to USD20mn now control 40% of addressable wealth, up from 18% in 2013, turning them from once-niche cohort to a market force. They are too small to deal with GPs directly and reach private markets through private banks, wealth managers and financial advisers. This is the most commercially mature corner of democratization, and one of the fastest-growing. Inflows to the major listed alternative-asset managers have surged over the past three years, topping USD25bn in the fourth quarter of 2025 alone (Figure 4). US financial advisers now allocate USD1.9trn to less-than-fully-liquid strategies, but the average adviser still has only 2.3% of client portfolios in alternatives, and roughly half use none at all. The runway is long.

Figure 4: Net inflows to wealth products offered by major listed alternative managers have surged since early 2023



Sources: JP Morgan, Allianz Research

The mass market is a much bigger prize, and a wholly different machine. Its investors have less investment expertise than institutions, shorter effective horizons and no practical capacity to select managers themselves, which is why the channel is built around fiduciaries, insurers and regulated structures that do the selection for them. The US mass affluent and middle-market segments, 46.9m households in all, hold USD25trn of financial assets. The retirement system holds another USD49trn, of which USD14trn sits in employer-sponsored DC plans. In Europe, household savings are mostly parked in low-yielding bank deposits, around EUR10trn across the EU; the next-largest pool is unit-linked insurance, notably France's assurance-vie, Italy's PIR and Germany's equivalents.

Three structural differences set wealth-channel investors apart from institutions. They lack the governance infrastructure, they are less tolerant to illiquidity and they reach private markets through intermediaries whose operating models demand standardized features that fit into model portfolios. The result is a generation of products engineered for accessibility, with performance pursued within those constraints rather than at any cost. Semi-liquid evergreens have become the workhorse: quarterly subscription and redemption windows, no capital calls, simplified fees and standardized tax reporting. Each feature addresses a specific constraint. Well-engineered evergreens, with adequate liquidity buffers and disciplined valuation, can deliver a meaningful share of the institutional return profile to investors who previously had no access at all.

Reaching the mass market requires a different paradigm. The end investor no longer chooses the product directly; instead, the structure has to satisfy a fiduciary or regulator that it is suitable for an average participant. The result is a new range of vehicles, each tailored to a specific regulatory category or savings product rather than evaluated on its own merits. New structures are launching in every major jurisdiction; the main channels are defined contribution (DC) plans, ELTIFs and LTAFs, interval funds, public-private hybrids and insurance wrappers.

The deepest pools sit inside existing retirement and savings systems. US DC plans, accessed mainly through target-date funds with embedded private-market sleeves, are the largest mass-market pool in the world: USD14trn of employer-sponsored DC assets, including roughly USD10trn in 401(k)s. European insurance wrappers play the same role in the savings sector, with the insurer handling manager selection, valuation oversight and liquidity management for the policyholder. Both work precisely because they do not demand investment expertise of the end investor: the fiduciary or insurer takes on the discipline that retail investors lack and would struggle to acquire. The private exposure is bundled inside a familiar structure and the design ensures that no single policyholder's liquidity needs can destabilize the underlying assets. The retail experience is unchanged at the surface; private-market access, properly governed, sits one layer below. For long-horizon savers, this is one of the few realistic routes to the diversification and return premium of private markets without bearing institutional-style governance demands themselves.

Regulated retail vehicles take a more direct path, swapping regulators for fiduciaries. ELTIFs in the EU, LTAFs in the UK and interval funds in the US are standalone retail products an investor or adviser can pick actively, subject to rules on eligible assets, concentration, mandatory repurchases and disclosure. The regulatory framework replaces the institutional governance these products would otherwise lack, which is a large part of their appeal: the regulator has effectively vetted the structure. The cost is product flexibility and the channel has tilted strongly toward private debt and lower-volatility strategies that fit cleanly inside the templates.

Hybrid public-private funds are the newest design, aimed squarely at the liquidity problem. DC plans absorb redemption pressure through fiduciary scale; regulated retail vehicles use prescribed gates; hybrids embed a 30-60% public-asset buffer into the product itself. The public sleeve provides exit capacity; the private sleeve provides the alternative exposure that justifies the fund. The trade-off is return: a 40% private allocation inside a fund that is itself only a slice of a retail portfolio leaves effective private exposure of just a few percent at the investor level. Hybrids are a design solution to mass-market liquidity rather than a route to maximized private-market exposure. Their merit is extending some private-market characteristics to investors for whom pure private vehicles are not workable, accepting a return trade-off in exchange for genuinely daily-liquidity behavior.

The four channels share a single design principle, with different machinery. None assumes the end investor actively chooses the product. The decision rests with a fiduciary (DC plans), an insurer (insurance wrappers), a regulatory template (ELTIFs, LTAFs, interval funds) or the structural design itself (hybrids). This shift in who decides ultimately drives the outcome for investors, and means the routes to mass-market capital are likely to fragment by jurisdiction rather than converge on a single global model.

Why private markets are going mainstream now

For retail investors, the case rests on diversification and better risk-adjusted returns. The evidence is plain: private assets diversify; traditional ones increasingly do not. Since 2022, correlations between public equities and bonds have risen sharply, eroding the protective property of the standard 60/40 portfolio. Modern portfolio theory says that adding assets with low or negative correlations shifts the efficient frontier up and to the left; the correlation matrix in Table 1 backs this up. The numbers are forward-looking, not historical, and reflect today's regime. The arrows mark the largest moves, concentrated among traditional assets that have grown more closely correlated under stagflationary pressure.


Table 1: Cross asset correlation matrix

	Euro Corporates	Euro High Yield	Euro Covered Bonds	USD Corporates	Equity BAA Aggregate	US Private Equity (Buyout)	US Private Debt (senior)	US Real Estate	US Infrastructure	Euro Government Bonds	Euro Agency bonds	Emerging Markets Bonds	Gold
Euro Corporates	1,00	0,82	0,86	0,77	0,63	0,20	0,31	0,24	0,17	0,80	0,83	0,87	0,47
Euro High Yield		1,00	0,53	0,54	0,82	0,24	0,33	0,20	0,30	0,42	0,48 ▲	0,83	0,43
Euro Covered Bonds			1,00	0,71	0,37	0,21	0,31	0,19	0,20	0,93	0,94	0,70	0,49
USD Corporates				1,00	0,43	0,17	0,23	0,30	0,12	0,72	0,73	0,76	0,53
Equity BAA Aggregate					1,00	0,14	0,26	0,09	0,24	0,28 ▲	0,33	0,74	0,46
US Private Equity (Buyout)						1,00	0,81	0,42 ▼	0,71	0,09	0,10	0,20	0,42
US Private Debt (senior)							1,00	0,43	0,62	0,22	0,25	0,31	0,37
US Real Estate								1,00	0,17 ▼	0,23	0,21	0,17	0,04
US Infrastructure									1,00	0,01	0,03	0,27	0,58 ▲
Euro Government Bonds										1,00	0,98	0,62	0,35
Euro Agency bonds											1,00	0,67	0,38
Emerging Markets Bonds												1,00	0,63
Gold													1,00

Source: LSEG Datastream, Allianz Research

Private assets are consistently uncorrelated with public ones, and barely correlated with each other. As public-market correlations climb, infrastructure stands out as the biggest beneficiary; pair correlations with euro government, agency and corporate bonds are all low. The more striking feature of the matrix is the heterogeneity within the private universe itself: real estate and infrastructure are not substitutes for private equity. The rolling four-year analysis in Figure 6 backs this up. From 2005 to 2026, all four American private asset classes, PE buyout, senior debt, real estate and infrastructure, have averaged correlations close to zero with a 60/40 public portfolio, none above 0.5. The dispersion has widened recently, with infrastructure the most decoupled, even as overall market correlations have climbed.


Figure 6: Cross-asset rolling correlations



Source: Pitchbook, Allianz Research

The most compelling quantitative argument for democratization sits in the efficient frontier itself. Figure 7 plots annualized volatility against total return for portfolios ranging from a pure traditional mix of 13 asset classes to ones with progressively larger alternative allocations (10-50%). As the alternatives weight rises, the frontier shifts unmistakably outward: the green alternatives frontier dominates the orange liquid frontier at every point, lower volatility and higher returns with no exceptions. Theory says low-correlation assets should improve the trade-off; the empirical gap is striking. Private assets earn an illiquidity premium without the concentration risk that comes from leaning on public equity, provided investors can stomach the illiquidity and keep enough liquid ballast for market stress. The classic 60/40 (orange dot) and the typical European household (green dot) sit on noticeably higher volatility, since bonds and equities have failed to diversify each other since 2022's stagflation. Owning a home does not fix this; a diversified slice of private markets does. The same pattern shows up in Conditional Value at Risk, which captures downside losses in bad markets and actually fits the non-normal distribution households better than volatility does.


Figure 7: Efficient frontier liquid - illiquid assets



Source: Pitchbook, Allianz Research

An optimizer sharpens the picture. Figures 8 and 9 turn the theoretical frontier into a realized-performance comparison. The optimized portfolio (blue), allocating across alternatives and traditional assets, posts a Sharpe ratio of 2.85, against 0.64 for the typical household benchmark (orange). The optimizer solves one trade-off: maximize risk-adjusted return while penalizing both single-asset concentration and large tail losses (CVaR). It uses market-cap weights, forward-looking assumptions for returns, volatility and correlations, and alternative-history scenarios rather than historical averages, producing an allocation that holds up across regimes. The result: a balanced portfolio, half in private markets, complemented by low-risk bonds, gold and cash. Smaller private allocations still make sense, since the diversification benefit persists across weightings. The right share for any investor depends on their tolerance for risk.

Figure 8: Portfolio allocation in % (left) and portfolio performance (right)




Source: Pitchbook, Allianz Research

For governments, democratization is more than an access question. It is a policy lever for retirement adequacy, distributional fairness and the funding of long-term economic priorities. With public finances constrained and traditional capital sources struggling to keep up, private markets are being repositioned as a way to mobilize household savings.

The most pressing motive is the retirement-adequacy crisis. The shift from defined-benefit to defined-contribution pensions has moved investment risk from employers and governments to individuals, leaving DC participants with portfolios dominated by public equities and bonds. DB plans have long held 15-30% in alternatives and benefited; DC participants have not, and the gap has compounded into a serious shortfall. The global retirement-savings gap is projected to reach USD400trn by 2050, up from USD70trn in 2023. With public-pension systems under strain, widening the DC investment universe is a way to narrow the gap without lifting public spending.

The second motive is inequality. As private markets stay walled off, a growing share of value creation slips out of reach of ordinary savers. The number of US-listed firms has roughly halved since the late 1990s, while the number with fewer than 500 employees has grown by nearly a fifth (Figure 10). Companies stay private for longer, so much of their growth, and the wealth that comes with it, is captured by sophisticated investors with broader access. Ordinary savers, stuck in public equities, watch from the sidelines.


Figure 9: US listed companies have nearly halved while mid-market firms have grown by one-fifth over the past three decades



Sources: US Census Bureau, World Bank, Allianz Research

Public-market concentration is a third motive. The top 10 S&P 500 firms now make up 40% of the index, and US equity benchmarks are more concentrated than at any point since 2000 (Figure 11). Retail investors who hold passive index funds are heavily exposed to a narrow set of large-cap tech names; private assets reach sectors and stages that public indices underrepresent. The case for retail access is not just about returns. It is also about reducing the concentration risk that the passive revolution has, unintentionally, amplified.

Figure 10: S&P 500 concentration has reached its highest level since 2000




Sources: Macrobond, Allianz Research

Private capital is also being marshaled to fill gaps public budgets cannot. The EU's investment needs for infrastructure, the energy transition and digital transformation are projected at roughly 5% of GDP a year, well beyond what fiscally constrained governments can fund. Europe's EUR10trn or so of household bank deposits is largely unproductive capital. Channeling those savings into infrastructure, growth equity and energy projects through ELTIFs or LTAFs would do two things at once: lift household returns over the long term and deliver real-economy capital that the public purse cannot.

Jurisdictional competition is shaping both the design and the timing of these reforms. US, the EU, the UK, Australia and Singapore are all moving at once, none willing to see domestic savings drift offshore in search of alternatives. Britain's decision to make LTAFs ISA-eligible from 2026 was driven by competition with Luxembourg and Dublin as fund-domicile hubs; France's Green Industry Act was designed to keep retail savings inside the French insurance sector. The race is not just to channel retail capital into private markets, but to do so on one's own terms before a rival captures the flow.

For asset managers, the move into wealth and retail is no longer optional. After decades of rising inflows, institutional fundraising has hit a wall. In 2025, closed-end private-equity fundraising fell to a five-year low (Figure 12), as subdued exits left more than 16,000 portfolio companies held for over four years and LP distributions slowed to a trickle. The 2022 public-market drawdowns produced a denominator effect that left many institutions over their allocation targets, prompting a pause in fresh commitments. Beyond the cycle, the institutions that began private-markets programs in the 2000s have now reached, or nearly reached, their targets. With the traditional capital pool maturing, the industry has pivoted to investor groups that, until recently, had limited access.

Figure 11: PE fundraising volume and fund count have fallen sharply from their 2021 peaks



Sources: Preqin, Allianz Research

Box 1. Policy in motion: bringing private markets into 401(k)s


For most of their existence, American 401(k) and other defined-contribution (DC) plans have excluded private-market assets. ERISA fiduciary constraints, litigation risk and operational frictions (illiquidity, valuation complexity, fee opacity) made them poor fits. The SEC's accredited-investor rules further limited direct retail participation, reinforcing the structural divide between DC and defined-benefit (DB) plans. The first significant move came in June 2020, when the DOL issued an Information Letter allowing private equity inside diversified DC options such as target-date and balanced funds, provided it was managed prudently. A December 2021 Supplemental Statement then signaled a pause, noting that most plan sponsors lacked the expertise to evaluate such products properly.

By 2025 the policy had shifted. On 12 August 2025, the DOL formally rescinded its 2021 guidance discouraging 401(k) alternatives. The move aligns with the White House Executive Order on "Democratizing Access to Alternative Assets for 401(k) Investors" and clarifies that exposure to private-market assets may be permitted within diversified options under ERISA. The DOL is now signaling a neutral, principles-based approach: Fiduciaries must judge investment options on their merits and suitability, a clear departure from the earlier practice of singling out specific asset classes for scrutiny.

Congress has lent legislative support. The SECURE Act 2.0 of 2022 facilitated alternative access indirectly, by letting 403(b) plans use Collective Investment Trusts (CITs). CITs are often the preferred vehicle for private-market exposure in employer plans thanks to lower costs and more customization than mutual funds. Together these moves create a coherent policy environment in which private-market exposure can expand within DC plans under existing fiduciary and securities-law frameworks. An alternative investment option that sits inside a well-diversified fund managed by qualified professionals, and that the plan sponsor has documented thoroughly, can sit comfortably within ERISA's duty of prudence. Further sub-regulatory guidance from the DOL is expected, and the SEC may add guidance on valuation and disclosure as more products come to market.

The base effects are substantial. As of Q1 2025, the American DC system held roughly USD12.2trn in assets, compared with USD3.8trn in private and public DB plans. Alternatives make up around 30% of DB allocations but just 0.1% of DC portfolios, so even a modest reallocation would be material in absolute terms. A baseline scenario of 3-5% DC penetration by 2030 would deliver cumulative inflows of USD300bn-400bn; an upside of 10% penetration by 2035 would mean USD1.2trn-1.5trn in new allocations. Adoption is expected to start with large-plan default options that wrap diversified PE and private credit inside multi-asset structures, then diffuse gradually into mid-sized plans as operational frameworks standardize.


Figure 12: US retirement market assets (USD tn)



Sources: Investment Company Institute, The US Retirement Market, Fourth Quarter 2024, Allianz Research

Product evolution and opportunity set. Capital will flow along the path of least operational resistance. CITs are gaining ground because of their flexibility on liquidity and fees, smooth integration into plan recordkeeping and exemption from certain retail-disclosure requirements under the Investment Company Act. Semi-liquid '40-Act interval and tender-offer funds serve as secondary conduits for smaller, retail-aligned platforms, though their quarterly-liquidity features and higher cost structures may limit early scale. Across both channels, managers are re-engineering fund economics toward evergreen structures with quarterly or semi-annual rebalancing to accommodate DC flows. Pitchbook estimates that roughly \$2.7 tn is now managed in various indefinite-life formats globally, projecting growth to \$4.4 tn by the end of 2029 (Figure 13). Recordkeepers, custodians and administrators are investing in capabilities to support partial redemptions, performance-fee accruals and standardized daily valuation reporting.

Figure 13: Pitchbook estimate for Evergreen fund AUM forecast (USD tn)



Source(s): Pitchbook, Allianz Research. Geography: Global. Note: Note: Forecasts were generated on April 14, 2025. "Institutional evergreen" includes insurance AUM from Blackstone, KKR, Blue Owl Capital, The Carlyle Group, Ares Management, Apollo Global Management, and Brookfield.

Within diversified default options such as target-date funds, private equity is expected to dominate the growth allocation for younger cohorts, mainly through fund-of-funds or secondaries to achieve vintage diversification and reduce dispersion risk. Private credit is well placed to stabilize income for pre-retirement cohorts, focused on senior secured lending, infrastructure debt and core real-estate credit. Real assets (core real estate, infrastructure equity) will play a secondary role as inflation hedges and diversifiers, supported by established appraisal-based

valuation practices. Allocations will stay modest relative to liquid assets, typically capped at 10-15% of the default fund, to maintain daily-liquidity compliance.

The primary concerns at this stage are fiduciary risk, liquidity and valuation governance. Fiduciary exposure under ERISA remains real: sponsors must justify private-market inclusion as prudent on a net-of-fee basis, with documented due diligence and monitoring. The mismatch between daily participant flows and illiquid underlying holdings calls for hard caps, liquid sleeves and robust queue procedures, and under stress redemption gates can still be tested. A related vulnerability is the need for valuation fairness across transacting participants, which requires consistent third-party appraisal cycles, interim model-based pricing and audit transparency. Operational dependencies, recordkeeping, performance-fee calculation and participant communications introduce a layer of execution risk that must be addressed before a large-scale rollout.

For investors, the policy shift creates a durable, multi-year reallocation theme, not a short-term trade. Early strategies will favor diversified private equity and senior private credit that fit established liquidity and valuation criteria, with real assets in a supporting role. As operational systems and legal precedents mature, DC allocations to private markets could move from pilot programs to a structural component of retirement portfolios. Early participation should bring stable, long-duration capital to managers; investors should expect risk premia in yield-oriented segments to converge gradually as the DC channel scales.

Do returns justify the hype?

Whether private assets truly outperform is less straightforward than the headline numbers suggest. At first glance, the case is compelling. Over the past two decades, buyout funds have delivered net annualized returns of roughly 12-14%, against 7-9% for broad public benchmarks like the Russell 3000¹. But this obscures significant variation. Private debt has yielded steady mid- to high-single-digit returns; venture capital and real estate have been more cyclical, with periods of underperformance against public benchmarks.

The critical adjustment is leverage, and it transforms the comparison. Buyout funds typically run 50-70% loan-to-value at the initial stage to amplify equity returns; public benchmarks are unlevered. Once public markets are adjusted to comparable risk, the excess return collapses. Simulations from 2000-2024 show that a 20% buyout allocation generates only about 64bps of annual excess return over an unlevered benchmark, and roughly nothing once the benchmark is levered to match. Most of the apparent alpha turns out to be financial engineering, not investment skill².

Table 2: Private markets vs public markets (leveraged comparison, 2000–2024)

Asset Class	Portfolio Return	Benchmark (Unlevered)	Excess Return	Benchmark (Leveraged)	Excess Return (Leveraged)
Private Equity (Buyout)	7.54%	6.90%	0.64%	7.58%	-0.03%
Private Debt	7.37%	6.80%	0.57%	7.09%	0.29%
Venture Capital	5.71%	6.92%	-1.21%	,	Negative
Real Estate	6.89%	7.18%	-0.29%	,	Negative

Source: PitchBook simulations

The picture is one of conditional opportunities, not structural outperformance. Private markets can generate excess returns, but those returns are neither uniform nor guaranteed. Outcomes diverge sharply by asset class. Private debt emerges as the most consistent contributor, holding on to positive excess returns even on a leverage-adjusted basis. The buyout edge largely disappears; venture capital and real estate typically underperform. Results hinge on cycle, entry point and vintage, and the post-crisis years show notable regime sensitivity. The notion of a stable, persistent illiquidity premium is overstated; the return advantage is more fragile and context-dependent than the brochures suggest.

¹ Preqin, Burgiss, Cambridge Associates long-term private equity performance vs public benchmarks.

² PitchBook (2025), Are Private Markets Worth It?

Manager selection is the whole game. Private markets are defined less by an illiquidity premium than by a dispersion premium. Even modest selection skill, a moderately higher probability of landing in top-quartile funds, can lift total-portfolio return by 20-60bps a year, which is roughly 100-300bps of alpha inside the private allocation itself. The effect is largest in buyout and venture capital, where dispersion is widest, and smallest in private debt, where outcomes cluster. The risks of getting it wrong fall on direct and indirect investors alike, particularly those allocating at scale. Private markets reward access, discipline and underwriting capability far more than they reward passive exposure.

Table 3: Impact of manager selection skill (Average annualized alpha by strategy)

Asset Class	Underperformer Avoidance	Moderate Skill	High Skill
Private Equity (Buyout)	0.41%	0.35%	0.64%
Venture Capital	0.38%	0.32%	0.61%
Private Debt	0.25%	0.21%	0.38%
Real Estate	0.32%	0.28%	0.47%

Source: PitchBook simulations

Do retail investors get the B-product?

Three structural factors place retail vehicles on weaker terms than institutional ones. Higher fees eat into the return premium; retail investors risk being handed inferior underlying assets and in an asset class where the gap between best and worst is unusually wide, retail investors cannot meaningfully pick managers themselves. The three disadvantages compound, dragging down net returns.

Retail access comes with stacked fees that eat into the return premium. Institutional drawdown funds typically charge 1.5-2% on committed capital plus 20% carry over an 8% hurdle, and the largest LPs negotiate down within those bands, with within-fund management-fee dispersion of about 90bps. Wealth and retail vehicles carry additional cost layers, reflecting the work of getting private-market product to a broader audience: placement fees at the distribution platform, advisory fees from the adviser and a feeder-fund management fee on top of the underlying fund's economics (Table 4). Add it up and the gap between retail and institutional access often exceeds 200bps a year, enough, over a decade, to consume a meaningful share of the very illiquidity premium that justified the asset class to institutions in the first place.

Table 4: Fee layers in wealth channels

Fee type	Recipient	Purpose	Typical range	Notes
Servicing fees	Adviser (sometimes shared with brokerage)	Annual compensation for account management and client servicing	0.25% - 0.85% per year of client investment	Brokerage houses sometimes take a cut
Placement fees	Brokerage	One-time fee for offering the fund on the brokerage platform	Up to 0.5% of initial investment	Some funds cap combined servicing + placement at 0.85%
Brokerage commissions	Brokerage	Front-end commission charged at the point of sale	~2% on average (up to 3.5%)	Advisers have discretion to waive
Independent adviser fees	Independent adviser	Flat advisory fee on overall client assets	Variable (typically % of AUM)	Generally do not charge commissions or significant placement fees

Sources: FT, Allianz Research


Notes: These fees sit on top of the underlying private capital fund's own management fee (typically 1.5%-2% of AUM) and performance fee (typically 20% of profits above an 8% hurdle)

Retail and institutional vehicles also tend to diverge in what they hold. Retail evergreens are naturally suited to perpetual-hold, lower-volatility assets that support NAV-based redemptions, while buy-to-exit strategies that need patient capital sit more comfortably in institutional drawdown funds. This is a structural feature of the architecture. The liquidity profile of the wrapper shapes what can sit inside it. The implication is that two vehicles carrying the same brand can deliver materially different return and risk profiles, and advisers and platforms will need to make that distinction legible. The rise of GP-led secondaries has created a channel for moving hard-to-exit assets between

funds. In these transactions, the GP is seller, buyer and valuation agent at once, an arrangement that puts the integrity of the price under strain.

The third disadvantage is the inability to choose. Manager selection is arguably the most important driver of private-market returns; the gap between top- and bottom-quartile private-equity managers is roughly 1,400bps, against around 300bps in public equity. For institutions, that dispersion is an opportunity. For retail investors, it is a risk. Mass-market participants in DC plans, insurance wrappers and platform-default products do not pick managers at all; intermediaries do, on the basis of their own incentives, including fee-sharing and distribution relationships, which can favor available managers over the most suitable ones (Figure 13).

Figure 14: Performance dispersion is especially wide for buyout funds



Sources: KKR, eVestment Alliance, Preqin, Allianz Research

Notes: Bars show median net IRRs; whiskers show first- and third-quartile thresholds. 15-year period through 31 December 2024. US equities include large- and small-cap indexes


The three disadvantages mentioned above are interconnected and have a compounding effect. In order to achieve the same returns as institutions, retail investors must be willing to pay higher fees. Inferior underlying assets can negatively impact the gross return received. The inability to select top-quartile managers anchors the median retail outcome closer to the bottom of the dispersion than the top. When the effects compound, the median retail experience in an average private-market product can deliver a meaningfully lower net-of-fee return than an institutional investor in the same asset class, even though both are nominally invested in private markets. The gap is a function of product design, and well-built retail vehicles can close a portion of it.

The democratization narrative, therefore, comes with an important qualifier. Retail investors are gaining genuine access to assets that were previously out of reach, but the conditions of that access can compress the return premium that justifies the asset class in the first place. Whether the move from public to private markets pays off on a risk-adjusted basis depends on several factors: the specific product, the quality of the manager, the fee structure and the ability of the investor or adviser to distinguish well-engineered retail vehicles from those that travel less well outside the institutional setting they were designed for.

Is the “semi-liquid” promise an illusion?

An industry built on locking capital up is now selling products that promise to let it out. Retail investors face less predictable liquidity needs than institutions, and semi-liquid products try to bridge the gap with quarterly redemption windows. The bridge is narrow. Withdrawals are permitted only on set dates and after a holding period. If aggregate requests breach a preset threshold, typically 5% of NAV, the manager can invoke gates and honor only part of the request, leaving investors locked in for longer than they expected (Figure 15).

Figure 15: Redemption pressure is building for non-listed BDCs



Sources: SEC filings, Allianz Research

Notes: the data represent the mean and median redemption request rates for the eight largest non-listed BDCs

Disclosure has not always kept up with product complexity. Semi-liquid evergreens are typically marketed as a route into previously exclusive strategies, with the added convenience of periodic redemptions. Gates and conditional liquidity terms appear in the fund documentation, but get less prominence in marketing materials than the headline features. The gap between how the product behaves in calm markets and how it behaves under stress therefore surfaces only when redemption pressure rises, often to investors who had not fully absorbed the small print.

The underlying problem is that demand for liquidity peaks precisely when supply is lowest. Semi-liquid mechanisms work smoothly in normal markets. In periods of market stress, liquidity tightness or broad portfolio rebalancing, just when investors most want their money back, managers find it hardest to provide. Gates prevent the bank-run dynamic that would force fire sales and worse outcomes for everyone, but they trap investors at the worst possible time. Retail investors who expect access to their capital may discover that access depends on the same market conditions that drove them to seek it.

Once redemption pressure builds, the dynamics turn self-reinforcing. Initial requests draw headlines, which prompt more redemptions from anxious investors. Gating works as a brake, but the act of gating itself can be read as a distress signal, accelerating the next wave. If gates persist for several quarters, managers come under pressure to sell holdings or draw on subscription lines, producing losses for those who stayed and validating the concerns that drove the run. The non-traded BDC turmoil has played out exactly that way: redemptions rose, sponsors invoked caps, the caps generated more requests and gating spread across the category.


The semi-liquid promise is not an illusion; it is conditional. In normal markets, gates rarely trigger and redemptions clear smoothly. The mechanics work as intended. The case for democratization, however, depends on these structures holding up over time. The unfolding BDC episode is the first real-world stress test at scale, and a cautionary tale for the wider effort.

Box 2. The private BDC reckoning. A proof of concept for retail private credit or the proof of limits?

The Business Development Company (BDC) sector grew from USD190 bn in 2021 to over USD500bn by mid-2025, fueled by post-crisis bank retrenchment, floating-rate income and the push to bring private credit to retail and wealth-channel investors (Figure 16). The first quarter of 2026 provided the first real test of that expansion.

The result was not the broad credit crisis some had feared. It was something more telling. Wrappers designed for slow-moving institutional capital buckled when retail investors hit their limits, lagged marks and ran into limited liquidity in vehicles whose reported portfolios looked sound. This is not evidence of a failure in private credit; it is a reminder that product design, valuation architecture and investor base now matter alongside underlying credit quality. Any democratization vehicle built on periodic liquidity, manager-determined valuations and illiquid assets will carry some version of the same risk.

Figure 16. Perpetual BDCs have captured impressive AUM growth



Source Cliffwater Direct Lending Index, Preqin. As of Q2 2025

Perpetual non-traded BDCs pose a persistent design challenge. In Q1 2026, Blue Owl's technology vehicle received redemption requests totaling 40.7% of shares against a 5% quarterly cap, even though reported non-accruals were only 0.6%. The gap matters because quarterly liquidity sits on top of manager-marked, illiquid loans, giving investors an incentive to redeem before future marks absorb losses. The implication is durable: even if this queue clears, the wrapper stays vulnerable whenever confidence in reported NAV weakens.

Investors should know that reported income and reported credit quality can diverge for longer than commonly assumed. In Q4 2025, the share of private-credit loans classified as bad PIK reached 6.4%, against headline non-accrual rates of around 2%; Lincoln estimated a shadow default rate near 6%. The mechanism is straightforward: non-cash income preserves earnings optics while borrower stress is deferred, not resolved. As a result, dividend coverage, NAV stability and the perceived quality of the management team can look stronger than they really are until the recognition event arrives.

The issue is not just credit losses; it is delayed price discovery. BlackRock TCP Capital's roughly 19% NAV cut in late 2025 showed that loans previously carried near par could be repriced sharply once underlying weakness was acknowledged. This creates a feedback loop: stale marks dampen observed volatility but raise the incentive to exit before the revaluation lands. Future stress episodes are therefore likely to show up first as redemption pressure and only later as formal credit deterioration.


The 2026 episode points to structural vulnerability, not merely idiosyncratic manager error. According to Fitch, non-traded BDCs entered the period with lower leverage than their traded peers, yet sector-wide redemption requests for the largest vehicles averaged 12.1% in Q1 2026. The pattern suggests the pressure came from distribution and product design as much as from weak balance sheets. The implication for democratization is clear: expanding retail private credit will require a redesign of the framework, not just sharper marketing or stricter manager selection.

A new vector for systemic risk?

Even before retail money enters, private markets already inject several systemic risks into the financial system.

Start with banks. By end-2024, the non-bank financial intermediation (NBFI) sector held 51% of global financial assets (Figure 17), and the linkages have tightened: banks' credit exposure to private credit funds alone is estimated at USD270bn-500bn. In times of stress, private funds tend to draw on their committed credit lines just as banks are absorbing market volatility and rising loan losses. The drawdowns force previously off-balance-sheet exposures onto bank balance sheets, lifting capital requirements at the worst possible moment.

Figure 17: The NBFI sector has grown to 51% of global financial assets, tightening linkages with the banking system




Sources: FSB, Allianz Research

Within the life sector, regulators have flagged a more specific concern: the rise of private-equity-backed insurers.

Private capital now controls roughly USD700bn of US life-insurance assets, and the market share of PE-backed insurers has climbed from 1% in 2012 to 13% by 2023. They have driven a broader shift toward private and structured credit, which now makes up 36% of North American life insurers' balance sheets (Figure 18). The IMF and the IAIS note that this subset of the industry tends to run thinner liquidity buffers and larger illiquid-credit positions than traditional, prudentially regulated life insurers. Stress at a PE-backed insurer can therefore feed pro-cyclical asset sales, as the 2023 collapse of Italy's Eurovita showed. The episode is a reminder that the case for retail private-market exposure rests on the quality of the wrapper as much as the quality of the asset, and that not all life-insurance balance sheets are built the same way.

Figure 18: Private credit accounts for 36% of North American life insurers' balance sheets



Sources: Moody's, Allianz Research

Democratization broadens these channels by tilting them toward a more reactive investor base. The institutional money that powered three decades of growth came with long horizons, contractual lock-ups and governance structures that resisted short-term liquidity pressure. Retail capital tends to chase peaks and redeem under stress, often on the same signals that drive broader volatility: news headlines, equity drawdowns, employment shocks. As retail's share of private-market AUM grows, each existing channel widens and the marginal investor becomes more skittish. The asset class may shift from absorbing market stress to amplifying it.

Democratization also introduces new risks. The mismatch between quarterly redemptions and decade-long assets revives run dynamics that the asset class had engineered away. Traditional institutional funds run on closed-end capital and multi-year commitments, which contain rapid redemption pressure: if a few funds suffer big losses, capital stays locked and sophisticated investors absorb the impact without a broader confidence shock. That mechanism weakens sharply once the investor base turns retail. Quarterly redemption rights on fundamentally illiquid assets reintroduce the maturity transformation that closed-end structures had removed. Gating can in principle prevent a bank-run dynamic, but the channels funds use to raise liquidity, credit lines and new loan issuance, are directly tied to banks and broader credit markets. If redemption pressure persists for several quarters, forced sales at distressed prices cannot be ruled out.


Valuation contagion is amplified when retail money is involved. Private assets are appraised by management, not by markets, and NAV marks can lag reality by several quarters. When a publicly traded vehicle holding similar private-credit loans trades at a persistent discount to its stated NAV, the market is signaling that it does not trust the valuation. That skepticism rarely stays local: it can prompt LPs to rethink commitments, lenders to tighten collateral terms and providers to pull subscription lines. Retail investors, who may not always tell well-marked vehicles from poorly-marked ones, tend to redeem indiscriminately when confidence breaks. Incentives in retail-focused funds can lean toward smoother, more optimistic NAV paths, since stable marks support both fee continuity and the credibility of the redemption mechanism. Where governance does not push back, this can build feedback loops of optimism that unwind sharply when tested.

Concentration among a handful of major managers can turn firm-level stress into a market-wide contagion. Six firms now control nearly half the retail-alternatives market, and the largest individual vehicles top USD100bn in assets. The GPs that have been most aggressive in raising retail capital also run multiple semi-liquid wrappers across product categories. Stress in one large manager's product can therefore spread quickly through correlated vehicles run by the same firm, and through the wider retail-alternatives universe, where so much AUM sits with so few entities.

What happens next?

Democratization has hit its first substantial obstacle. Inflows to wealth products at the major listed alternative managers have fallen sharply from their December 2025 peak (Figure 19). BDC sales were down roughly 40% year-on-year in early 2026, and redemption caps were triggered across major non-traded vehicles. The narrative has gone from "growth drivers" to "discipline and selectivity" in a single year.

Figure 19: Inflows to wealth products have retreated from their December 2025 peak



Sources: JP Morgan, Allianz Research

This looks like a stress test, not a fundamental shift. The BDC episode exposed real vulnerabilities in semi-liquid product design, and the fundraising slowdown reflects multi-year exit constraints and investor over-allocation. The policy scaffolding put in place over the past three years – the DOL safe-harbor rule, ELTIF 2.0 and LTAF ISA eligibility – is built and unlikely to be dismantled as the pressures that drove it – retirement-adequacy gaps, infrastructure funding needs and the hunt for broader investment opportunity – remain unresolved.

The next phase will be recalibration. Product designs are evolving in three visible directions. First, larger liquid-asset buffers inside semi-liquid wrappers, with some new launches reported to hold 25-40% in public securities to support redemptions. Second, tighter gating, longer notice periods, smaller per-quarter caps and clearer disclosure of restrictions at the point of sale. Third, a migration toward asset classes better suited to retail liquidity profiles.

The most likely structural outcome is consolidation. Mega-GPs are well placed to ride the coming wave of retail capital; smaller managers without distribution scale, brand recognition or the operational capacity to run semi-liquid wrappers will struggle to compete for retail flows. The projected parity of retail and institutional fundraising by 2030 is therefore likely to coincide with a sharp concentration of AUM among a few large managers, deepening the concentration risk noted above.

The regulatory environment has been benign so far, but the BDC turbulence will probably trigger a new round of scrutiny. Adding liquidity has resurfaced concerns that closed-end structures had previously sidestepped. Lock-ins removed deposit-like risk by stopping investors from pulling capital on demand, which limited the channels through which fund stress could spread. Open-ended structures bring back the duration mismatch and tilt the investor base toward less sophisticated participants more likely to react to market volatility. As private-markets funds push into territory once held by bank lending, regulators face a familiar dilemma: the same maturity transformation that justified the post-2008 reforms is being recreated outside the bank-regulatory perimeter, and now sold to retail investors.

Regulators themselves are divided. The SEC has explicitly rejected the idea that private credit is a systemic threat, prioritizing capital formation. The IMF, the Bank of England and the IAIS take a different view, flagging

interconnections, valuation opacity and less sophisticated investors. The disagreement is essentially about how much weight to place on tail risks that have not yet materialized.

Four developments will shape the next 12-24 months. The DOL's safe-harbor rule for 401(k) alternatives, in public comment through mid-2026, will decide whether plan sponsors get the legal certainty to offer alternatives broadly in DC plans. *Anderson v. Intel*, before the US Supreme Court in the autumn 2026 term, will decide whether ERISA fiduciary-breach claims require "meaningful benchmarks", with direct implications for plan-sponsor litigation risk. ELTIF 2.0 will get its first proper stress test as semi-liquid launches multiply and distribution widens across the EU. And the BDC gating events of early 2026 will set the tone for retail investor protection for years, depending on whether the SEC or FINRA pursues formal enforcement on point-of-sale disclosure.


Whatever its form, the regulatory response will add cost and complexity. The BDC episode has already pushed the SEC to tighten oversight of how BDCs value illiquid holdings in falling markets. New disclosure rules, leverage limits, restrictions on retail-targeted marketing and mandatory liquidity buffers are all on the table in the major jurisdictions. Each addresses a real concern; each adds operational burden on GPs, costs that will be passed on to LPs and will erode the net-of-fee returns that justified the asset class in the first place. The deeper worry is that regulation will arrive after retail losses, eroding four decades of structural advantage.

What does this mean for investors?

Not all retail private-market vehicles are built the same way, and the distance between well-engineered and poorly-engineered product is likely to widen as the channel scales. Four markers separate the two. The first is liquidity architecture: adequate public-asset buffers, conservative gating thresholds and redemption windows calibrated to the underlying asset, rather than to the distribution channel. The second is clean separation between retail and institutional pools, so that allocation conflicts and cross-subsidies do not arise. The third is valuation governance: independent pricing, regular external review and transparency on NAV methodology. The fourth is manager selection capability, both at the GP level and, where the wrapper allows it, at the underlying-asset level. Investors and advisers will increasingly need to look past the brand and the headline strategy to these design features. For the industry, the opportunity is to build retail product that is genuinely institutional in its discipline, rather than institutional only in its marketing.

If retail capital arrives at scale, the illiquidity premium is likely to compress. Broader participation will erode the conditions that historically underpinned private-market outperformance, with more capital chasing an opportunity set that is already increasingly competitive. The pressure is acute now: managers are sitting on more than USD4trn of dry powder (Figure 20) to deploy in an environment of constrained exits and intense competition for quality deals. Institutions have historically captured the illiquidity premium through patient capital. That advantage is diluted as retail capital enters at scale.

Figure 20: Managers are sitting on USD4trn of dry powder



Sources: Preqin, Allianz Research

More fundamentally, the rise of retail products creates a structural conflict of interest in deal allocation. GPs must now spread a finite pipeline of deals across products with very different investor bases and liquidity profiles. That complicates life for institutional LPs, whose visibility on whether their interests are being protected, whether assets are being transferred between vehicles on unfavorable terms, and whether their original strategies are being reshaped to accommodate retail liquidity needs.

If GPs prioritize institutional vehicles, institutional returns are not directly compromised, but new risks emerge. Retail products receive systematically secondary deal flow while institutional LPs continue to see the most attractive transactions. A more troubling possibility is that retail vehicles end up serving as liquidity outlets for the institutional side, especially in tough exit environments, with hard-to-sell assets channeled into retail evergreens. In these trades the GP is seller, buyer and valuation agent at once, an arrangement whose integrity depends on a valuation discipline that is hard to sustain when the same firm is on both sides.

If GPs allocate deals on broadly equivalent terms across institutional and retail products, the critical question is whether retail and institutional capital is cleanly separated. In pooled vehicles, retail capital flows directly into vehicles with the same exposure as institutional LPs, which exposes institutional capital to the risk of more reactive retail behavior. In deployment, retail inflows tend to peak when valuations are high and competition for deals is fiercest, pushing GPs to deploy into transactions that would otherwise pace through. In the holding period, retail redemption pressure during market stress can force premature sales at distressed prices, with losses that institutional capital would have absorbed and recovered through patience. In distribution, the same pressure can disrupt the orderly exit timing institutional LPs rely on, with assets sold into weaker markets to meet retail liquidity needs rather than held until the optimal exit window. The institutional return premium earned from patient, locked-up capital is diluted in this scenario.

The corollary is that the firms best positioned in the next phase will be those that combine institutional underwriting with retail-aware product design: large balance sheets that can carry liquidity buffers without diluting returns, integrated insurance and asset-management capabilities that internalize valuation and governance discipline and distribution that reaches savers without adding unnecessary cost layers. For long-horizon retail investors, the structures that meet these tests are not a watered-down version of institutional product. They are a route to diversification and return outcomes that public markets, increasingly concentrated and increasingly correlated, can no longer deliver on their own.

These assessments are, as always, subject to the disclaimer provided below.

FORWARD-LOOKING STATEMENTS

The statements contained herein may include prospects, statements of future expectations and other forward-looking statements that are based on management's current views and assumptions and involve known and unknown risks and uncertainties. Actual results, performance or events may differ materially from those expressed

or implied in such forward-looking statements.

Such deviations may arise due to, without limitation, (i) changes of the general economic conditions and competitive situation, particularly in the Allianz Group's core business and core markets, (ii) performance of financial markets (particularly market volatility, liquidity and credit events), (iii) frequency and severity of insured loss events, including from natural catastrophes, and the development of loss expenses, (iv) mortality and morbidity levels and trends,

(v) persistency levels, (vi) particularly in the banking business, the extent of credit defaults, (vii) interest rate levels, (viii) currency exchange rates including the EUR/USD exchange rate, (ix) changes in laws and regulations, including tax regulations, (x) the impact of acquisitions, including related integration issues, and reorganization measures,

and (xi) general competitive factors, in each case on a local, regional, national and/or global basis. Many of these factors may be more likely to occur, or more pronounced, as a result of terrorist activities and their consequences.

NO DUTY TO UPDATE

The company assumes no obligation to update any information or forward-looking statement contained herein, save for any information required to be disclosed by law.

Allianz Trade is the trademark used to designate a range of services provided by Euler Hermes.

Thinking fast, building slow: The energy cost of the US AI boom

12 May 2026

Content

Page 3-4

Executive Summary

Page 5-10

Can power markets handle the AI surge?

Page 11-12

Beyond the baseline: Agentic AI, adoption velocity and the rebound effect

Page 13-17

The hidden bill?: AI's impact on power prices

Page 18-19

Preparing AI for energy and energy for AI: policy recommendations

Executive Summary



Patrick Hoffmann
Economist, ESG & AI
patrick.hoffmann@allianz.com



Katharina Utermoehl
Head of Thematic and Policy Research
katharina.uteramoehl@allianz.com

- **Artificial intelligence is about to impose the largest sustained demand shock on US electricity infrastructure in decades.** By 2030, data-center power consumption is expected to nearly double, lifting the sector's share of total US electricity demand from roughly 5% to around 9%. Although planned generation additions look sufficient on paper, data centers may absorb nearly half of projected new capacity, leaving thin margins if electric-vehicle adoption or industrial electrification accelerate faster than expected. Yet demand itself is evolving faster than forecasts can capture: generative AI reached 53% population-level adoption in just three years, agentic systems consume far more energy per interaction than conventional workloads and a 280-fold decline in inference costs since 2022 is driving a rebound effect that efficiency gains alone cannot offset.


- **The critical bottleneck, however, is not generation capacity but the grid itself.** Data centers can be built in under two years, but grid connections can take up to seven years in congested markets such as Northern Virginia. Nationwide interconnection requests now total 1.84 terawatts, exceeding total installed US generating capacity. Supply-chain shortages have pushed lead times for critical grid equipment to several years, while projected demand would require building roughly 8,000 km of high-voltage transmission lines annually – around ten times the current pace. The strain is already showing: In early 2026, the Department of Energy invoked emergency powers to shift data centers onto backup generation during peak demand periods. Rising public opposition and legislative scrutiny add a further layer of uncertainty that conventional supply forecasts have yet to fully capture. The strain is already visible in project pipelines, with half of the 12 GW of US data-center capacity planned for 2026 being delayed or cancelled.

- **Aggregate electricity prices have not yet fully reflected these pressures, but a growing „data-center premium“ is emerging.** States hosting the highest concentration of data-center activity have so far seen price inflation below the national average, reflecting favorable grid conditions, economies of scale and the lagged structure of utility rate-setting. Beneath the surface, however, the impact has become increasingly visible since 2023. US residential customers are already paying USD1.4bn more per year on their electricity bills as a direct result of data-center demand, with just five utilities serving 4.4mn households in Northern Virginia, the Pacific Northwest and Arizona accounting for more than 40% of that total. The sales-weighted average price effect across all utilities sits at just 0.6%, but for the most exposed utilities roughly 7.8pps of a 24.5% cumulative price increase between 2020 and 2024 are directly attributable to data-center demand, adding 0.19pp to headline inflation over four years through the direct electricity channel alone. These markets historically benefited from below-average electricity costs, a

gap that has already narrowed from 5% to 3.7% since 2020 and is set to close further. Data-center investment grew 32% in 2025 and is set to rise a further 75% in 2026 alone, pointing to an additional electricity price effect of close to 14pps for the most exposed utilities over 2025–2026, nearly doubling the cumulative four-year effect in just two years. Meanwhile investor-owned utilities filed USD18bn in rate-increase requests in 2025, the highest since the mid-1980s, with costs largely falling on existing rate-payers rather than the facilities driving them.

- **Preparing energy infrastructure for AI and addressing community concerns is as urgent as building the AI infrastructure itself.** The immediate priority is interconnection reform: binding timelines, penalties for speculative queue filings and priority treatment for shovel-ready projects with firm power commitments would help relieve the most acute bottlenecks. Cost allocation is equally important. Unless data centers bear a proportionate share of the infrastructure costs they create, public opposition and permitting delays will intensify further. Mandatory energy-use disclosure, incentives to redirect investment away from saturated regions, stronger efficiency standards, demand-flexibility mechanisms and stricter additionality requirements for power-purchase agreements would fill the most critical gaps in a policy framework that remains largely inadequate to the scale of the challenge – and lay the foundation for AI ambitions that the grid can actually support.





Can power markets handle the AI surge?


The AI revolution is reshaping global economies, starting with power markets. In the IEA's baseline scenario, AI-fueled growth in data-center power demand is expected to increase between 58% and 137% across major world regions between 2025 and 2030, supported by annual investment that now rivals the scale of the entire global oil and gas industry (Figure 1).¹ The shift is particularly pronounced in the US, where data-center power demand is projected to reach 426 TWh by 2030, nearly double the 2025 level. Based on long-term EIA projections, our base case estimates suggest this figure could double again by 2050, reaching roughly 810TWh. While the precise share attributable

to AI remains difficult to isolate, bottom-up modeling suggests AI inference already accounts for 12–14% of total consumption, with total AI data-center capacity reaching 29.6 GW by end-2025.² The implications for power systems are already visible in the US, where the data center share of total electricity consumption could jump from around 5% today to between 7% and 12% by 2028, with our estimate, based on current supply pipelines and baseline demand assumptions, pointing to a share of around 9%.³

¹ [Key Questions on Energy and AI \(IEA\)](#)

² [The 2026 AI Index Report | Stanford HAI](#)

³ [2024 United States Data Center Energy Usage Report | Energy Technologies Area](#)

Figure 1: Data-center power demand growth by region (in TWh)

Source: IEA baseline projection

This significant demand growth is fueling a massive expansion in generation capacity. According to the Energy Information Administration's (EIA) Short-Term Energy Outlook, power generation in the US is expected to grow by 441 TWh between 2023 and 2027, an increase of 10.5%.⁴ This would mean the five years since the breakthrough of generative AI see a larger increase than the preceding two decades, over which power generation grew by only 7.7%. 76% of the increase in power generation is expected to be driven by wind and solar (21.9% wind and 54.1% solar) with the remaining 23% supplied by gas (15.3%), nuclear (5.2%) and hydropower (3.0%). Coal remains on its post-2007 declining trend, though the pace has slowed markedly from -4% annual reductions before the AI era to just -0.5% today.

On paper, the capacity additions currently included in the EIA's pipeline of planned and under-construction projects appear broadly sufficient to accommodate rising data-center demand, with data centers expected to absorb around 47.1% of projected new supply.⁵ That apparent adequacy, however, rests on the fragile assumption that competing sources of electricity demand remain relatively contained. In reality, the grid is coming under pressure from

multiple directions simultaneously. Beyond data centers, electrified transport alone could materially tighten the balance. Under the IEA's baseline scenario, US electricity consumption from electric vehicles is projected to rise from 25.2TWh in 2024 to 91.4TWh by 2030. Near-term signals are mixed, but the risks remain tilted to the upside: Secondary-market EV sales continue to rise despite the expiry of the federal tax credit, while gasoline prices elevated by the escalation in the Middle East are reinforcing the economic case for adoption. Historically, sustained fuel-price pressure has been one of the strongest catalysts for accelerating EV uptake.⁶ Industrial electrification adds a further layer of uncertainty: Manufacturing reshoring and the electrification of industrial processes point towards structurally higher power demand that remains difficult to forecast with precision. Taken together, these competing demand vectors suggest that the supply cushion that seems comfortable could narrow far more quickly than current projections imply.


The bigger issue lies on the supply side and is rooted in US power grids, which struggle to keep pace with data-center interconnection demands. While physical construction typically takes under 24 months, grid-connection queues can more than double total

⁴ IEA STEO, April 2026 update⁵ Based on EIA-860M (February 2026 update)⁶ Global EV Outlook 2025 - IEA

project timelines, with hotspots like Northern Virginia reporting waiting times of up to seven years.⁷ This bottleneck reflects a deeper structural problem as both new generation capacity and new demand face the same overloaded interconnection queues (Figure 2). In Texas (ERCOT region), generation and storage projects queuing for grid connection nearly quadrupled between 2020 and 2024, while nationally active interconnection requests now stand at 1.84 TW, exceeding total US installed generating capacity. That figure overstates the genuine pipeline as only around 20% of requests ultimately reach commercial operation, with the remainder abandoned due to speculative filings, rising costs and mounting delays. Resolving these constraints requires grid expansion on a scale the US has not attempted in decades: In their 2024 National Transmission Planning study, the DOE estimates that in a lowest-cost scenario the total transmission system would increase by 2.1-2.6 times until 2050 and up to 3.3 times in a high demand case. This would imply building around 8000km of high voltage transmission lines per year, a pace that appears far beyond current construction rates, which have averaged well below 1,000 km annually in recent years.⁸

Getting there is complicated by constraints that span the entire supply chain. The grid equipment needed (transformers, switchgear and cables) faces lead times of up to four years alongside a 30% supply shortfall.⁹ Building new transmission lines compounds the problem further, taking four to eight years in advanced economies, a timeline fundamentally incompatible with the pace of data-center deployment. Beyond power infrastructure, the AI build-out is also running into semiconductor supply constraints: According to the IEA, HBM chip shortages have emerged as a binding bottleneck over the past six months, with insufficient supply of this specialized memory threatening to slow AI data-center deployment independently of grid availability.¹⁰ A construction workforce shortage of around 439,000 skilled workers and increasingly complex permitting processes compound the problem further. As a result, almost half of the approximately 12 GW of US data center capacity planned for 2026 is expected to be delayed or cancelled, with only around 5 GW currently under active construction.¹¹

Figure 2: US power capacity in queue for grid connection by electricity market region (in GW)



Source: LBNL, interconnection.fyi

⁷ [Epoch AI and Energy and AI](#)

⁸ [DOE and Grid Strategies](#)

⁹ [Wood Mackenzie](#)

¹⁰ [Key questions on energy and AI - IEA](#)

¹¹ [Slightline Climate Data and Bloomberg](#)

Meanwhile grid stress from data centers has moved from projection to reality, helping explain the prolonged integration assessments that slow their expansion. Unlike traditional commercial or residential load, data centers represent large, concentrated blocks of power demand that can disconnect or transfer to backup power in a coordinated manner during grid disturbances. Compounding this, AI-optimized data centers can exhibit sharper and more frequent load swings than conventional facilities: As workloads shift between high-intensity processing and lower-utilization periods, power draw can fluctuate rapidly, placing greater demands on frequency regulation and reserve markets. When a disturbance triggers a simultaneous disconnect, as occurred in Virginia in 2024, where 60 facilities representing 1.5 GW dropped off the grid at once, the sudden loss of load forces operators to curtail generation equally fast to avoid a frequency imbalance that can cascade into widespread blackouts.¹² The reverse risk is equally real: When data centers draw at full capacity during peak periods, they can overwhelm local transmission capacity, forcing power to reroute through congested lines and driving up congestion costs across the broader grid. This strain is increasingly forcing active intervention, as seen in January 2026 when the DOE invoked emergency powers to authorize PJM and ERCOT to shift data centers onto backup generators to prevent residential blackouts during a severe winter storm. Both failure modes reflect the same underlying mismatch: a grid designed for distributed, flexible demand is being asked to absorb loads that are geographically concentrated and operationally rigid.

These physical bottlenecks are now increasingly intersecting with political and social constraints.

Organized opposition to data-center development has emerged as an additional factor shaping the outlook. In 2025, an estimated USD156bn worth of data-center projects were blocked or delayed by local opposition, moratoriums and litigation in the US, with project cancellations more than quadrupling from six in 2024 to 25 in 2025.¹³ The opposition is largely rooted in concerns about electricity costs as the infrastructure required to accommodate data-center load growth is typically funded through broader rate-payer tariffs rather than charged directly to the facilities driving demand. In PJM alone, the grid's independent market monitor estimates that data-center load growth added USD16.6bn in capacity costs across the 2025/26 and 2026/27 delivery years, spread across 67mn consumers.¹⁴ Legislative responses are following, with moratorium bills introduced in at least 11 states in 2026, with some like Maine considering statewide construction pauses.¹⁵ While unlikely to halt the AI infrastructure buildout, rising opposition introduces permitting delays and regulatory uncertainty that compound existing supply-side constraints and underscores a broader question about who ultimately bears the cost of AI infrastructure.

¹² DCD

¹³ [Data Center Watch and Construction Dive](#)


¹⁴ [Comment of the Independent Market Monitor for PJM](#)

¹⁵ [Axios](#)

Where are data centers built and what determines their location?

Data center location reflects a specific set of infrastructure, cost and regulatory requirements that have historically concentrated capacity in a small number of markets. Today, Texas and Virginia alone account for 42% of total US data center capacity, though the nature of that concentration differs markedly between the two. Virginia, and Northern Virginia's Loudoun County in particular, dominates current installed capacity and represents the largest data-center cluster in the world. Texas, by contrast, carries the largest pipeline of planned and under-construction capacity, reflecting its abundance of land, a deregulated electricity market, and aggressive state-level incentives (Figure 3).

Figure 3: US data center capacity by county (in MW)



Source: National Laboratory of the Rockies.


Four factors dominate siting decisions. Power availability and cost are the most fundamental as cheap and reliable electricity underpins the economics of large-scale compute. Fiber connectivity to major internet exchange points determines latency and operational viability. Land availability and permitting speed determine how quickly large campuses can be developed. Tax incentives, particularly data-center exemptions offered by Virginia, Texas and several other states, have actively directed investment toward specific markets.

The consequence of this concentration is both operational and systemic risk. When grid capacity tightens in saturated markets, the AI infrastructure buildout slows not just regionally but globally, given the role these clusters play in serving international demand. A regional grid failure in these clusters would ripple globally, translating into higher prices for AI and cloud services, slower model deployment and capacity-driven rate limiting for end users. Overclustered markets also face compounding vulnerabilities across power supply, water availability, cooling infrastructure and skilled labor, making geographic diversification of future capacity an increasingly urgent priority.

While aggregate power-supply growth appears sufficient on paper, the regional picture is more nuanced. In the two most exposed markets, ERCOT (Electric Reliability Council of Texas) and PJM (PJM Interconnection), power generation growth has accelerated sharply since 2022, jumping from 10 TWh p.a. to 22 TWh p.a. in Texas and from 4 TWh p.a. to 31 TWh p.a. in PJM (Figure 4 a). This momentum is expected to continue, with the EIA's Short-Term Energy Outlook projecting an additional 99 TWh of supply in Texas and around 60 TWh in PJM by end-2027. On the demand side, Texas and Virginia each carry approximately 30 GW of data center capacity in their pipelines, which, if fully implemented, would imply an increase in annual electricity demand of around 140 TWh per region. In practice, however, given typical build-out timelines and project attrition, a more realistic estimate puts incremental data center demand at 30–60 TWh per region by 2030 (Figure 4 b). At those levels, both markets can realistically absorb expected AI-driven load growth, provided their supply pipelines materialize.

Delivering on that supply potential is, however, not without obstacles. In ERCOT, the supply pipeline is theoretically comfortable, but interconnection queues have grown sharply in recent years, raising the risk that capacity additions lag behind the pace of data center buildout. In PJM, the challenge is more structural. Generation growth has been slower and surrounding regions offer limited import relief, given mostly flat supply growth expectations. Data centers are also not the only source of incremental demand, with electrification of transport and industry adding further pressure on available grid capacity. With 105 GW currently queued in PJM's interconnection queue, processing delays alone could push a significant share of planned capacity beyond the 2030 horizon. Should capacity additions in either market fail to keep pace with demand, persistently tight grid conditions would translate into additional upward pressure on regional electricity prices.

Figure 4: Estimated power generation growth by grid region vs 2015 (lhs, TWh) and estimated data center power demand in Texas and Virginia (rhs, TWh)



Source: Allianz research based on EIA (STEO), IEA, LBNL and National Laboratory of the Rockies; Note: Planned power demand was estimated based on state level NLR capacity data for facilities planned or under construction. IEA, EIA and LBNL country level estimates were downscaled to regions based on planned capacity shares.



Beyond the baseline: Agentic AI, adoption velocity and the rebound effect

The future trajectory of AI's impact on power markets is shaped by two opposing forces. Rapid efficiency improvements are materially reducing the energy cost per token, while surging adoption, growing usage complexity and the emergence of agentic workflows are pushing demand in the opposite direction. Which force dominates will determine whether current supply projections prove adequate or significantly understated.

The energy efficiency of AI systems has improved dramatically across all layers of the technology stack... Compute performance per watt has increased by around 34% per year since 2008, with successive GPU generations delivering compounding gains.¹⁶ At the model level, architectural innovations such as Mixture-of-Experts designs, quantization, speculative decoding and model distillation have each contributed further reductions in energy per task, in some cases cutting compute

requirements by several multiples without meaningful loss in output quality. On the infrastructure side, hyperscale and AI-specialized data centers have become significantly more efficient at converting electricity into useful computation, with overhead losses from cooling and power distribution falling from around 40–50% of total facility energy in older facilities to just 10–20% in modern ones.¹⁷ According to the IEA, widespread adoption of existing AI applications across industry, buildings and transport could deliver aggregate energy savings of over 13 EJ by 2035, equivalent to around 3% of global final energy consumption.¹⁸ Realizing this potential, however, is far from guaranteed, as widespread adoption remains contingent on overcoming persistent barriers including fragmented data, limited digital skills and inadequate regulatory incentives.

¹⁶ [Epoch AI](#)


¹⁷ [LBNL](#)

¹⁸ [Key Questions on Energy and AI - IEA](#)

...yet demand is expanding at least as fast, driven by rapid adoption, growing model capability and a structural shift in how AI is used. Today, generative AI has reached 53% population-level adoption within three years of its mass-market introduction, faster than the personal computer or the internet, with 88% of companies reporting AI use in 2025 according to Stanford's AI Index.¹⁹ The models powering this growth have scaled at an equally striking pace, with the compute required to train frontier AI models growing at roughly 4–5x per year since 2020 and GPT-4 estimated at around 1.8trn parameters, roughly 10 times larger than GPT-3 released just three years prior.²⁰ Usage

intensity has followed the same trajectory. According to OpenRouter's platform data, average prompt length has grown roughly fourfold since early 2024 as users shifted from simple queries to context-heavy, multi-step workflows, with reasoning models growing from a negligible share to over 50% of token usage by late 2025.²¹ This growth is reflected in token consumption figures across major AI companies. OpenAI's API alone grew from 300mn tokens per minute in 2023 to over 6bn by late 2025, a 20-fold increase in under two years, while Google's Gemini token consumption increased approximately sevenfold between February and September 2025 (Figure 5).²¹

Figure 5: Google Gemini monthly token consumption (trillion tokens)



Sources: Google DeepMind; Alphabet Q2 2025 earnings.

The primary driver behind this acceleration is a fundamental shift in how AI is being used. Unlike conventional chat interactions, agentic systems plan and execute multi-step tasks autonomously, dramatically increasing token consumption per interaction. The energy implications are substantial, with an agentic task with reasoning consuming around 50Wh compared with roughly 0.3Wh for a standard text query, roughly 150 times more energy per interaction. Another reinforcing factor is the sharp decline in inference

costs, which fell more than 280-fold between late 2022 and late 2024.²³ As lower prices stimulate broader and more intensive use, efficiency gains at the model and hardware level risk being offset by a classic rebound effect, whereby falling costs per unit of consumption drive total consumption higher. Taken together, these trends suggest that AI-driven electricity demand is likely to continue growing faster than efficiency improvements can offset in the near term, reinforcing the supply-side pressures on regional power markets.


¹⁹ [The 2026 AI Index Report - Stanford HAI](#)

²⁰ [Epoch AI](#)

²¹ [OpenRouter](#)

²² Gemini's token growth partly reflects its integration into Google Search, which significantly expanded query volumes beyond standalone AI product usage.

²³ [The 2025 AI Index Report - Stanford HAI](#)




The hidden bill?: AI's impact on power prices

Data-center demand affects electricity prices through several compounding channels. Large concentrated loads drive up capacity market auction prices as grid operators must procure additional firm generation to meet anticipated peak demand, with costs socialized across all ratepayers. Grid-connection infrastructure represents a second channel as transformers, switchgear and dedicated transmission lines are typically recovered through broader utility tariffs rather than charged to the facilities driving demand. These pressures land on grids already strained by decades of underinvestment, amplifying the cost impact of each new interconnection. A third channel operates through generation mix: Data centers require firm, around-the-clock power that renewables alone cannot provide, pushing utilities toward more expensive gas and nuclear capacity. Finally, the renewable energy that data centers rely on for sustainability commitments is often located far from load centers, creating congestion costs that ripple across the broader grid. Network charges have so far remained relatively stable, but are structurally bound to rise in the regions where AI infrastructure is most concentrated. Because each channel operates on a different timescale and affects different consumer groups, the overall price impact is neither immediate nor evenly distributed across consumers and regions.

So far, aggregate state-level data tell a surprisingly benign story. Despite hosting more than 40% of total US data-center capacity, Texas and Virginia have seen power price inflation of only 18% and 20% respectively since 2022, below the 24% US average (Figure 6). This is consistent with research from Lawrence Berkeley National Laboratory (2026), which finds that state-level load growth from 2019 to 2025 was generally associated with lower all-sector average retail electricity prices as rising fixed infrastructure costs are spread across a larger base of consumers and kilowatt-hours sold.²⁴ A further moderating factor is that data centers have historically clustered in regions with favorable grid conditions and low base energy costs, insulating those markets from the price pressures that concentrated demand might otherwise create. Regulatory lag compounds this dynamic: Utility rate cases can take well over a year from filing to approval, and capital investments are recovered over asset lifetimes rather than immediately, meaning that retail prices today largely reflect costs incurred years earlier. That lag is, however, beginning to unwind: Investor-owned utilities filed USD18bn in rate increase requests in 2025, the highest level since the mid-1980s, with regulators approving 64% of requests.

²⁴ [Energy Markets & Planning](#)

Figure 6: Data-center capacity and retail electricity price inflation since 2022, by state

Source: Allianz Research based on EIA and NLR.

State-level data mask considerable heterogeneity

across individual utilities. A utility directly serving a major data-center cluster faces fundamentally different demand dynamics than a rural cooperative in the same state, yet both are folded into the same state-level price average. To move beyond the aggregate picture, we estimate the effect of data-center demand growth on residential electricity prices at the utility level, exploiting the fact that utilities entered the AI investment boom with very different pre-existing exposure to data center capacity.²⁵ The intuition is straightforward. Utilities that happened to serve counties with significant data center infrastructure were more directly in the path of the subsequent investment wave than those that did not. By interacting each utility's data center footprint with the national surge in construction investment – a design known as a Bartik shift-share instrument – we can estimate the price effect of data center demand while accounting for time-invariant utility characteristics and aggregate macroeconomic trends. The analysis covers approximately 1,200 US utilities over 2020–2024, drawing on EIA utility data matched to county-level data center capacity data from the National Laboratory of the Rockies (NLR) and national construction investment figures from the Census Bureau.

The results suggest that price pressures are more visible at the utility level than aggregate data imply, with three findings standing out. First, the price effect is concentrated entirely in the post-2023 period, precisely when AI-driven investment accelerated sharply, while the pre-2023 period shows no significant relationship between data center exposure and price changes (Table 1, column 2). This timing is difficult to explain by anything other than the AI demand shock itself, and provides confidence that we are capturing a genuine effect rather than a coincidental correlation with pre-existing trends. Second, the demand-side results validate the approach: utilities with higher data center exposure show markedly stronger commercial electricity consumption growth post-2023, consistent with hyperscale facilities connecting to utility grids and driving up load (Table 1, column 3). The effect is broad-based across utilities rather than driven by any single outlier. Third, in terms of magnitude, a utility with 4% of national data center capacity in its service territory, broadly representative of the most exposed markets outside the largest hubs, is associated with a cumulative residential price increase of approximately 3% over 2020–2024, based on the USD22bn growth in national data-center construction investment observed over that period. Modest in isolation, this effect is best understood

²⁵ Data center capacity shares reflect 2025 installed capacity, the large majority of which was planned and permitted before the generative AI boom took hold in 2023, making them a reasonable proxy for pre-existing exposure rather than a response to the shock we are trying to measure.

as an early signal rather than a final estimate – given the regulatory lag discussed above, much of the cost pressure already absorbed by utilities is still working its way through to consumers.


Table 1: Estimated effect of data-center demand on electricity prices

	Dep. var: log(Residential Price)		Dep. var: log(Commercial Sales/Customer)
	(1) Baseline	(2) Pre/Post Split	(3) Post-2023
Bartik IV	0.0342* (0.018)	—	—
Bartik IV × Post-2023	—	0.0321** (0.0156)	0.2416*** (0.0556)
Bartik IV × Pre-2023	—	0.0145 (0.021)	—
Observations	6,586	6,586	6,388
Within R²	0.0098	0.0115	0.033

Source: Allianz Research; Note: All specifications include utility fixed effects and year fixed effects. Standard errors are clustered at the utility level and reported in parentheses. The Bartik instrument is $B_{ut} = z_u \times G_t$, where z_u is utility u 's predetermined share of national data-center capacity and G_t is annual US data center construction investment (US Census Bureau). Specifications (1) and (2) use log residential prices as the dependent variable; (1) uses the full-period IV while (2) splits it at 2023 to test for pre-trends. Specification (3) uses log commercial sales per customer as the dependent variable with the post-2023 IV, providing demand-side validation of the instrument. Significance: *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Data-center demand is leaving a measurable imprint on electricity prices, with effects concentrated in the most exposed markets. Figure 7a) shows that while all utility groups experienced meaningful price increases over 2020–2024, the pattern is far from uniform. The no-exposure group saw the smallest increase at 14%, while the high-exposure group ended the period 24.5% above its 2020 level, a gap that widened noticeably after 2023 as AI-driven investment accelerated. Of that 24.5% rise, approximately 7.8pps are directly attributable to

data-center demand based on our regression estimates (Figure 7b)), with the remainder reflecting the broader energy cost pressures that affected all utilities over the period, including fuel-price volatility, rising infrastructure investment costs and general inflation. The difference between the two groups of around 10pps therefore reflects both the data-center premium and other structural differences between high and low-exposure utility markets.

Figure 7: a) Price index by exposure group (2020=100) and b) Observed vs data center-implied price change 2020-2024 (in %)

Source: Allianz Research based on EIA and NLR; Note: Exposure groups defined by utility share of national data center capacity: no exposure (n=718), low 0–1% (n=442), medium 1–5% (n=73), high >5% (n=5). Data center-implied effect = $\beta \times \text{avg. group capacity share} \times \text{DC investment growth (USD21.91bn)}$

Across the US, residential customers are already paying around USD1.4bn more per year on their electricity bills as a direct result of data-center demand. Just five utilities, serving 4.4mn households – roughly 4% of residential customers – in Northern Virginia, the Pacific Northwest and Arizona, account for more than 40% of the total, with their customers facing an average extra annual bill of USD139, attributable to data-center growth. A wider group of 73 medium-exposure utilities pays around USD28 per customer per year, while the vast majority of US utilities are not affected. Strikingly, the most affected markets are not pricier than the national average to begin with. Their historical 5% cost advantage has already narrowed to 3.7% by end-2024, and the gap is set to close further as investment accelerates, eroding a buffer that has so far masked the true scale of the data-center premium.



Table 2: Estimated financial impact of data center demand on residential electricity prices by utility exposure group in 2024

Metric	No exposure	Low (0–1%)	Medium (1–5%)	High (>5%)
Utility characteristics				
Number of utilities	718	442	73	5
Avg. exposure share (%)	0	0.17	2.3	10.47
Market size				
Avg. residential price (USD/kWh)	0.143	0.161	0.159	0.147
Total sales (TWh)	203.2	724.5	247.4	52.5
Total customers (mn)	17.3	71.42	24.08	4.37
Estimated price impact				
Estimated price effect (%)	0	0.13	1.72	7.84
Extra bill per customer (USD/yr)	0	2.1	28.05	138.65
Extra bill total (USD mn)	0	150	675	605
Inflation impact				
Direct CPI contribution (pp)	0	0.0032	0.0427	0.1945


Source: Allianz Research based on EIA and NLR

Translated into aggregate consumer prices, our findings confirm an overall modest but highly skewed data-center price premium. Across the full sample, the sales-weighted average residential price increase attributable to data centers amounts to around 0.6% over 2020–2024, translating into a direct headline CPI contribution of just 0.015pp. For the most exposed utilities, however, the same channel adds roughly 0.19pp to headline CPI, around thirteen times the full-sample figure.²⁶ Because these estimates capture only residential electricity and exclude commercial pass-through, they should be read as a lower bound on the true inflation impact, with the full effect on aggregate consumer prices likely materially larger once higher input costs for electricity-intensive sectors are reflected in final goods prices.

This premium is likely to grow as the investment wave continues. Data-center investment already grew by 32% in 2025 and is set to increase by a further 75% in 2026 alone.²⁷ As pending rate cases translate accumulated cost pressures into retail tariffs, our estimates point to an additional price effect of close to 14pps for the most exposed utilities over 2025–2026, equivalent to a further CPI contribution of around 0.34pp, nearly doubling the cumulative four-year effect in just two years. With infrastructure costs still largely socialized across all rate-payers rather than borne by the facilities driving demand, the question of who ultimately pays for the AI buildout is set to become increasingly central to the policy debate.

²⁶ Direct CPI contribution estimates assume an electricity share of 2.48% in the consumer price basket, based on [BLS CPI](#) relative importance weights (December 2024)

²⁷ [Key Questions on Energy and AI - IEA](#)



Preparing AI for energy and energy for AI: policy recommendations

The scale and pace of AI-driven power demand growth points to a set of policy priorities that are urgent, tractable and largely absent from current frameworks, which would help lay the foundation for AI ambitions that the grid can actually support.

1) Expand and accelerate

- **Reform interconnection queues.** Interconnection queues are the single most binding near-term constraint on both generation capacity and data-center deployment. Streamlining queue processes, introducing binding timelines for interconnection studies, strengthening financial penalties for speculative filings that inflate backlogs and delay genuine projects and prioritizing shovel-ready projects with firm power commitments would significantly accelerate the pace at which new capacity reaches the grid.

2) Govern & allocate

- **Improve transparency and demand tracking.** Mandatory disclosure of data-center energy consumption, water use and interconnection requests would materially improve capacity planning. Incorporating large-load interconnection requests into publicly accessible data systems would make the pace and geography of data-center expansion visible to regulators, utilities and the public in real time.
- **Address overclustering risks in saturated markets.** The concentration of approximately 42% of US data-center capacity in Texas and Virginia creates systemic risks that extend beyond local grid stress. Overclustered markets face compounding vulnerabilities across power supply, water availability, cooling infrastructure, and skilled labor, while their position as critical nodes in

global AI value chains means that a regional grid failure or prolonged supply disruption could have disproportionate consequences for AI deployment globally. Policymakers and grid operators should assess concentration thresholds beyond which new approvals in saturated markets require enhanced scrutiny, consider incentive structures that actively redirect investment toward underutilized regions with available grid capacity, and prioritize locations with proximity to renewable energy resources to reduce long-run dependence on firm fossil fuel generation.

- **Require data centers to internalize their infrastructure costs.** The socialization of capacity costs onto existing ratepayers is both economically inefficient and politically unsustainable. Policymakers should establish frameworks requiring large new loads to bear a proportionate share of the transmission and capacity infrastructure their demand necessitates. Time-of-use and peak-pricing tariffs that reflect system stress would more accurately align data-center costs with the grid impacts they create, reducing the cross-subsidization of large loads by residential ratepayers. Where market-based pricing mechanisms are insufficient, direct compensation requirements, obliging data centers to offset the costs their demand imposes on host communities, provide an alternative route to addressing the equity dimension and reducing the political friction that is increasingly delaying project approvals.

3) Optimize & decarbonize

- **Implement mandatory energy-efficiency standards.** No binding energy performance standards currently apply to data centers in most jurisdictions despite their growing share of national electricity consumption. Minimum PUE thresholds, water-usage effectiveness requirements and mandatory energy intensity reporting would establish a baseline accountability framework, with stricter requirements for new hyperscale builds where best-in-class efficiency is already technically and economically achievable.²⁸
- **Develop demand-flexibility frameworks.** Full curtailment of data center operations is economically impractical given facility capital intensity, but partial demand flexibility is more tractable and largely untapped. Batch workloads such as model training and non-time-sensitive inference can in principle be shifted to off-peak periods, reducing peak demand pressure and congestion costs without compromising operational continuity. Regulators and grid operators should develop demand-response frameworks that reward this flexibility through capacity market mechanisms or dynamic pricing signals.
- **Ensure AI powers the energy transition rather than slowing it.** The risk that AI-driven data-center growth locks in years of additional fossil-fuel generation is real and underappreciated. Behind-the-meter generation, power purchase agreements tied to additionality requirements, and prioritization of nuclear and long-duration storage for firm power needs would reduce this risk. AI infrastructure and decarbonization are not inherently in tension. The former should be designed to accelerate the latter.

²⁸ [Data Centers and Their Energy Consumption - Congressional Research Service](#)

A photograph showing a variety of hands of different skin tones stacked on top of each other, resting on a thick, textured tree branch. The background is a lush green forest with sunlight filtering through the leaves. The text 'Our team' is overlaid in the center, with 'Our' in white and 'team' in orange.

Our team

Chief Investment Officer
& Chief Economist
Allianz Investment Management SE



Ludovic Subran
ludovic.subran@allianz.com

Head of Economic Research
Allianz Trade



Ana Boata
ana.boata@allianz-trade.com

Head of Macroeconomic & Capital
Markets Research
Allianz Investment Management SE



Bjoern Griesbach
bjoern.griesbach@allianz.com

Head of Outreach
Allianz Investment Management SE



Arne Holzhausen
arne.holzhausen@allianz.com

Head of Corporate Research
Allianz Trade



Ano Kuhanathan
ano.kuhanathan@allianz-trade.com

Head of Thematic & Policy Research
Allianz Investment Management SE



Katharina Utermoehl
katharina.uterhoehl@allianz.com

Macroeconomic Research



Lluís Dalmau Taules
Economist for Africa & Middle East
lluis.dalmau@allianz-trade.com



Maxime Darmet Cucchiari
Senior Economist for UK, US & France
maxime.darmet@allianz-trade.com



Jasmin Gröschl
Senior Economist for Europe
jasmin.groeschl@allianz.com



Françoise Huang
Senior Economist for Asia Pacific
francoise.huang@allianz-trade.com



Maddalena Martini
Senior Economist for Southern Europe & Benelux
maddalena.martini@allianz.com

Outreach



Luca Moneta
Senior Economist for Emerging Markets
luca.moneta@allianz-trade.com



Giovanni Scarpato
Economist for Central & Eastern Europe
giovanni.scarpato@allianz.com



Heike Baehr
Content Manager
heike.baehr@allianz.com



Maria Thomas
Content Manager and Editor
maria.thomas@allianz-trade.com

Corporate Research



Guillaume Dejean
Senior Sector Advisor
guillaume.dejean@allianz-trade.com



Maria Latorre
Sector Advisor, B2B
maria.latorre@allianz-trade.com



Maxime Lemerle
Lead Advisor, Insolvency Research
maxime.lemerle@allianz-trade.com



Sivagaminathan Sivasubramanian
ESG and Data Analyst
sivagaminathan.sivasubramanian@allianz-trade.com



Pierre Lebard
Public Affairs Officer
pierre.lebard@allianz-trade.com

Thematic and Policy Research



Michaela Grimm
Senior Economist,
Demography & Social Protection
michaela.grimm@allianz.com



Patrick Hoffmann
Economist, ESG & AI
patrick.hoffmann@allianz.com



Simon Krause
Economist, ESG & Insurance
simon.krause@allianz.com



Hazem Krichene
Senior Economist, Climate
hazem.krichene@allianz.com



Kathrin Stoffel
Economist, Insurance & Wealth
kathrin.stoffel@allianz.com



Markus Zimmer
Senior Economist, ESG
markus.zimmer@allianz.com

Recent Publications

[12/05/2026 | Behind the gate: The promises and perils of private markets democratization](#)
[06/05/2026 | US large banks: The peak of the cycle is not the time to be complacent](#)
[05/05/2026 | Automotive: Will the Middle East crisis supercharge EV momentum?](#)
[29/04/2026 | Staycation summer? Jet-fuel crunch reshapes the peak holiday season](#)
[23/04/2026 | Energy shock and policy response: Once bitten, twice shy?](#)
[22/04/2026 | Global Insolvency Outlook: Brace for Middle East spillovers](#)
[31/03/2026 | Economic outlook 2026-27: The Fog of War](#)
[25/03/2026 | AI capex cycle: war-proof for now](#)
[23/03/2026 | Signal without response: Why the EU ETS needs resolve, not redesign](#)
[17/03/2026 | Not all Emerging markets are equal: Hormuz, triple deficits, and the new energy risk premium](#)
[16/03/2026 | Warsh's double dilemma: when the Middle East rewrites the Fed's playbook](#)
[11/03/2026 | Allianz Social Resilience Index 2025: The Middle-Resilience Trap](#)
[10/03/2026 | The second energy shock: Why Europe still isn't energy secure](#)
[05/03/2026 | Closing the gender income gap: from paycheck to pension](#)
[03/03/2026 | Conflict in the Middle East: Implications for markets and macro](#)
[25/02/2026 | Mining for the future: Addressing liabilities and unlocking sustainable transition opportunities](#)
[24/02/2026 | Variable geometry for European trade: Building resilience and diversification](#)
[23/02/2026 | Schroedinger's tariffs](#)
[20/02/2026 | Private equity in transition: from distribution drought to selective recovery](#)
[19/02/2026 | Eurobonds – A window of opportunity for a strategic necessity](#)
[16/02/2026 | Country Risk Atlas 2026: Under the surface](#)
[13/02/2026 | Fragmentation tests US market primacy and reshapes the global investment landscape](#)
[11/02/2026 | High prices, thin buffers: America's affordability crisis persists](#)
[10/02/2026 | Europe's households after the rate shock: A windfall for some, a squeeze for others](#)
[04/02/2026 | Team Italy: An economic performance worthy of a gold medal?](#)
[04/02/2026 | Five things that could derail the ECB](#)
[02/02/2026 | A new decade high for major insolvencies driven by services, retail and construction](#)
[29/01/2026 | From Japan with love: New policy stance creates both market opportunities and liquidity risks](#)
[27/01/2026 | Trade receivables in a fragmented world: Navigating collection complexity](#)
[26/01/2026 | EU-India trade deal: EUR30bn of combined yearly exports gains in a fragmented world](#)
[23/01/2026 | Eyes back on the Fed \(and on interventionist financial policies\)](#)
[21/01/2026 | Tackling the insurance protection gap](#)
[20/01/2026 | The heat is on: Unlocking Germany's heat-pump potential](#)
[16/01/2026 | Geopolitics heats up from Venezuela, to Greenland to Iran, but investors shrug. For how long?](#)
[14/01/2026 | Allianz Risk Barometer - Identifying the major business risks for 2026](#)
[17/12/2025 | Economic Outlook 2026-27: Stretching the limits](#)
[11/12/2025 | What to watch](#)
[10/12/2025 | Convertible bonds: The Financial Roadster for Dynamic Markets](#)
[05/12/2025 | What to watch](#)
[02/12/2025 | High hopes, heavy footprint: Aviation's quest for climate-neutral skies](#)

Discover all our publications on our websites: [Allianz Research](#) and [Allianz Trade Economic Research](#)

Director of Publications
Ludovic Subran, Chief Investment Officer & Chief Economist
Allianz Research
Phone +49 89 3800 7859

Allianz Group Economic Research
https://www.allianz.com/en/economic_research
<http://www.allianz-trade.com/economic-research>
Königinstraße 28 | 80802 Munich | Germany
allianz.research@allianz.com

 @allianz

 allianz

Allianz Trade Economic Research
<http://www.allianz-trade.com/economic-research>
1 Place des Saisons | 92048 Paris-La-Défense Cedex | France

 @allianz-trade

 allianz-trade

About Allianz Research

Allianz Research encompasses Allianz Group Economic Research and the Economic Research department of Allianz Trade.

Forward looking statements

The statements contained herein may include prospects, statements of future expectations and other forward-looking statements that are based on management's current views and assumptions and involve known and unknown risks and uncertainties. Actual results, performance or events may differ materially from those expressed or implied in such forward-looking statements. Such deviations may arise due to, without limitation, (i) changes of the general economic conditions and competitive situation, particularly in the Allianz Group's core business and core markets, (ii) performance of financial markets (particularly market volatility, liquidity and credit events), (iii) frequency and severity of insured loss events, including from natural catastrophes, and the development of loss expenses, (iv) mortality and morbidity levels and trends, (v) persistency levels, (vi) particularly in the banking business, the extent of credit defaults, (vii) interest rate levels, (viii) currency exchange rates including the EUR/USD exchange rate, (ix) changes in laws and regulations, including tax regulations, (x) the impact of acquisitions, including related integration issues, and reorganization measures, and (xi) general competitive factors, in each case on a local, regional, national and/or global basis. Many of these factors may be more likely to occur, or more pronounced, as a result of terrorist activities and their consequences.

No duty to update

The company assumes no obligation to update any information or forward-looking statement contained herein, save for any information required to be disclosed by law.

Allianz Trade is the trademark used to designate a range of services provided by Euler Hermes.

Allianz Research | 6 May 2026

US large banks: The peak of the cycle is not the time to be complacent

In Summary

Ziqi Ye
Investment Strategist
ziqi.ye@allianz.com

Maxime Darmet
Senior Economist US-UK-France
maxime.darmet@allianz-trade.com


Romain Poncet
Research Assistant
romain.poncet@allianz-trade.com

- **US large-cap banks posted record Q1 2026 earnings, and both earnings growth and asset quality sits well above trend. Yet, investors seem wary about how long the good times can last, for at least four reasons.** Bank equities underperformed the S&P 500 and credit spreads widened versus non-financial peers. The market is looking beyond balance sheets and pricing a gap between reported solvency and true loss-absorption capacity.
- **First, buybacks offer short-term support but have limits.** The Big Six returned over USD110bn in 2025 and ~USD32bn in Q1 2026 alone via share repurchases, mechanically inflating EPS levels. But buybacks support EPS growth only as long as they continue: once excess capital is exhausted, the growth tailwind disappears with a cliff effect, leaving valuations resting entirely on business growth. Elevated trading income and the M&A cycle are both volatile and unlikely to persist at current levels.
- **Second, regulatory softening reduces buffers.** The March 2026 Basel III/GSIB/stress-test proposals will cut CET1 requirements for the largest banks by ~5–6% cumulatively. Banks have already positioned for this: CET1 ratios for the Big Six fell ~100bps y/y, and this direction will persist as big banks still have rich buffer of around 2pps against the minimum required CET1 ratio.
- **Third, SRTs artificially lift reported CET1.** Synthetic risk transfers (SRT) have grown five-fold since 2016 (~EUR800bn outstanding), allowing banks to shed risk-weighted assets and boost CET1 ratios without reducing actual exposure. “Organic” solvency, which is not at the mercy of high-risk-buyers, is lower than headline ratios suggest – an estimated gap of 43bps from the CET1 ratio, plus the procyclical and liquidity risk that have not been tested yet.
- **Last, long financial cycles analysis points to weaker fundamentals for the banking industry in the years ahead.** Based on BIS methodologies, potential financial strains can be detected by examining credit-to-GDP, real credit growth, and residential property prices over 8–30 year periods. This framework, which successfully identified the two last U.S. financial or banking crises (the S&L crisis and the GFC), helps spot out cycle reversals that typically precede stress, though timing remains uncertain. Current signals indicate the U.S. long financial cycle has passed its peak and entered a downturn phase between 2021 and 2023. Emerging strains may extend beyond banks to non-bank financial institutions, given their larger role in the economy. Although real corporate credit growth has been picking up recently, and credit standards have loosened, most other segments (such as mortgages) have been subdued for several years now, suggesting that the US economy is indeed entrenched in the downward phase of the long financial cycle.
- **The conclusion is asymmetric.** Equity prices may capture the upside created by continuous share buybacks in the short term. But the valuations, which already embed rate-driven net interest income (NII), buyback-fueled EPS growth and M&A super cycle tailwinds, are fragile and leave no room for disappointment. Credit spreads have widened modestly but do not compensate for structurally thinner capital buffers. The largest banks may be well-positioned to weather a downturn; the question is whether their investors are being paid enough for the risk that one arrives.

Is it the peak of the cycle for US banks?

America's large banks entered Q2 2026 in their strongest collective shape since before the financial crisis. Five out of the biggest six banks beat Q1 consensus estimates, supported by high trading income and the M&A and IPO cycle. Bank of America's equities trading income jumped 30%, representing its best trading quarter in 15 years. JPMorgan's markets revenue hit a record USD11.6bn, with fixed income up 21% and equities up 17%. Goldman Sachs' investment banking fees grew 48% year-over-year to USD2.8bn. However, Wells Fargo stood out, with the NII revenue miss sending its stock price down 5-6% on results day. Looking past the decade, the current cycle is clearly well above the long-term trend. Plotting the average earnings per share of the Big Six against its long-run trend shows that the group's average EPS now sits materially above trend at levels that have historically preceded compression rather than continuation.


Figure 1: Earnings cycle goes above long-term trend



Sources: Bloomberg, Allianz Research


Nevertheless, markets are wary about how long the good times can last. The recent market reaction demonstrated more concerns than confidence. Since earnings season kicked off on 13 April, US large banks (KBW Bank Index) have dropped 0.6%, and underperformed the broader index (S&P 500) by more than 5pps. In credit, the divergence is subtler but more structurally revealing: US bank and broker spreads have widened relative to non-financial investment-grade issuers (keeping ratings and maturities equal). Investors have started to demand more compensation to hold bank credit than comparable corporates.

Figure 2: US bank stocks performance year to date (rebased to 100)



Sources: Bloomberg, Allianz Research


Figure 3: Credit performance - US banks credit premium is picking up



Sources: LSEG Datastream, Allianz Research

In the short term, the large scale of share buybacks is a cause for concern. America's six largest lenders returned more than USD110bn to shareholders in 2025, and Q1 2026 suggests the pace is accelerating with around USD32bn total repurchased. Share buybacks are, mechanically, among the most direct levers available to support earnings per share and the stock price: the fewer the shares outstanding, the higher earnings per share – an immediate signal of management confidence. But this is capital optimization, not capital creation. As buyback programs slow – whether from unstable trading revenue or simply from the pool of excess capital being exhausted – the mechanical EPS support they have provided will diminish. The equity story then rests entirely on the underlying earnings power of the business. At current multiples, that is a thin margin of safety because the elevated trading income and investment banking income is inherently volatile and not guaranteed to persist.


Figure 4: US Big Six banks distributed more than USD120bn capital in 2025 by share repurchases



Sources: Bloomberg, Allianz Research

For credit investors, it hurts more. The equity on a bank's balance sheet is the buffer that absorbs losses before bondholders are impaired. Each repurchase, executed at a premium to book, reduces that buffer in absolute terms. The Tier 1 capital ratio across the US banking system, which reflect banks' credit quality, is now turning for the first time after the interest rate shock in 2022. Tier 1 capital adequacy fell 13bps in the last quarter of 2025, reflecting the capital return program and signaling the peak of the credit cycle. And the new Basel III proposal announced last month was explicitly designed to push it further down.


Figure 5: US banks regulatory Tier 1 capital to risk-weighted asset



Sources: LSEG Datastream, Allianz Research

In the medium term, deregulation is reducing buffers. The turning trend of bank Tier 1 capital adequacy brings us to the mid-term concern of deregulation. On 19 March, the Fed, OCC and FDIC jointly issued three proposals that cover Basel III Endgame, the GSIB surcharge and stress testing. The Basel III component revises capital requirements specifically for the largest banks (Category I and II banks) – setting higher requirements on trading activity while deliberately lowering risk weights on traditional lending. However, the GSIB surcharge revision reduces additional capital charges on the eight largest systemic institutions. Combined with proposed stress-testing changes, CET1 requirements for the largest banks will decline by 4.8% on a cumulative basis, while Tier 1 requirements will fall by approximately 6%. Banks had anticipated the direction and positioned accordingly. Ahead of formal confirmation, the Big Six had built an average CET1 buffer of 2pps above regulatory minimums to guard against the threat of the original, more punitive 2023 rules. With that threat now formally lifted, the incentive to maintain the excess is gone. CET1 ratios fell by an average of nearly 100bps y/y – Goldman Sachs and Wells Fargo led the compression – with excess capital flowing directly into the buyback programs detailed above.

Figure 6: Big Six CET1 Ratio Q1 2026 vs Q1 2025, y/y change in bps




Sources: Bloomberg, Corporate reports, Allianz Research

The second engine is less visible but runs parallel. Even before deregulation formalized the relief, banks had been quietly optimizing regulatory capital through synthetic risk transfers (SRTs). Through these structures, banks transfer the credit risk of a tranche of their loan book to private investors – typically private credit funds, insurers or hedge funds – retaining the loans on their balance sheets while the CET1 ratio improves. Globally, SRT issuance has grown fivefold since 2016, with outstanding loan portfolios protected reaching almost EUR800bn as of end-2024¹. European banks have traditionally dominated the market, but US banks expanded this market rapidly in 2023 and

¹ International Association of Credit Portfolio Managers (IACPM) (2025): Global SRT Bank Survey 2016–2024.

2024 and account for roughly 20% of global issuance volume by the end of 2024.² While SRTs volume is still small relative to bank balance sheet, procyclicality risk and liquidity risk that haven't been tested by a credit downturn are the core problems. In a broad economic downturn, banks may find credit protection expensive or unavailable when capital needs are greatest – forcing a simultaneous pullback in lending that amplifies the very downturn it coincides with. The liquidity dimension compounds this when open-ended funds investing in SRTs face redemption pressure in stress.

Figure 7: Underlying pool size by issuer region (EUR bn)




Source: Basel Committee on Banking Supervision (BCBS), Allianz Research

Regulatory relief lowered the floor, while SRTs removed the risk from headline ratios. Combining those two engines, the financial system is not as well-capitalized as its headline ratios suggest: everything that has been moved into the shadows should also be deducted. When stress surfaces, the gap that is invisible on paper will become visible in prices.

The opacity of fast-growing private lending has created another risk in the longer term. In response to growing market concerns on private credit, the Big Six have voluntarily disclosed more detail on their non-bank financial institution (NBFI) lending in recent quarters. Without a certain reporting standard, banks brought similar reassurance – NBFI exposure is growing slowly but the size is still relatively small.

Figure 8: Increasing amount of US bank loans to NBFIs



Source: FED, Allianz Research


² Basel Committee on Banking Supervision (BCBS) (2026): *The rise and risks of synthetic risk transfers*.

Detecting downturn risks

Long financial cycle analysis offers valuable insights to detect potential banking strains before they occur. We build on the work from the BIS³ which seeks to characterize empirically the booms and busts of the financial cycle since the 1970s. The analysis combines two analytic approaches: turning-point analysis and frequency-based filters. It looks at three macro-financial variables that are found to best capture the medium-term financial cycle, lasting between 8 and 30 years: the credit-to-GDP ratio, the growth of real credit and the growth of real residential property prices. Only credit to the private credit is included. When the statistical analysis identifies a consistent turning of the cycle across these macro-financial variables, including past the trough, financial strains are likely to occur in the next couple of quarters or years – the timing being uncertain. The long financial cycle has been a good indicator to anticipate the two banking crises in the US over the past 40 years: the L&S savings crisis starting in 1988-89 and the Global Financial Crisis (GFC) starting in 2007. Unfortunately, a big caveat is that there have been only two occurrences of banking or financial crises over the past couple of decades, so the empirical regularity of the signal has to be taken cautiously. Nevertheless, the long financial cycles have also been good predictors of banking and/or financial strains in other markets, which leads us to take these signals seriously.

The financial cycle now signals weakness for the finance industry ahead. Figure 9 depicts the financial cycle for the US since the 1970s. It shows that we are past the peak of the cycle recently, with both signals converging, and signaling the beginning of a downturn phase in the cycle, which had begun between 2021 (turning-point analysis) and 2023 (filter method). Figure 10 shows the evolution of the three underlying variables. Since 2021, the year-on-year growth slows down, or even becomes negative since 2024-25. The strains looming in the finance industry are not necessarily bank-driven. The analysis looks at aggregate measures of credit. Supply of credit by banks makes up only 30% of the total supply of credit. Hence, strains could also emerge in non-bank financial institutions (NBFIs) rather than banks, though the links between the two are closely tied.


Figure 9: Long financial cycle in the US



Sources: BIS, Allianz Research

³ Characterizing the financial cycle: don't lose sight of the medium term!, BIS Working Papers No 380, June 2012.


Figure 10: Growth y/y (%) of the components of the financial cycle and GDP



Sources: BIS, Allianz Research

The turning of the long financial cycle has just begun and signals further rises in delinquencies and charge-offs ahead. The turning of the financial cycle reflects the combination of potentially several factors but are mostly determined by rising Fed interest rates when the central bank strives to reduce inflationary pressures or financial exuberance (loose credit standards and increased supply of credit to risky products). Higher interest rates are gradually feeding into higher rates to the private sector and tightening credit standards as borrowers become weaker amid rising interest expenses. In turn, curtailed supply of credit and higher interest rates lead to decelerating real property prices, pressuring down the value of collateral for households and some corporates, which further exacerbate the credit downturn. This is related to the financial accelerator theory. Eventually, it means typically slower growth of banks' revenues, and increasing costs related to provisions as more borrowers struggled to repay their debt. Figure 11 shows that after the financial cycle turns down (with the turning point method sending a signal before the filter method in the last episode), delinquencies and charge off rise. In the current environment, private credit defaults would appear much higher if borrowers held to public market standards. Figure 12 shows that including distressed exchanges realized default rates on leveraged loans would indeed have been much higher, corresponding well to the current trend of the financial cycle.


Figure 11: Financial cycle compared to the evolution (basis 100) of the charge-off rate (on all loans) and delinquencies (on all loans and leases, all commercial banks)



Sources: FED, BIS, Allianz Research

It is important note that this analysis captures long-term trends: within this long cycle, we are currently witnessing easier credit standards and a pickup of credit creation, suggesting that strains are not imminent. For instance, the Fed SLOOS survey indicates that banks have loosened credit standards for corporate loans since early 2023 and that real credit growth has accelerated from end-2025. But those are unlikely to be sustained if the downturn phase of the long financial cycle is entrenched. Our analysis suggests that this short term cycle will reverse soon.

Figure 12: Financial cycle compared to default rates on leveraged loans including distressed exchanges



Sources: FED, PitchBook, BIS, Allianz Research

Note: The default rate on leveraged loans including distressed exchanges is calculated as the number of issuers in default or distressed exchange over the past 12 months divided by the total number of issuers at the beginning of the 12-month period. These loans are issued to highly leveraged companies and are more sensitive to changes in the financial cycle.

What does this mean for investors?

US large-cap banks have entered mid-2026 in their strongest operational position since 2008 – earnings at cycle highs, balance sheets well-capitalized, capital returns at a record pace. However, as discussed above, the forces pushing the cycle turning are cumulating. Beyond the cycle, structural headwinds form from digitization, where low-cost disrupters are manifold. The American Bankers Association warned that a stablecoin market scaling toward USD2trn could drain up to 10% of core bank deposits, raising funding costs by roughly a quarter of a percentage point – a modest but directionally unwelcome addition to an NII outlook that is already plateauing⁴. At the same time, current equity valuations and credit spreads offer investors almost no cushion against disappointment.

Share buybacks and high ROEs will keep benefiting equity holders in short-term. Current valuations price in three tailwinds continuing simultaneously – rate-driven NII, a buyback-fueled EPS ramp-up and a M&A super-cycle. But none are guaranteed in the long-term. The market has already paid for the good news but has not priced in the medium-term and long-term cycle turning pressure.

With tight spreads, the upside is more limited for credit investors. Downside risk would be amplified if the large amount of capital distributed to shareholders and the persisting downtrend of solvency has been billed to them, as lower required capital means a thinner equity buffer sitting above bondholders. Though solvency is still healthy and far from triggering concerns even after the proposed capital release, credit investors are not being asked to absorb such structurally lower protection at historical thin compensation. Senior bank spreads have widened relative to non-financial investment-grade peers, but only modestly, and not in a way that reflects the structural shifts.

The banks themselves, in launching CDX Financials index, have built another instrument to make the shadow system they created legible and hedgeable — and declined to include themselves in its perimeter. The largest banks, at their peak, are still well-positioned to weather the downturn. **The question the market is beginning to ask in the gap between record earnings and cautious share prices is whether their investors are being adequately paid to stand in the rain with them.**

⁴ Huther, J. and Wang, Y. (2025). "How stablecoins could affect borrowing costs for the government, businesses and households." ABA Banking Journal, August 2025.

These assessments are, as always, subject to the disclaimer provided below.

FORWARD-LOOKING STATEMENTS

The statements contained herein may include prospects, statements of future expectations and other forward-looking statements that are based on management's current views and assumptions and involve known and unknown risks and uncertainties. Actual results, performance or events may differ materially from those expressed

or implied in such forward-looking statements.

Such deviations may arise due to, without limitation, (i) changes of the general economic conditions and competitive situation, particularly in the Allianz Group's core business and core markets, (ii) performance of financial markets (particularly market volatility, liquidity and credit events), (iii) frequency and severity of insured loss events, including from natural catastrophes, and the development of loss expenses, (iv) mortality and morbidity levels and trends,

(v) persistency levels, (vi) particularly in the banking business, the extent of credit defaults, (vii) interest rate levels, (viii) currency exchange rates including the EUR/USD exchange rate, (ix) changes in laws and regulations, including tax regulations, (x) the impact of acquisitions, including related integration issues, and reorganization measures,

and (xi) general competitive factors, in each case on a local, regional, national and/or global basis. Many of these factors may be more likely to occur, or more pronounced, as a result of terrorist activities and their consequences.

NO DUTY TO UPDATE

The company assumes no obligation to update any information or forward-looking statement contained herein, save for any information required to be disclosed by law.

Allianz Trade is the trademark used to designate a range of services provided by Euler Hermes.

Allianz Research | 5 May 2026

Automotive: Will the Middle East crisis supercharge EV momentum?

In Summary

Guillaume Dejean
Senior Sector Advisor
guillaume.dejean@allianz-trade.com

Hazem Krichene
Senior economist Climate
hazem.krichene@allianz.com


Rozalia Papaioannou
Research assistant, Climate
rozalia.papaioannou@allianz.com

- **An electric tilt boosted by energy volatility.** After a bruising 2025, Q1 2026 data reveal a striking reversal. BEV market share hit 19% EU-wide (+4pps vs; Q1 2025) and surged to 28% in France and 23% in Germany. The 30% oil spike fueled by Middle East tensions has sent pump prices back above EUR 2.0/L, reviving the cost-of-ownership debate and accelerating a shift that was already structurally underway. This is notably thanks to price compression over other powertrain models on primary markets and rising supply on second-hand markets.
- **Consumers are more sensitive than ever before to energy costs.** Fuel costs are hitting European households already squeezed by a post-Covid cost spiral across all car-related services with repair, maintenance and parts costs up between 20-30% since 2021, 10% higher than the increase in average disposable income. Car-usage costs absorb 7–8% of disposable income on average in France and Germany but rise to 11% for lowest-income quintiles during periods of volatility (e.g. 2022), making fuel-consumption costs a critical issue for low-income households. Currently, switching to BEVs would offer material energy consumption savings – equiv. to a 4–5% purchasing-power gain per capita on average in Western Europe – with an energy-cost differential range estimated at 5-7x regardless of price volatility.
- **But turning current momentum into a durable energy shift will require firing on all four policy cylinders in parallel: increasing local battery production, building sufficient grid infrastructure implementing effective carbon pricing and consistent subsidies.** While there have been positive signs of structural progress (e.g. average battery range crossing 500km psychological barrier, faster charging times), Europe remains exposed to China's technological dominance on batteries and powertrains. Its infrastructure deficit gets structural, with only 1.1mn charging points in Q1 2026 (mostly concentrated in 4 countries) – far from the European Commission's 3.5mn target by 2030. Moreover, AI-driven data-center expansion will compete directly with EV electrification for constrained grid capacity: EU consumption is projected to grow from 70TWh in 2024 to 115TWh by 2030. To reach its electrification targets by 2030, Europe needs carbon pricing that makes the fossil-fuel cost signal permanent, purchase subsidies that bridge the affordability gap for mass-market buyers, technological acceleration that makes BEVs the cheaper option without government support and grid decarbonization that ensures electrification delivers on its promise.
- **We find that even the most generous subsidy scenario (at least EUR5,000) reaches only 70% BEV share by 2030.** Carbon pricing under NDC commitments lifts the EU BEV share from roughly 29% to 42% by 2030 – meaningful, but still not enough to be on the Net Zero target which estimates 79% of BEV share in the EU by 2030. On the technology side, battery costs have fallen -93% since 2010 and are heading toward USD 60–70/kWh by 2030, the level at which EVs become cheaper than combustion vehicles in most segments without any grant. But ultimately an EV is only as clean as the electricity it runs on: even full fleet electrification falls short of its environmental potential without a cleaner grid. Moving Europe's low-carbon electricity share from 70% to 80% by 2035 would cut per-vehicle EV emissions by over 40%.
- **What if energy volatility continues beyond Q2?** EV resiliency is likely to be tested, notably as the progressive phase-out of subsidies reduces public support buffers against energy shocks. However, relatively strong purchasing power among EV buyers, ongoing technology improvements and intensified competition from Chinese OEMs should continue to underpin demand. The main downside risk lies upstream: renewed semiconductor supply disruptions could disproportionately impact EV production and pricing, given their higher chip intensity.

An electric tilt boosted by energy-price volatility?


Another energy crisis or another reason to switch to EVs? The recent 30% spike in Brent crude triggered by the Middle East crisis has sent retail fuel prices in several major European economies like France, Germany or Netherlands surging back above EUR 2.0/L, approaching historical peaks last seen when Russia's invasion of Ukraine unleashed an unprecedented shock across oil, gas and power markets. For most European economies – structurally exposed as net energy importers – such volatility is far from abstract: A decade of data reveals a stubborn negative correlation between rising pump prices and new car registration volumes in Germany, confirming that fuel costs remain a decisive purchasing factor. Each price shock becomes a powerful advertisement for electrification. Indeed, in Q1 2026, Europe's BEV market share reached a record 19%, up 4pps year-on-year, confirming a structural shift already underway. While Nordic countries lead the EV penetration trend – Norway records almost 100% – some large markets like France and Germany have seen a strong impulse for EVs, with market share increasing by +10% and +6%, respectively, over the past 12 months to 28% and 23% in Q1 2026. Indeed, Germany hit the 1mn milestone in EV sales amid the energy crisis of 2022 and just crossed the 2mn mark in early 2026. Nevertheless, overall demand remains relatively moderate (in absolute terms) and fully electric vehicles account for only around 3% of the total European fleet. Moreover, the momentum remains highly uneven: Nearly half of EU countries still post BEV shares below 10%, while only a small core – around 15% of members – exceeds one-third penetration. The growing popularity of EVs also reflects a broader decline in ICE appeal, a movement amplified by energy price volatility but rooted in deeper market dynamics. Falling retail prices, driven by intensified competition – particularly from Chinese OEMs – and a growing secondary market fueled by off-lease vehicles are easing affordability constraints and stabilizing residual values, unlocking new demand segments. Still, headline gains should be interpreted cautiously: they occur in a stagnant market where total registrations remain roughly 30% below pre-Covid levels, meaning higher shares partly mask weak volumes. Even leading BEV markets such as Norway, Belgium or the Netherlands have recently seen declining sales in absolute terms.

Figure 1 & 2: Oil price benchmark in Europe (Brent) vs retail car fuel price in France / 12-month trailing cumulative registration of new passenger cars in EU-27



Sources: INSEE, ACEA, LSEG Datastream, Allianz Research


Figure 3 & 4: Annual powertrain car market share in EU-27 / BEV market share in top EU markets (Q126 vs. Q125)



Data as of Q1 2026. Sources: ACEA, Allianz Research


Car-running costs have become a genuine financial burden for European households – and the current oil-price spike is exposing once again how structurally vulnerable consumer budgets remain to energy volatility, which could accelerate the shift to electrification. In France and Germany, car-usage spending (fuel, insurance, maintenance – excluding the purchase price) accounted for 7-8% of annual disposable income on average, with the lowest earners hit hardest (11% of annual disposable income). This share visibly surged during the 2022 energy crisis and is trending sharply upward again in early 2026 (Figure 5). Car-usage spending is among the heaviest fixed-cost categories in the family budget and one of the least compressible, given the structural dependency on personal mobility in both urban peripheries and rural areas, notably for commuting to work. In the past, social movements such as the Yellow Vest protests in France erupted due to the rising burden of energy costs on personal transport. The pressure has intensified after the pandemic and the 2022 crisis as car-related services were entangled in a broader cost spiral: Across the EU-27, insurance prices have climbed +37% since end-2021, while maintenance, parts and repairs are up 20-27%, driven by the increasing electronic complexity of modern vehicles and the costlier battery components embedded in new powertrains (Figure 6). The 2021-2022 period has left a permanent scar in the mindset of European households, with inflation being by far the most deterrent downside risk in the main European consumer barometers since then.

Figure 5: Share of annual disposable income allocated to car usage spending* in France, per income distribution quintile



*This includes all costs except car purchase (fuel & lubricant, maintenance, repairing, insurance). 2025 and 2026 figures are estimates calculated on the basis of HICP statistics for fuel, repairing & maintenance and insurance for personal transport. Sources: SDES (Chiffres clés des transports, April 2026), Eurostat, Allianz Research


Figure 6: EU-27 HICP index for personal transport index (2015 = 100), trailing 12-month average growth



Sources: Eurostat, Allianz Research

Powertrain technology arbitrage offers a 4-5% purchasing power boost. At today's fuel price of around EUR 2.0/L across Western Europe, the annual energy bill for an ICE car runs roughly 7–8 times higher than for an equivalent BEV – 7x based on the most efficient models available, rising to 8x when taking into account the current European fleet average, of which EVs still account for less than 3%. That spread is wide by any historical standard but even before the Middle East crisis, when fuel hovered between EUR1.50 and EUR2.00, the multiplier was already sitting at 5–7x – enough to settle the budget debate in favor of EVs regardless of geopolitical tailwinds. Country dispersion adds another layer: Depending on local electricity tariffs, the cost advantage ranges from a modest 5x to a striking 14x, meaning the weakest case for EVs still represents a substantial saving. When translated into purchasing power, that energy differential amounts to a 4–5% income boost for the average Western European driver – not negligible in an inflationary environment. Add the narrowing total cost of ownership (BEVs now show roughly twice fewer recalls than ICE models per ADAC's latest figures, offsetting higher upfront battery costs) with public purchase subsidies closing much of the residual price gap and the picture is clear: the structural economics of the technology shift do not hinge on an oil-price spike. The total cost of ownership equation is quietly, but decisively, tilting toward electric.


Figure 7 & 8: Annual car energy consumption cost comparison between ICE and BEV models in Europe* (based on most advanced energy-efficient model and average consumption of current fleet in 2025)



* based on a benchmark of 10 European countries (Belgium, Denmark, France, Germany, Hungary, Italy, Netherlands, Poland, Spain & Sweden). Calculations based on average power price in April in each country (as of 29 April 2026) and assumption of an average car distance of 12000 km per inhabitant (ACEA). Sources: ACEA, EPEX, Alternative fuels observatory (European Commission), Allianz Research

Infrastructure deficit and grid pressure: The invisible brake on Europe's EV momentum. Beyond technology and price dynamics, the charging infrastructure gap remains one of the most tangible and underestimated structural constraints on EV mass adoption across Europe. Figure 12 illustrates the scale of the challenge starkly: despite a consistent ~20% CAGR in charging stations installed since 2020, reaching 1.1mn points in Q1 2026, the current trajectory falls materially short of the European Commission's 27% CAGR target – projecting a gap of approximately 0.8mn stations by 2030 relative to the EC's 3.5mn objective (Figure 9). Critically, only 16% of installed stations qualify as ultra-fast DC chargers, the very infrastructure that would most directly address consumer range anxiety and charging convenience – arguments that, as discussed, weigh heavier than price premium in many consumer resistance surveys. Dismissing the infrastructure deficit as a consequence of still-insufficient EV penetration misreads the causality: it is the perceived inconvenience of charging that actively deters consumers from making the technology shift in the first place. The distribution problem is equally structural: four countries – the Netherlands, France, Germany, and Belgium – account for roughly 65% of all EU charging stations, with France and Germany alone concentrating 40% of ultra-fast DC points. Rural areas across the broader EU remain critically underserved, amplifying adoption barriers precisely where private vehicle dependency is highest. Compounding these infrastructure pressures is an emerging and formidable competitor for Europe's electrical grid: the rapid expansion of data center capacity driven by AI investment. The IEA estimates EU data center electricity consumption at 70 TWh in 2024, with projections pointing toward 115 TWh by 2030 – a near 65% surge that will compete directly with transport electrification for grid capacity and investment priority. EU electricity demand growth remains comparatively modest, at around 1.1% in 2025 before a modest acceleration to 1.5% in 2026, leaving limited headroom to absorb simultaneously the electrification of mobility and the capex-heavy buildout of AI infrastructure. In this context, grid modernization and targeted infrastructure investment are not secondary enablers of Europe's EV transition – they are its most urgent prerequisite.

Figure 9: Total number of charging stations installed across the EU and projections based on current trend and European Commission target for 2030



Data as of Q1 2026. Sources: Alternative fuels observatory (European Commission), Allianz Research

From crisis to structural shift: Four pillars for a durable EV transition in Europe

There is robust evidence that fossil fuel prices influence the adoption of electric vehicles. In a study of California between 2014 and 2017, Bushnell et al. (2022)¹ show that when consumers decide whether to go electric, the price of gasoline matters four to six times more than the price of electricity, suggesting that what pushes people toward EVs is less the promise of cheap charging than the pain of expensive fuel. The pattern holds in China: Fei et al. (2025)² track monthly sales across 36 cities and find that a 1 CNY/L rise in gasoline prices lifts overall EV sales by 4.67%, with pure battery electric vehicles surging by 9.04%. Most recently, Zhang et al. (2026)³ confirm the relationship in a European context, showing that a 1% increase in gasoline prices raises EV registrations by 0.85% across Denmark, Finland, Norway and Sweden, with budget-friendly models benefiting the most.

Yet the historical record shows that without structural reinforcement, oil-shock-driven shifts in vehicle preferences tend to reverse. During the 2007–08 oil spike, US demand for SUVs and pickup trucks fell sharply, before recovering in subsequent years as fuel prices declined and macroeconomic conditions improved⁴. The pattern repeated after 2014: When crude prices collapsed from USD115 per barrel to below USD30 by early 2016, lower fuel prices coincided with renewed consumer demand for SUVs and pickups, encouraging manufacturers to prioritize those segments⁵. Analysis by Resources for the Future confirms that gasoline prices accounted for a substantial share of the shift toward smaller vehicles between 2003 and 2007, yet that trend reversed once prices fell post-2014⁶.

Crucially, where EV adoption remained resilient through past oil price downturns, strong policy frameworks appear to have played a decisive role. The International Council on Clean Transportation (ICCT) observed that during the 2014 crash, EV sales continued to grow in Norway, the UK and China, markets that combined robust purchase incentives, charging infrastructure deployment and regulatory support, while markets without comparable policy frameworks showed stagnation⁷. The IEA's Global EV Outlook 2025 reinforces this point looking forward: lower oil prices reduce the fuel-cost savings of EVs, and this effect is particularly pronounced in countries with low fuel taxation⁸. The distinction between cyclical and structural displacement is critical: the demand lost during the pandemic was temporary and rebounded once economies reopened, whereas displacement driven by EVs becomes durable only when underpinned by lasting changes in technology, policy and consumer behavior.

This is precisely where carbon pricing enters the picture as a mechanism to make the fossil-fuel cost signal permanent and predictable. While geopolitical shocks produce sharp but temporary price spikes, a well-designed carbon tax raises the floor price of fossil fuels structurally, removing the cyclicity that has historically allowed consumer preferences to revert. Under such a regime, every barrel of oil carries its environmental cost regardless of what OPEC decides or whether a Middle Eastern conflict escalates or de-escalates.

To explore how different carbon pricing trajectories might shape the pace of BEV adoption in Europe, we consider three oil-price scenarios over the 2025–2035 horizon, each reflecting a distinct policy ambition, following the NGFS framework. Under a Baseline scenario – current policies, no additional carbon pricing – oil prices stabilize around USD77–86/barrel through 2035, despite short-term possible price fluctuations related to exogenous factors such as geopolitical instabilities. EU pump prices remain flat at approximately EUR1.73–1.95/L, and BEV adoption grows primarily through the autonomous trend at roughly 1.75 pps per year, driven by technology improvement rather than price signals⁹.

¹ [Energy Prices and Electric Vehicle Adoption | NBER](#)

² [Does high gasoline price spur electric vehicle adoption? Evidence from Chinese cities - ScienceDirect](#)

³ [Electric vehicle adoption and energy prices: Empirical evidence from four Nordic countries - ScienceDirect](#)

⁴ [Causes and Consequences of the Oil Shock of 2007-08 | NBER](#)

⁵ [Automobile manufacturers, electric vehicles and the price of oil - ScienceDirect](#)


⁶ [How Do Gasoline Prices Affect New Vehicle Sales?](#)

⁷ [2014 fuel price turbulence didn't pull the plug on EVs - International Council on Clean Transportation](#)

⁸ [Global EV Outlook 2025 – Analysis - IEA](#)

Under a Nationally Determined Contributions (NDC) scenario, carbon border adjustments and emissions trading begin to bite, pushing oil toward USD100–114/barrel by the early 2030s and EU pump prices to approximately EUR2.10–2.55/L, enough to make fuel expenditure a meaningful push factor, with BEV share reaching approximately 50% in Germany by 2030. Under a Net Zero 2050 scenario, aggressive carbon pricing drives oil above USD190/barrel by 2028–2031 and pump prices to EUR3.50–5.60/L, at which point ICE ownership becomes economically untenable for most households and the mass switch to BEVs accelerates sharply. Figure 10 shows the different pathways of BEV share related to a permanent shift in oil prices.

Figure 10: BEV adoption in EU-26 under different carbon pricing pathways (%)




Sources: NGFS, Allianz Research

Carbon pricing raises the long-run cost of fossil-fuel driving. But making the old technology expensive is not the same as making the new one affordable, particularly for low- and middle-income households, who are the mass market. These are distinct economic problems, because carbon pricing internalizes the emissions externality, while purchase subsidies address the separate barriers – such as upfront cost, learning-by-doing, network effects around charging infrastructure – that hold back adoption of a technology most consumers still perceive as novel and risky⁹. The two instruments are often best deployed as complements rather than alternatives. European policymakers appear to have absorbed this logic, as the number of active national EV incentive policies across EU-27 countries has grown from just two in 2010 to 74 by 2024, with a sharp acceleration after 2019 as governments layered purchase grants, tax exemptions, charging infrastructure support and fleet incentives on top of one another (Figure 11).

⁹ This estimate is derived from a linear probability model of BEV sales share across EU-26 countries over 2010–2024: $ev_sales_share = \alpha(u) + \beta_1 \cdot pump_price + \beta_2 \cdot \ln_income + \beta_3 \cdot (\ln_income \times price_range) + \beta_4 \cdot year_trend$. The autonomous trend coefficient ($\beta_4 = 0.0175$, $p < 0.001$) captures the annual increase in BEV market share attributable to technology diffusion and shifting consumer preferences, independent of price and income dynamics. The pump price coefficient ($\beta_1 = 0.161$, $p < 0.001$) implies that a EUR0.10/L increase in fuel prices lifts BEV share by approximately 1.6 percentage points, while a EUR0.50/L spike — comparable to the 2022 energy crisis — adds roughly 8.1 percentage points. Income enters negatively at the baseline ($\beta_2 = -0.292$, $p = 0.027$) but is moderated by an interaction with within-year price volatility ($\beta_3 = 0.0075$, $p < 0.001$), indicating that higher-income consumers respond more strongly to fuel price uncertainty when choosing EVs.

¹⁰ [Carbon pricing in climate policy: seven reasons, complementary instruments, and political economy considerations - Baranzini - 2017 - WIREs Climate Change - Wiley Online Library](#)


Figure 11: Number of EV subsidies schemes in the EU



Sources: IEA, Allianz Research

How far can subsidies alone take Europe? To answer this, we model BEV adoption across EU-27 countries under three policy scenarios, calibrating the subsidy effect from Martins et al. (2024)¹¹. Their estimates show that a modest tax exemption worth around EUR1,000 adds 3.2% annual growth to BEV market share, a moderate grant below EUR5,000 adds 11.6% and a generous grant above EUR5,000 adds 26.4%. To analyze the BEV pathways as a response to subsidies, we built different scenarios (Figure 12). In the most optimistic scenario, every EU country raises grants to at least EUR5,000 through 2031, then gradually phasing them out by 2035 as BEV-ICE cost parity approaches, i.e., market share reaches 70% by 2030 and 93% by 2035. If instead governments simply maintain current programs through 2028 and then let them wind down, BEV share reaches only 59% by 2030. If political backlash or fiscal pressure leads to abrupt subsidy removal by 2029, adoption stalls at just 48% by 2030, slowly recovering to 76% by 2035 as the market falls back on cost economics and regulation alone. Therefore, even the most generous pathway falls 10pps short of the 80% BEV share by 2030 that would be consistent with meeting European decarbonization targets.

Figure 12: BEV adoption in EU-26 under different subsidies pathways (%)




Sources: Allianz Research

¹¹ [Assessing the effectiveness of financial incentives on electric vehicle adoption in Europe: Multi-period difference-in-difference approach - ScienceDirect](#)

The evidence shows that subsidies accelerate adoption, but fiscal constraints might make it unrealistic for several economies to sustain large-scale purchase incentives through the full duration of the transition. Muehlegger and Rapson (2022)¹² show that low- and middle-income households are highly price-sensitive when it comes to EVs: a 10% price reduction can increase sales in this segment by roughly 21%. But no government can subsidize its way to 80% market share indefinitely. The fiscal cost escalates as adoption scales, and as Archsmith et al. (2022)¹³ argue, what ultimately sustains mass adoption is not the subsidy itself but what they call "intrinsic demand", the point at which vehicle quality, model availability, charging convenience and total cost of ownership make the EV the obvious choice without a grant, which would depend on speeding up the technological transition.

The technological case for EVs rests, above all, on one number: the price of a lithium-ion battery pack. Figure 13 shows that price has fallen from USD1,474 per kilowatt-hour in 2010 to USD108 in 2024, a 93% decline in 14 years. This trajectory follows a remarkably stable learning rate of approximately 18%: every doubling of cumulative manufactured capacity reduces costs by roughly a fifth. What makes this learning rate so consequential is that global battery demand is now entering an exponential phase. Passenger EVs already dominate, but demand from electric buses, commercial vehicles, two- and three-wheelers and stationary storage is compounding rapidly. It is projected that total demand will rise from roughly 1,000 GWh today to over 5,000 GWh by the early 2030s. Each of these segments feeds the same manufacturing base, driving cumulative production volumes higher and pulling costs down further regardless of what happens in any single market. If the 18% learning rate holds, pack prices are on track to reach approximately USD60–70/kWh by 2030 and could fall below USD55 by the mid-2030s. At those levels, the upfront purchase price of a BEV falls below that of an equivalent ICE vehicle in most segments without any subsidy at all. This is the point at which "intrinsic demand" takes over: the EV becomes the default choice not because governments pay people to buy it, but because it is simply the cheaper, better car.

Figure 13: Lithium-ion battery pack prices and demand



Sources: New Energy Outlook BNEF 2025, Allianz Research

¹² [Subsidizing low- and middle-income adoption of electric vehicles: Quasi-experimental evidence from California - ScienceDirect](#)

¹³ [Future Paths of Electric Vehicle Adoption in the United States: Predictable Determinants, Obstacles and Opportunities | NBER](#)


There is, however, a final condition that determines whether all of this – the carbon pricing, the subsidies, the falling battery costs – actually delivers on its environmental promise. An EV is only as clean as the electricity it runs on. Across the EU, grid carbon intensity has fallen from 0.45 tCO₂/MWh in 2005 to 0.22 in 2025, driven by a low carbon energy share that has grown from 47% to 70% of total generation (renewables grew from 16% to 50%). Under the EU NDC scenario, that trajectory continues: carbon intensity drops to 0.18 by 2030, 0.13 by 2035 and 0.07 by mid-century as low-carbon sources rise from 70% to 86% of generation. Each reduction compounds the emissions benefit of every EV already on the road. Moving from today's low-carbon electricity share of 70% to roughly 80% by 2035 would, on its own, reduce the per-vehicle carbon footprint of an EV by over 40%.

Germany illustrates the benefits of the grid transition. With a low-carbon share of 57% in 2025, an EV driven 14,000 km per year emits approximately 0.98 tCO₂, already 46.6% less than a new ICE car. But the gains from grid greening are steep: as Germany's low-carbon share rises to 64% by 2030 and 76% by 2035 under NDC commitments, the annual CO₂ saving per vehicle climbs from 0.85 tCO₂ today to 1.05 by 2030 and 1.39 by 2035, a 75.6% reduction relative to ICE. Therefore, grid decarbonization alone, without any change in driving behavior or vehicle efficiency, cuts the German EV's carbon footprint by 54% over a single decade. This creates a powerful feedback loop: the more EVs on the road, the stronger the case for grid investment; the greener the grid becomes, the larger the climate return on every vehicle already switched. Conversely, stalling grid decarbonization erodes the credibility of the entire transition. An 80% BEV fleet running on a 'brown' grid delivers far less than a 50% BEV fleet running on a clean one. Grid decarbonization is a critical complement of green transport in Europe.

Downside scenario: what if Middle East effects extend beyond Q2?

#1 - What if power price spikes? European EV demand proved unexpectedly resilient during the European energy crisis of 2022, with registrations rising ~28% while the total car market declined. This resilience reflected two temporary buffers: aggressive public intervention to cap electricity prices and still-generous purchase subsidies. Today, both buffers are weaker, notably subsidy schemes that broadly phased out over 2022-2023 in large markets like France and Germany. Fiscal space is more constrained and subsidy schemes are being scaled back or better targeted. Meanwhile, gas storage levels and geopolitical uncertainty—particularly around the Strait of Hormuz—reintroduce upside risks to power prices. In such a scenario, EV total cost of ownership (TCO) advantages could narrow, especially for households relying on spot-priced electricity. That said, structural mitigants are stronger than in 2022. Europe's power mix has shifted toward renewables, nuclear availability has improved (notably in France), and occasional negative wholesale prices in 2025 highlight underlying capacity surplus. In addition, diversified gas supply from Norway, the U.S. and LNG markets reduces extreme tail risks. Net-net, EV demand is no longer insulated but remains less exposed than often assumed—suggesting downside risk is real but likely contained rather than systemic.


Figure 14 & 15: Natural gas price and storage rate in Europe / Annual passenger car registration trend & BEV market share in EU-27 since 2021



Market data as of 1st May 2026. Sources: LSEG Datastream, ACEA, Allianz Research

#2 - What if consumer sentiment deteriorates? EV demand has historically shown relative resilience during macro slowdowns, largely because early adopters skew toward higher-income households with lower marginal sensitivity to energy and financing costs. This partially explains why EV penetration continued to rise even as broader European car demand remained roughly 30% below pre-COVID levels in recent years. However, this resilience has limits. EVs remain a discretionary, high-ticket purchase and global uncertainty—geopolitical tensions, energy volatility or tighter monetary policy—typically triggers postponement behavior. Rising interest rates directly affect monthly leasing costs, a key driver of EV affordability in Europe, while weaker confidence can delay fleet renewals by corporates. Policy can smooth the cycle but is becoming less supportive. With subsidy schemes being reduced or retargeted, the burden increasingly shifts to private demand elasticity. Even if new instruments (e.g. social leasing in France) provide targeted support, they are unlikely to fully offset a broad-based confidence shock. The key risk is not a collapse in penetration, but a slower-than-expected ramp-up, as consumers defer adoption despite improving technology fundamentals. In this context, EV demand should remain relatively more resilient than ICE demand in a downturn—but not immune, so a potential extension of the conflict in Iran could downsize EV sales growth to 10-15% this year (vs. almost -35% in Q1 2026) while total registration would contract again.

Figure 16: Profile ID of EV purchasers in Europe (2025)



Sources: EU drivers' view on electric cars : European Alternative Fuels Observatory Consumer Monitor 2025, ACEA, Allianz

#3 - What if automakers are lacking chips? EVs are structurally more exposed to semiconductor supply disruptions than ICE vehicles. Semiconductor content is roughly 2x higher in EVs—and up to 3x for advanced models integrating L2+/L3 functions—driven by power electronics, battery management systems and centralized computing architectures. This amplifies sensitivity to any disruption across the semiconductor value chain. A key vulnerability lies in upstream inputs such as helium, essential for advanced lithography. A supply shock linked to Middle East tensions could disrupt production at major foundries like TSMC or Samsung Electronics, with immediate spillovers to automotive chip supply. This risk compounds an already tight environment, as AI-driven demand continues to absorb leading-edge capacity. The consequences would be twofold: higher input costs and renewed production bottlenecks. Unlike 2021–2022, OEMs have improved inventory management and sourcing diversification, but buffers remain limited for advanced nodes and power semiconductors. In a severe scenario—particularly if disruptions coincide with tensions around the Strait of Hormuz—EV supply could tighten more than ICE, widening price differentials at the retail level. Based on a price sensibility matrix, we estimate that on average BEV retail prices would be twice as high as ICE models if chip price double or triple whatever state incentives allocate to boost sales (assuming a maximal subsidy rate at EUR 10,000). Bottom line: semiconductor risk is asymmetric. Any supply shock would disproportionately weigh on EV output and margins, reinforcing short-term downside risks to adoption.


Table 1: BEV vs. ICE retail price spread matrix based on chip price volatility and state EV subsidy policy

BEV vs ICE - average retail price spread (End-2025 = ~30%)		Chip price - Multiplier ratio				
		2.0	2.5	3.0	4.0	5.0
EV subsidies on new vehicle purchase (in EUR)	1000	114%	118%	122%	129%	136%
	2000	109%	114%	118%	125%	132%
	3000	105%	109%	113%	121%	128%
	4000	100%	105%	109%	117%	124%
	5000	96%	101%	105%	113%	121%
	6000	92%	96%	101%	109%	117%
	7000	87%	92%	96%	105%	113%
	8000	83%	87%	92%	101%	109%
	9000	78%	83%	88%	97%	105%
	10000	74%	79%	83%	92%	101%

Sources: Allianz Research

#4 - What if China slows down? An extended Middle East conflict could push energy costs higher in Asia, weighing on China's already fragile domestic demand. Combined with intense price competition and margin compression among OEMs, this would likely accelerate the international expansion of Chinese manufacturers. Europe stands out as the primary target market, especially as access to the U.S. remains restricted. Chinese OEMs have already increased their European market share to ~7% in 2025, leveraging strong cost competitiveness and increasingly high product quality. In a context of domestic overcapacity, exports become a key adjustment variable. Trade policy remains the swing factor. While the European Commission has initiated anti-subsidy investigations, the current regulatory framework still allows significant market access. Notably, local content requirements remain partial, with battery sourcing rules tightening only progressively toward the end of the decade. If no additional tariff barriers emerge, a China-led supply push could compress EV price premiums versus ICE, accelerating mass-market adoption in Europe. Paradoxically, a negative external shock (China slowdown) could therefore support EV penetration in Europe through intensified competition and lower prices—at the cost of increased pressure on domestic OEM margins.

Figure 17 & 18 : Chinese automakers' passenger car sales breakdown (domestic & foreign market) / Net profit margin rate & estimated market share in EU-27 (all powertrains)



Sources: CAAM, Eurostat, Allianz Research

These assessments are, as always, subject to the disclaimer provided below.

FORWARD-LOOKING STATEMENTS

The statements contained herein may include prospects, statements of future expectations and other forward-looking statements that are based on management's current views and assumptions and involve known and unknown risks and uncertainties. Actual results, performance or events may differ materially from those expressed

or implied in such forward-looking statements.

Such deviations may arise due to, without limitation, (i) changes of the general economic conditions and competitive situation, particularly in the Allianz Group's core business and core markets, (ii) performance of financial markets (particularly market volatility, liquidity and credit events), (iii) frequency and severity of insured loss events, including from natural catastrophes, and the development of loss expenses, (iv) mortality and morbidity levels and trends,

(v) persistency levels, (vi) particularly in the banking business, the extent of credit defaults, (vii) interest rate levels, (viii) currency exchange rates including the EUR/USD exchange rate, (ix) changes in laws and regulations, including tax regulations, (x) the impact of acquisitions, including related integration issues, and reorganization measures,

and (xi) general competitive factors, in each case on a local, regional, national and/or global basis. Many of these factors may be more likely to occur, or more pronounced, as a result of terrorist activities and their consequences.

NO DUTY TO UPDATE

The company assumes no obligation to update any information or forward-looking statement contained herein, save for any information required to be disclosed by law.

Allianz Trade is the trademark used to designate a range of services provided by Euler Hermes.

Allianz Research | 30 April 2026

Happy Labor Day? How geopolitics, immigration and AI will reshape work

In Summary

Ludovic Subran
Chief Investment Officer & Chief
Economist
ludovic.subran@allianz.com

Maxime Darmet
Senior Economist US-UK-France
maxime.darmet@allianz-trade.com

Bjoern Griesbach
Head of Macroeconomic and Capital
Markets Research
bjoern.griesbach@allianz.com

Jasmin Gröschl
Senior Economist for Europe
jasmin.groeschl@allianz.com

Simon Krause
ESG and AI Economist
simon.krause@allianz.com

Maddalena Martini
Senior Economist Southern Europe &
Benelux
maddalena.martini@allianz.com

Katharina Utermöhl, CFA
Head of Thematic & Policy Research
katharina.uterhoehl@allianz.com

Thomas Loubeyres
Research Assistant
thomas.loubeyres@allianz-trade.com


Lucie Trollé
Research Assistant
lucie.trolle@allianz.com

- **US and European labor markets appear to be in good shape, with headline unemployment rates near historic lows. But three strong undercurrents (immigration policy, energy-price shock and AI) are churning beneath the calm surface.** Demand for workers had already begun to soften before the shock of the Middle East crisis: vacancy rates were easing and hiring rates cooling from uncertainty and the trade war. The shift in immigration policy in the US, the UK and Germany is adding to the mix: Fading immigration inflows have quickly translated into lower employment numbers. In the US, immigration has gone from contributing over half of job creation in 2024 to near-absent in 2025 – weighing on potential growth and risking sectoral labor shortages, while offering a short-term offset against rising unemployment.
- **The Iran shock will land unevenly on labor markets depending on countries' energy exposure and the duration of the conflict. Still, we see only a limited increase in unemployment rates in the range of 0.1-0.3pp.** If the crisis is resolved by end-May, unemployment rates would rise only in the most exposed European economies, and by +0.1pp at most, with 102,000 jobs lost in the Eurozone and half that in the US. Unlike in 2022, when labor hoarding led to falling unemployment, leaner buffers today could see any shock frontloaded into unemployment. In case of a prolonged closure of the Strait of Hormuz, the energy shock would hit Europe harder than the US, though higher labor protection standards provide some cushion: an estimated 225,000 jobs would be lost (0.13% of employment) in the Eurozone versus 126,000 (0.08%) in the US, with higher losses in Germany, Poland, Italy, France and Spain due to high energy exposure.
- **AI's labor-market effects are emerging, pointing to a K-shaped pattern, with youth and mid-level white-collar workers most at risk.** Early evidence shows pressure on younger and less experienced white-collar workers in routine cognitive tasks, while gains accrue to higher-skilled, AI-complementary roles. Since late 2022, higher AI adoption has been associated with larger increases in youth unemployment, with AI exposure explaining about 40% of cross-country variation (excluding high-unemployment economies). AI may therefore appear first not as job loss but as fewer entry points, weaker wage growth, and sharper polarization, with reallocation driven mainly by shifts in job composition rather than wage pressures in labor-intensive activities predicted by Baumol.
- **In the medium term, the AI labor-market impact will be substantial, uneven across countries and unprecedented in the scale of workforce reorganization required.** Over the next 1-3 years, AI is expected to affect 23.3% of jobs across major economies, with reorganization (10.4% of jobs) dominating augmentation (5.3%) and outright displacement (7.6%). The share of jobs affected ranges from 9.2% in Italy to 28.7% in the US, with the UK (17.7%), Germany (16.2%), France (14.7%) and Spain (12.4%) in between. That is equivalent to 52.5mn jobs in the US and 21.8mn across the major European economies. Our analysis does not account for potential AI-related job growth, which we expect to at least partially offset adverse employment effects. However, job displacement is likely to outpace job creation in the medium term, as firms adjust faster than workers, creating a temporary gap. Ultimately, whether – and how quickly – AI leads to job losses, reorganization, or new job creation will depend less on technology than on policy choices. AI-proofing policy frameworks will be critical: labor-market policies, including re- and upskilling, active labor-market programs, and social protection, will shape worker transitions. Taxation (including the relative treatment of labor and AI capital), firm incentives, and competition policy will determine whether AI is deployed to augment or replace labor and how broadly productivity gains are shared.

The buffer has thinned: labor markets and the new energy shock

Labor markets on both sides of the Atlantic have so far remained resilient despite successive shocks. In the Eurozone, the unemployment rate remained close to historical lows at 6.2% in March 2026, reflecting the job-rich post-Covid recovery and the strong employment response relative to modest growth. In the US, labor-market conditions have been weaker, with net job creation grinding to close to zero since early 2025, after dynamic outturns in 2024. However, stalling job creation mostly reflects sharply weakening labor supply amid very tight immigration policy, though plunging labor demand for federal employees has contributed, too. As a result of lower immigration inflows, the US unemployment rate has not crept up by much: from 4.1% in January 2025 to only 4.3% in March 2026.

Figure 1: Unemployment rate (%) including Allianz Research baseline forecasts.




Sources: LSEG Datastream, Allianz Research

However, beneath the headline resilience, Eurozone labor market buffers have been progressively eroding. Even before the latest energy shock, labor demand had already cooled visibly (Figure 2): vacancy rates declined from 3.4% in Q4 2022 to 2.3% in Q4 2025; reported labor shortages eased across sectors, with firms seeing labor as a less limiting factor to production, and adjustment increasingly took place through hours worked rather than new hiring. Compared to the previous energy shock in 2022, firms no longer face acute rehiring constraints, reducing incentives to hoard labor. As a result, the Eurozone enters the current shock with a labor market that is less tight and with less buffers, despite still-low unemployment. Against this backdrop, the cooling of labor-market conditions should reduce the risk of second-round effects from higher commodity prices, dampening underlying wage and price pressures.

The US labor market also looks fragile, though the US economy is more insulated from surging energy prices. As a net energy exporter, the US has been far less exposed to energy-import-driven cost shocks since 2022, limiting the need for margin or labor adjustment and allowing firms to preserve buffers to a much greater extent. US corporates retain large cash buffers, structurally higher profit margins and a historically low interest-coverage ratio. Even so, labor demand looks increasingly soft. The job-hiring rate recently dipped to its lowest level since the pandemic (at 3.1%), signaling increasing hiring caution in an environment of persistent uncertainty, while the rollout of AI may have also deterred new hiring. Besides, the bulk of job creation has been concentrated in non-cyclical sectors, i.e. education & healthcare. Private jobs excluding education & healthcare have dipped by a cumulative -420k since January 2025.


Figure 2: Beveridge curve (%)



Sources: LSEG Datastream, Allianz Research

Job growth was heavily supported by immigration over the past few years, but this dynamic is fading quickly, particularly in the US. Lower labor supply offers a short-term reprieve against rising unemployment rates but will weigh on potential growth. Figure 3 depicts the percentage annual change in employment (net job creation) since 2023 in the largest European countries and the US, with a breakdown of the contribution from immigration vs the resident. In Germany and the UK, employment of residents has decreased and net job creation has been concentrated among the new immigrants. Spain has seen strong employment growth of residents but new immigrants have contributed almost as much. In total, elevated immigration inflows have contributed to positive job creation in many European countries, alleviating labor shortages. However, job creation has stalled significantly in 2025, particularly in the UK, France and Italy. Firms have prioritized efforts to regain productivity rather than hiring to produce output. Consistent, harmonized working immigration inflows data are not available in European countries, but national data suggest a fading contribution in Germany, France and the UK. The US has seen a similar pattern of rapidly weakening job-creation momentum since 2025, driven almost entirely by reduced immigration inflows. Softer immigration inflows reduce labor supply and will thus weigh on potential growth. However, in the short term, scarcer labor is probably acting as a near-term brake on unemployment, even as it drags on growth potential.

Figure 3: Employment change (% employed)




Sources: OECD, Census Bureau, Brookings, Allianz Research. Note: Permanent migration is defined by legal status, not duration of stay by the OECD. It includes migrants with a pathway to remain indefinitely such as refugees, family migrants, and some labor migrants, while those without such rights (e.g. students or seasonal workers) are not considered

Initial conditions are important for understanding how the renewed rise in energy prices propagates through labor markets. Although energy prices have fallen since their 2022 highs, industrial electricity prices in Europe remain significantly higher than pre-crisis levels and those in the US (Figure 4), even after adjusting for recent

wholesale price stabilization. In contrast, US industrial electricity prices have remained stable over time due to abundant gas supply. Consequently, energy costs pose a persistent structural challenge for European firms rather than a temporary cyclical shock.

Figure 4: Industrial electricity prices, EUR/MWh including all taxes



Sources: EIA, Eurostat, Allianz Research. Notes: EU price range for consumption bands above 2,000 MWh.

There is far less margin to absorb renewed cost pressures than in 2022. Total energy costs are high in manufacturing, utilities and transport due to the scale of these sectors. The most energy-intensive industries, such as chemicals, steel and metals, paper, glass and ceramics, and cement, are concentrated in large industrial economies, particularly Germany, followed by France, Italy and Poland. Since margins have not fully recovered from the last energy crisis, firms have limited capacity to absorb new shocks through profitability alone.

As a result, the labor-market adjustment mechanism is shifting. While the 2022 energy shock was largely absorbed through margins, working-time reductions and labor hoarding in an overheated labor market, firms today face stronger incentives to adjust earlier if high energy prices persist. This time it might be different due to labor markets in a looser state where less labor hoarding will be seen. The immediate risk is not a sharp unemployment spike but a gradual erosion of labor-market resilience, in particular in countries that are suffering from structural economic weakness such as Germany.

Consistent with this gradual erosion, baseline projections already point to non-negligible job losses driven by broader macroeconomic headwinds. For the Iran war, we consider two scenarios (Figure 5). In our baseline scenario, traffic through the Strait of Hormuz gradually normalises by the end of May. This implies employment losses of around 102,000 in the Eurozone (with German, France and Spain already accounting for 68% of the total), and 51,000 in the US by 2026 – so in both cases the impact on the unemployment rates stemming solely from the Iran war under the baseline scenario remains very small (<0.1pp). These baseline dynamics provide the reference against which additional energy-related labor-market effects should be assessed.


Figure 5: Middle East escalation scenario

Description	Feb 28th	Baseline (45% probability): Prolonged conflict with less than three months of global oil and gas disruption ***		Downside (35% probability): Long sustained escalation with prolonged disruption of global oil and gas supply ***	
		2026		2026	
		A deal allowing for a progressive normalization of oil and gas trade in May. In between partial re-routing, additional production (US, Russia) and strategic reserves being tapped covers >50% of Hormuz gap in oil and gas; no material destruction of oil/gas industry.		Hormuz closed for more than three months. Oil, gas and other commodity supply shortages (fertilizers, helium) transmits to confidence & equity shock exacerbating downturn.	
Oil Brent (USD/bbl), Q2 average / eoy 2026	73\$	90\$ / 78\$ (\$140)*		130\$ / 85\$ (\$180)*	
Gas prices, annual avg.	32€	49€ (150€)*		65€ (200€)*	
GDP	EZ	1.3%**	0.8%	0.2%	
	US	2.6%**	2.1%	1.5%	
Inflation	EZ	1.9%**	3.0% (3.4%)*	3.9% (4.9%)*	
	US	2.5%**	3.2% (3.7%)*	4.0% (4.6%)*	
EURUSD		1.18	1.15 (1.11)*	1.12 (1.05)*	
Monetary Policy	ECB	2.0	2.25	2.75	
	Fed	3.75	3.75	4.25	
10y	DE	2.6	2.8 (3.3)*	3.1 (3.7)*	
	US	4.0	4.5 (4.8)*	5.0 (5.7)*	
Equity market	EZ	6	5 (-10)*	-25 (-30)*	
	US	1	6 (-10)*	-20 (-25)*	

Sources: Allianz Research. (*) max daily draw-down in brackets, (**) macro forecast for 2026 before the start of the Iran war, (***) probabilities add up to 80% with remaining upside and downside tail-risks (≤5%) omitted

Country-level exposure further shapes how energy shocks translate into labor-market outcomes. Energy intensity – measured as energy use relative to gross value added (GVA) – varies widely across Europe and is not, on its own, a reliable indicator of employment risk. Some highly energy-intensive economies, such as Hungary, the Czech Republic and Romania, account for a relatively small share of total Eurozone employment, despite requiring more energy per unit of output. By contrast, larger economies such as Germany, France, Italy and Spain combine more moderate energy intensity with much larger employment shares, making them more exposed to labor-market adjustment driven by energy prices (Figure 6). Poland stands out in this respect, combining a sizable employment share with high energy costs relative to GVA, which increases its vulnerability. This composition explains why employment risks are likely to materialize unevenly across countries, rather than through a synchronized rise in Eurozone unemployment.

Figure 6: Energy exposure of employment (index) and energy intensity (MWh per current 1,000 EUR), by country




Sources: Eurostat, SBS, Allianz Research. Notes: Energy exposure are energy costs relative to gross value added (GVA) relative to EU employment shares by country. Energy intensity is calculated as energy use divided by gross value added; it gives an indication how many MWh are needed to generate 1,000 EUR of GVA.

In a downside scenario involving a prolonged closure of the Strait of Hormuz (Figure 6), employment effects would intensify but remain gradual. Rising fossil-fuel prices would weigh on value added through energy

dependencies and supply bottlenecks in intermediate goods. Eurozone unemployment rates would rise by around +0.14pp on average, with the largest increases in Spain and Belgium (around +0.20pp). Eastern European economies – Czechia, Poland and Romania – would experience the steepest output losses, with unemployment rising by +0.12pp, +0.17pp and +0.27pp respectively. Germany and Italy, despite accounting for a large share of Eurozone employment, would see more contained increases (+0.12pp and +0.09pp), reflecting their comparatively higher energy efficiency. In absolute terms, the Eurozone would lose an estimated 225,000 jobs (0.13% of employment) versus 126,000 (0.08%) in the US. This is the subset of job losses attributed to the current energy shock in a downside scenario, compared to job losses due further economic factors, such as uncertainty, structural country-specific challenges and the ongoing trade war, to mention only a few, in our baseline scenario quantified in our latest insolvency report¹. For comparison, for the US, about 11.9% of the total expected job losses of 428,000 in 2026 may be attributable to the energy-priceshock in the baseline scenario, while for Germany it would be 12.4% of the 209,000 job losses and for France 7.1% of the 283,000.

The relatively narrow dispersion in unemployment outcomes, however, masks a wider divergence in how the shock actually lands on workers. More energy-intensive economies – particularly in Eastern Europe – will face sharper inflation and steeper margin compression, but much of their labor-market adjustment will occur through real wage erosion rather than outright job losses. Unemployment figures alone therefore understate the true dispersion of worker hardship: where prices rise fastest and margins are thinnest, it is purchasing power, not headcount, that takes the first hit.

Figure 7: Change in GDP and unemployment rate in a downside Iran war scenario compared to the baseline, in pp



Sources: Oxford Economics, Allianz Research


AI is reshaping adjustment pathways, not driving the shock

AI is altering how firms adjust to the energy shock, not replacing energy costs as its source. With AI already embedded in production processes, firms have gained greater scope for internal substitution, particularly in routine cognitive tasks and entry-level and mid-level white-collar roles. In a context of elevated energy costs and softer demand, this lowers the threshold at which firms move from margin and working-time adjustment to labor reallocation, primarily through hiring restraint and weaker entry-level absorption rather than outright job destruction. Energy prices determine the need to adjust; AI shapes the speed, intensity and distribution of that adjustment.

Accordingly, more AI-intensive sectors have displayed muted or zero employment growth in the Eurozone since late 2022, consistent with adjustment taking place through subdued hiring and task reorganization and slower entry-level inflow rather than outright job destruction (Figure 8). This reinforces the view that AI dampens measured employment dynamics even as underlying restructuring progresses.

¹ See also: [Allianz Trade Insolvency report 2026](#)

Figure 8: Eurozone employment growth by sector and AI usage




Sources: Eurostat, Allianz Research.

From a Baumol perspective, uneven productivity growth tends to generate upward pressure on relative wages in less automatable, labor-intensive activities. In the current environment, however, this channel is muted: with AI dampening aggregate labor-demand growth and firms adjusting primarily through hiring restraint and task reorganization, rebalancing is taking place mainly through employment composition rather than prices. Wage adjustment is therefore asymmetric, with compression in AI-exposed white-collar and entry-level roles, while less-AI-intensive sectors absorb labor without broad-based wage acceleration.

AI has not yet raised aggregate unemployment, but it is already reshaping labor-market adjustment in a k-shaped manner. Since late 2022, countries with higher AI diffusion have experienced larger increases in youth unemployment, with AI exposure explaining around 40% of the cross-country variation (excluding economies with structurally high youth unemployment), even as overall unemployment rates remain contained. At the same time, countries with comparable levels of AI uptake display markedly different labor-market outcomes, underscoring the decisive role of institutions, labor-market structures and adjustment channels. This divergence suggests that AI diffusion does not mechanically translate into net job losses, but neither is it benign: it shapes who adjusts, how quickly and at which margin, depending on how easily firms can substitute tasks, how flexibly workers can be reallocated and how effective retraining and activation policies are. Where adjustment mechanisms are rigid and transition support is weak, the risk is not an immediate rise in aggregate unemployment but persistent scarring at the margin, concentrated among younger workers and in routine-intensive entry-level occupations. The policy challenge therefore lies less in protecting aggregate employment than in shortening transitions and reducing polarization, through portable skills, flexible wage-setting and active labor-market policies that allow economies to capture the productivity upside of AI without incurring lasting distributional costs.

Figure 9: Global AI adoption and youth unemployment rate



Sources: Microsoft Global AI Adoption Index, OECD, Allianz Research. *The regression line and R^2 exclude Spain and Italy. Both economies have structurally high youth-unemployment levels since the GFC and the sovereign-debt crisis. Bubble size corresponds to total unemployment rate.

AI will have a more substantial impact on the labor market over the medium term

As AI moves from initial deployment to broader integration in business processes, the focus shifts to how it shapes labor markets in the near- to medium-term. While it remains difficult to predict how the technology will evolve and diffuse, the next one to three years are likely to be critical in determining the extent to which current capabilities translate into tangible changes in jobs and tasks. Against this backdrop, we assess the potential impact of AI on the current workforce across the Eurozone's four largest economies – Germany, France, Italy and Spain – as well as the UK and the US.

AI job exposure varies significantly across sectors. As a starting point for the analysis, we identify AI exposure across sectors, measuring the share of tasks that AI can perform significantly faster or more efficiently. Based on the recent analysis of Falck and Röhe (2026)², it ranges from around 14% in agriculture and 22–25% in construction and manufacturing to over 40% in finance and real estate and nearly 48% in information and communication. These differences in AI exposure are driven by task composition and are broadly stable across countries. Sectors dominated by cognitive work are substantially more exposed, while manual and physical tasks remain less affected.

Table 1: Heatmap of AI exposure across sectors (% jobs exposed) and countries' sectoral employment composition (% total employment)

	AI Exposure	Employment Shares by Sector					
	by Sector	USA	UK	France	Germany	Italy	Spain
Wholesale and Retail Trade	34,5%	11,8%	13,7%	12,6%	12,5%	13,9%	14,5%
Manufacturing	24,7%	6,9%	7,2%	10,8%	18,2%	18,2%	12,0%
Administrative and technical activities	34,0%	12,3%	17,5%	11,2%	10,1%	10,9%	10,9%
Public Administration and Defence	38,3%	25,5%	4,8%	8,3%	7,2%	5,0%	6,6%
Education, Health and Social Work	30,5%	15,2%	8,6%	7,7%	6,9%	7,0%	7,1%
Accommodation and Recreation	29,4%	9,3%	10,2%	6,4%	5,0%	7,9%	10,6%
Construction	22,3%	4,6%	5,1%	6,6%	6,3%	6,7%	6,8%
Transportation and Storage	28,7%	3,6%	5,2%	5,0%	4,8%	4,9%	5,6%
Financial Activities	41,3%	5,0%	5,1%	4,7%	4,0%	3,2%	2,8%
Information and Communication	47,8%	1,5%	4,3%	3,6%	4,1%	3,2%	3,8%
Other Services	27,0%	3,3%	1,9%	2,5%	3,1%	3,1%	2,4%
Agriculture	13,5%	0,4%	0,6%	2,3%	1,1%	3,4%	3,5%
Electricity, Gas, Water supply	32,0%	0,3%	1,1%	1,7%	1,6%	1,7%	1,2%
Mining	25,5%	0,3%	0,1%	0,1%	0,2%	0,1%	0,2%
Total Economy AI Exposure		33,1%	32,1%	30,5%	31,0%	29,5%	30,0%

Sources: Falck and Röhe (2026), BLS, ONS, Eurostat, Allianz Research.

Cross-country labor market exposure to AI is driven by economies' sectoral employment composition (Table 1).


While AI affects sectors similarly across countries, economies differ in how their workforce is distributed across those sectors. Translating sectoral exposure into country-level outcomes therefore depends on national employment structures, making economic composition the key driver of aggregate exposure. For example, high-exposure sectors such as ICT (48%), finance and real estate (41%), public administration and defense (26%), and health and education (15%) account for a larger share of employment in the US and partly the UK than in continental Europe. By contrast, Italy and Spain are more concentrated in lower-exposure sectors such as manufacturing (18% and 12%) and agriculture (3%), reducing aggregate exposure. Germany occupies a middle position among all countries, combining a sizable manufacturing base (18%) with more moderate service exposure. As a result, even with identical sectoral exposure, total AI impact differs substantially across countries, with service-oriented economies more exposed than industrial or agriculture-based ones.

Among AI-exposed jobs, we apply a job-transition framework that distinguishes between three adjustment channels: augmentation, reorganization and automation. These categories capture how firms are likely to

² For the estimates of the AI exposure across sectors, see: Falck, E., and O. Röhe (2026), "From potential to practice: AI exposure and adoption in European industries", SUERF Policy Brief 1358.

restructure tasks within the existing workforce rather than simply whether jobs disappear. Based on the analysis by Martin Richmond (2026), 22% of AI-exposed jobs are expected to be augmented, 44% reorganized and 33% automated.³ Augmented jobs expand as productivity gains stimulate labor demand, reorganized jobs retain their core function but undergo task-level changes while for automated jobs humans are fully replaced by AI.

Figure 10: Workforce readiness for AI (left) and AI adoption among workers (right), in %




Sources: IMF, Eurostat, Bick et al. (2026), Allianz Research.

How these job-level effects translate into labor-market outcomes depends on country-specific conditions, which determine how quickly and to what extent AI exposure materializes over the next 1–3 years. We capture this through four scaling factors that amplify or cushion the impact. In the short term, AI adoption and workforce readiness for AI determine how quickly theoretical exposure causes actual impact (Figure 10). Countries such as the US and the UK lead both in current AI adoption and workforce readiness for AI, while Italy ranks the lowest among the countries considered, with Germany, France and Spain placed in the middle. In addition, we consider more structural factors such as unit labor costs and employment protection legislation that will shape firm incentives in adopting AI as well as influencing whether AI is used to complement or replace workers (Figure 11). High labor costs clearly increase the incentive to automate, while stronger employment protection legislation tends to slow displacement, at least in the near- to medium-term. For instance, the US has the weakest labor market regulations among the countries considered, leading to a higher expected pass-through of the theoretical AI exposure into actual labor market impact. By contrast, Italy has the strongest employment protection legislation, which shields the economy from the AI-related labor market disruption at least temporarily but also limits the economy’s ability to benefit from AI-enabled productivity gains.

³ These figures refer to AI-exposed jobs in the US economy. While we assume that the split between augmentation, reorganization and automation is broadly similar across countries, we adjust for country-specific structural factors, which results in lower estimated impacts in our analysis. Source: Martin Richmond, A. (2026), “The AI jobs transition framework: Mapping AI’s near-term impact on jobs”, OpenAI Research Paper.


Figure 11: Employment protection legislation (indicator range from 0 (weak) to 6 (strong), left), and average annual hourly labor costs per employee (USD PPP, right)



Sources: OECD, ILO, Allianz Research.

The AI impact on the current workforce across major European economies and the US is substantial but far from uniform (Figure 12). To gauge the overall AI impact on the current workforce, we combine multiplicatively at the country level the sectoral AI exposure estimates, sectoral employment structures as well as the four scaling factors. The results are country-specific shares of jobs affected by AI, reflecting both technological potential and real-world constraints. Over the next 1–3 years, AI is expected to affect 23.3% of jobs across major economies, with reorganization (10.4% of jobs) dominating augmentation (5.3%) and outright displacement (7.6%). The share of jobs affected ranges from 9.2% in Italy to 28.7% in the US, with Spain (12.4%), France (14.7%), Germany (16.2%) and the UK (17.7%) in between. These differences closely mirror the distribution of employment across sectors. Economies with larger shares of finance, ICT and professional services exhibit substantially higher exposure than those with more manufacturing, construction or agriculture.

Figure 12: Medium-term AI impact on the current workforce across countries, in % and number of jobs



Source: Allianz Research.

Across all countries, the most important adjustment channel is job reorganization, while automation remains more material than augmentation. This reflects the fact that in most AI-exposed jobs, a human in the loop remains necessary, leading firms to reorganize tasks rather than replace workers, while only a smaller share of jobs meets the conditions for full automation and even fewer benefit from demand-driven expansion. In the US, 12.9% of jobs fall into the reorganization category, compared to 7–8% in Germany and the UK and 4–6% in Southern Europe. This finding is consistent with AI adoption currently targeting auxiliary or “side” tasks rather than core functions, reinforcing a complementary rather than substitutive dynamic. The share of jobs at risk ranges from 3.4% in Italy to

9.0% in the US, with most countries clustering around 5–6%. This gap between exposure and actual automation reflects the role of labor-market regulation, human-interaction requirements and demand effects, which limit the immediate substitutability of labor even where tasks are technically automatable. With respect to the expected real world labor impact of AI over the next 1-3 years, between 2% and 7% of jobs benefit from AI-driven expansion, thanks to augmentation, with the strongest effects in the US and UK. These gains arise where productivity improvements outweigh labor costs and hence increase employment.

The transatlantic divide in AI labor-market exposure is substantial and systematic: the US can expect more short-term AI-related labor market pain coupled with more long-term gain. In the US, 28.7% of jobs – equivalent to around 52.5mn positions – are affected by AI, compared to 14.5% (21.8mn jobs,) across the major European economies (Germany, France, Spain, Italy, UK). This pattern largely reflects structural differences in economic composition. The US economy is more heavily weighted toward high-exposure sectors such as information, finance and professional services, where AI can be deployed at scale, whereas Europe – particularly Southern Europe – retains a larger share of lower-exposure sectors such as manufacturing, construction and hospitality. The comparison between the US and Italy illustrates the extremes of this distribution: while the US combines high exposure (28.7%) with strong augmentation (6.9%), reorganization (12.9%) and automation (9.0%), Italy shows low exposure (9.2%) and correspondingly limited adjustment. As a result, the US faces greater transition pressures but also stronger potential productivity gains, whereas Europe’s lower exposure provides short-term stability but risks slower productivity growth and weaker AI-driven job creation over time.

A key limitation of these results is that they focus on the current workforce and do not account for AI-induced job creation. Our estimates should therefore not be interpreted as direct changes in headline employment. Historically, technological change has followed a consistent pattern: job destruction by automating tasks is offset over time by job creation through new tasks and demand effects.⁴ However, this adjustment is neither immediate nor evenly distributed. Job creation typically lags behind displacement due to reallocation frictions and skill mismatches. The time mismatch may be more pronounced for AI, as diffusion and adoption rates are unprecedented. Taken together, AI-driven job losses are likely to be partially offset over time, but with meaningful transition costs.

The scale of required job creation is material. Our results indicate that around 25mn jobs across the US and major European economies are subject to automation. Fully offsetting displacement would therefore require the creation of a comparable number of new jobs over the medium term. Early evidence shows strong growth in AI-related roles with some 1.6m new AI-related jobs like AI engineers, forward-deployed engineers and data annotators created over the past two years plus 600,000 jobs driven by the data center boom with the US the key beneficiary⁵ However, diffusion across sectors and countries remains uncertain. Net employment outcomes will depend on the speed of job creation relative to displacement and on the efficiency of worker reallocation. In ageing economies, declining labor supply may additionally cushion the employment impact as labor shortages persist in many sectors.

Preparing policy frameworks for the AI age

AI does not determine labor-market outcomes on its own. The technology sets the direction, but policy and institutions determine the speed, scale and distribution of its effects including the balance between job augmentation, reorganization and automation.⁶ What distinguishes the current wave of AI is not only its breadth across sectors, but the pace and depth of task-level transformation, amounting to one of the most significant labor-market transitions in recent decades. Existing policy frameworks, designed for slower, occupation-based change, are therefore ill-suited to the AI age, where jobs are continuously reconfigured. To remain effective, education systems, active labor-market policies, social protection systems, tax structures and firm-level incentives must be updated and rethought to become more forward-looking. The main challenge is therefore not preventing

⁴ See: Acemoglu, Daron, and Pascual Restrepo (2019), “Automation and New Tasks: How Technology Displaces and Reinstates Labor”, *Journal of Economic Perspectives*, 33 (2), 3-30; Hötte, K. Somers, M., and A. Theodorakopoulos (2023), “Technology and jobs: A systematic literature review”, *Technological Forecasting and Social Change*, 194, <https://doi.org/10.1016/j.techfore.2023.122750>.

⁵ See: LinkedIn (2026), *LinkedIn-labor-market-report-building-a-future-of-work-that-works-jan-2026.pdf*

⁶ For an excellent overview of the importance of policies and on how to make AI “pro-worker”, see: Acemoglu, D., Autor, D., and S. Johnson (2026), “Building Pro-Worker Artificial Intelligence”, NBER Working Paper 34854, <https://www.nber.org/papers/w34854>.

adjustment, but to shape its direction and distribution, ensuring that productivity gains translate into broad-based improvements in employment, wages and job quality rather than displacement or increased inequality.

First, labor-market policy will determine how smoothly workers move between tasks, roles and sectors.

Traditional active labor-market policies will need to evolve from reactive systems focused on unemployment toward anticipatory, data-driven and task-based approaches. This includes early-warning systems to identify AI exposure, individual learning accounts, and rapid retraining mechanisms that can be deployed before displacement occurs. Public employment services should be strengthened with real-time labor-market intelligence and AI-enabled matching tools that connect workers to adjacent roles based on transferable skills. Wage insurance, portable benefits and unemployment support linked to retraining will be essential to reduce income risk and support mobility. Since most AI-exposed jobs are likely to be redesigned rather than eliminated, the priority is not protecting existing tasks, but enabling workers to transition quickly into new ones. Since most AI-exposed jobs are likely to be redesigned rather than eliminated, the priority is not protecting existing tasks, but enabling workers to transition quickly into new ones.⁷

Second, education and skills will become more binding constraints. As AI increasingly complements cognitive work, the ability to work alongside AI systems will determine whether workers benefit from productivity gains or face substitution.⁸ This requires embedding AI literacy as a core competence across education systems, alongside stronger emphasis on critical thinking, problem-solving, adaptability and communication. Education systems will need to become more modular and continuous, with expanded mid-career training, short-cycle credentials and closer alignment with labor-market needs. Governments should also support sector-specific training ecosystems and local partnerships between firms, unions and education providers to accelerate adjustment in regions and occupations most exposed to disruption.

Third, firm incentives, including taxation and regulations will shape whether AI is used to replace labor or to augment it.⁹ A key issue is that labor is typically taxed more heavily than capital, which creates a structural bias toward automation even when augmentation would be more socially efficient. Tax systems therefore need to be adjusted to incentivize augmentation rather than the substitution of labor. One proposal is a “robot tax,” aimed at reducing this bias by taxing automation more in line with labor—for example, by limiting preferential tax treatment of capital that substitutes for workers or linking levies to displacement. Implementation challenges would need to be addressed, including defining what constitutes a “robot,” avoiding disincentives to innovation, and limiting adverse competitiveness effects if policies are not coordinated internationally. Whereas broad-based taxation of automation may be difficult to implement in practice, more targeted approaches could be more effective. These include tax incentives to support augmentation, such as credits for training and reskilling, incentives for job redesign, and support for firms investing in human capital alongside technology. Frameworks can also encourage sharing of productivity gains through wages, profit-sharing, or reduced working time. Policies to support AI adoption among SMEs, strengthen competition, and prevent excessive concentration of productivity gains will also be critical.¹⁰ Even small differences in these incentives can lead to large differences in business and labor-market outcomes.

Finally, the broader societal response will matter: Public concerns around economic disruption, trust and safety, and the distribution of gains may lead to calls for stricter AI regulation and shape how AI diffuses across economies.¹¹ Maintaining social legitimacy will therefore be critical. This requires stronger transparency in the use of AI in the workplace, appropriate safeguards around its deployment, and mechanisms for worker voice and participation in how AI is introduced and used. Ensuring that the gains from AI are broadly shared will also be central to sustaining support for adoption. This includes profit-sharing mechanisms and competition frameworks that prevent excessive concentration of AI-driven rents. A complementary priority is to ensure that AI adoption is broad-based rather than concentrated. This calls for policies that lower barriers to adoption, including support for SMEs, promote competition, and facilitate the diffusion of AI capabilities across firms and sectors. Such measures will be critical to ensure that productivity gains translate into wider economic benefits and employment growth.

⁷ See: Martin Richmond, A. (2026), “The AI jobs transition framework: Mapping AI’s near-term impact on jobs”, OpenAI Research Paper.

⁸ See: Jaumotte, F. et al. (2026), “Bridging Skill Gaps for the Future: New Jobs Creation in the AI Age”, IMF SDN/2026/001.

⁹ See: Acemoglu, D., Autor, D., and S. Johnson (2026), “Building Pro-Worker Artificial Intelligence”, NBER Working Paper 34854, <https://www.nber.org/papers/w34854>.

¹⁰ See: Ferrando, A. et al. (2026), “Adopting and Investing in AI: Evidence from Euro Area Firms”, ECB Economic Bulletin, Issue 2/2026.

These assessments are, as always, subject to the disclaimer provided below.

FORWARD-LOOKING STATEMENTS

The statements contained herein may include prospects, statements of future expectations and other forward-looking statements that are based on management's current views and assumptions and involve known and unknown risks and uncertainties. Actual results, performance or events may differ materially from those expressed

or implied in such forward-looking statements.

Such deviations may arise due to, without limitation, (i) changes of the general economic conditions and competitive situation, particularly in the Allianz Group's core business and core markets, (ii) performance of financial markets (particularly market volatility, liquidity and credit events), (iii) frequency and severity of insured loss events, including from natural catastrophes, and the development of loss expenses, (iv) mortality and morbidity levels and trends,

(v) persistency levels, (vi) particularly in the banking business, the extent of credit defaults, (vii) interest rate levels, (viii) currency exchange rates including the EUR/USD exchange rate, (ix) changes in laws and regulations, including tax regulations, (x) the impact of acquisitions, including related integration issues, and reorganization measures,

and (xi) general competitive factors, in each case on a local, regional, national and/or global basis. Many of these factors may be more likely to occur, or more pronounced, as a result of terrorist activities and their consequences.

NO DUTY TO UPDATE

The company assumes no obligation to update any information or forward-looking statement contained herein,

save for any information required to be disclosed by law.

Allianz Trade is the trademark used to designate a range of services provided by Euler Hermes.

Allianz Research | 29 April 2026

Staycation summer? Jet-fuel crunch reshapes the peak holiday season

Ludovic Subran
Chief Investment Officer and
Chief Economist
ludovic.subran@allianz.com

Ano Kuhanathan
Head of Corporate Research
ano.kuhanathan@allianz-trade.com

Maria Latorre
Sector Advisor, B2B
maria.latorre@allianz-trade.com

In Summary

- **The Middle East crisis is squeezing airlines' jet-fuel supplies.** Unlike previous oil crises, the main bottleneck lies in refining capacity and product logistics. The Strait of Hormuz accounts for roughly 25% of global seaborne jet-fuel trade, making it a critical artery for aviation markets. As regional refinery operations have been disrupted, jet-fuel prices have doubled since the start of the conflict, while refining crack spreads have moved above USD 100/bbl. Fuel is the largest cost line (30–35%), meaning price shocks translate almost immediately into higher fares, capacity cuts or margin erosion.
- **Europe is among the most structurally exposed regions.** Europe produces only around 50% of its kerosene needs domestically, with the remainder met through imports. Gulf producers account for approximately 70% of imported volumes, affecting markets such as the UK, Germany, France and Italy, all of which run persistent deficits. Middle East kerosene flows to Northwest Europe fell -90% m/m in March, while April flows were effectively nil. Although US shipments surged +782% m/m, total combined inflows from the US and Middle East were still down -82% m/m in April, suggesting tightening physical availability into late spring and summer. More worryingly, even if Hormuz reopens soon, a full refinery ramp-up would likely take 3–6 months.
- **Airlines are reacting through fare increases (+5-15% for international routes) and tighter capacity management (2-5% cuts in Europe).** Across markets, carriers are attempting to preserve margins via pricing power. Beyond the fare hikes, specific fuel surcharges now range from USD20-60 on short/medium-haul routes and USD80–150 on long-haul tickets. If conditions worsen, further fare increases of 10–15% are likely. In Europe, announced capacity reductions remain selective, concentrated on lower-yield short-haul routes and secondary airports. Low-cost carriers are especially vulnerable because of thin margins, short-haul concentration and strong competition from high-speed rail alternatives.
- **Will it be the summer of staycations? Some substitution benefits are emerging in Southern Europe, but the upside is limited.** Equity markets have repriced airlines sharply while rewarding Western Mediterranean hospitality firms, with listed hotel stocks up +36% since the onset of the conflict. Booking trackers point to demand gains of +32% y/y for Spain and around +20% for Italy, Greece and Portugal. However, weakening sentiment in the US and Eurozone means many households may reduce total leisure spending rather than fully replace international trips with long domestic holidays.
- **Meanwhile, tourism in the Middle East faces a sharp reversal after strong post-pandemic growth.** Before the conflict, international arrivals were expected to increase by +13% y/y in 2026, after +51% in 2025 versus 2019, thanks to visa reforms and heavy tourism investment. If hostilities persist for another month, arrivals could instead decline by -35-40% y/y, implying approximately USD70-75bn in lost tourism receipts. Smaller tourism-dependent economies – i.e. Lebanon (where tourism accounts for 9.1% of GDP), Bahrain (6.2%), the UAE (6.2%) and Jordan (5.9%) – are most exposed to sudden declines in arrivals, foreign-exchange receipts and hospitality demand. Countries with weaker external balances may face stronger pressure on reserves and exchange rates. Beyond the Middle East, island tourism destinations where aviation is the primary gateway (Seychelles, Maldives and Mauritius) will also feel the pinch from disrupted connectivity through Middle Eastern air hubs.

Turbulence ahead: Jet-fuel crunch threatens airline operations in peak season


A supply shock beyond crude: the central role of refining constraints. The current crisis differs materially from previous oil price shocks in that it is not limited to crude supply disruptions. While the ongoing closure and uncertainty surrounding the Strait of Hormuz has constrained global crude flows, the more persistent bottleneck lies in refining capacity, for which the Middle East represents a huge market share (see Figure 1). Jet-fuel availability actually depends on complex refining processes and disruptions across Middle Eastern facilities have created a lag between the normalization of crude supply and the restoration of refined product output. The Strait of Hormuz accounts for around 25% of the global seaborne volume of transported jet fuel. However, its outsized importance stems less from sheer volume than from its role as a core export hub supplying structurally import-dependent markets, particularly in Asia and Europe, where domestic refining capacity is insufficient to meet aviation demand. This imbalance is further exacerbated by export restrictions from key refining hubs such as China and South Korea, as well as limited spare capacity in alternative regions. As a result, jet-fuel prices have doubled since the beginning of the conflict, driven not only by crude benchmarks such as Brent crude oil, but increasingly by elevated refining margins (see Figure 2), with the crack spread surpassing the USD 100/bbl threshold. With jet fuel typically accounting for approximately 30–35% of airlines total operating costs under normal conditions (making it the single largest cost item across much of the industry), airlines are now bracing for a challenging summer, increasingly forced to balance pricing power, capacity discipline and margin protection.

Figure 1: Oil refineries across the world, operational vs permanent closure, bubble size = capacity



Sources: Bloomberg, Allianz Research

Figure 2: Jet fuel and Brent crude oil prices (USD/bbl)




Sources: Bloomberg, Allianz Research

The decoupling of crude-oil and jet-fuel prices makes traditional fuel-hedging mechanisms significantly less effective. In parallel, regional hedging approaches are diverging. Most hedging programs are linked to crude oil benchmarks rather than the refined jet fuel that airlines actually consume. As a result, their effectiveness is reduced in periods of unusually wide crack spreads. More broadly, hedging practices vary significantly across regions. US airlines largely moved away from systematic fuel hedging about a decade ago, following adverse experiences in previous cycles. Large carriers instead rely on pricing flexibility and capacity management to absorb fuel-cost volatility. By contrast, European airlines have historically maintained more extensive hedging programs, often covering a significant share of fuel needs over a 12- to 24-month horizon. While this provides greater short-term cost visibility, some European carriers have recently signaled a reduction or cessation of hedging activity, reflecting a deteriorated risk-reward profile and a growing preference for market exposure combined with more dynamic fare adjustments. Asian airlines generally adopt an intermediate approach, typically layering hedges six to 12 months ahead, though strategies vary considerably by carrier and jurisdiction.

How much longer the Strait of Hormuz will remain closed is still uncertain. But once opened, the ramp-up of crude production and refinery activities across the region is expected to take about three to six months. We estimate that about half of crude production is currently curtailed in the Middle East, most of it via precautionary and stock-management measures rather than physical damage to fields. However, the shape of the recovery will depend on how long the Strait remains closed. Under our baseline scenario, a deal allowing for a progressive normalization of energy trade should happen in May, with a downside scenario being a closure of the Strait for more than three months. In either one, Gulf production is likely to mostly recover within a few months once the Strait safely reopens. The publicly reported physical damage to oil fields remains limited, in contrast to LNG assets. Saudi Arabia and the UAE together hold an estimated 2 Mbpd of spare capacity, enough to meaningfully stabilize markets in the first weeks. Kuwait production could take 3-4 months to return to full output even if the war ended today – a reminder that a "fast ramp up" applies to Saudi/UAE spare capacity, not to the wells that have been forced into precautionary curtailment. We should also underline that the longer the closure, the slower the production ramp up. We estimate a 70% recovery of lost production three months after Hormuz reopens, reaching close to 90% after six months. This could leave a long tail of structural underproduction that keeps the downstream product supply tighter than markets are pricing. However, downstream infrastructure such as refineries can come back online much quicker, in a matter of days.

The looming jet-fuel shortage will mainly affect import-dependent aviation hubs with tight refining capacity and heavy exposure to Middle Eastern supply flows. Asia-Pacific, Africa and Europe remain structurally dependent on oil and oil-derivative products from the Middle East, albeit for different reasons. On the one hand, Asia-Pacific has significant refining capacity but limited domestic crude supply, and with demand growth (particularly in China and India) continuing to outpace supply, dependence on Gulf crude remains high, with roughly 40–50% of APAC crude imports coming from the Middle East. For Africa and Europe the dynamic is different. Africa exports crude (notably from Nigeria and Angola) but lacks refining capacity, leaving it reliant on imported middle distillates (such as jet fuel and diesel) from suppliers such as Saudi Arabia and the UAE (see Figure 3). In Europe, reliance is largely driven by substitution effects. Since the Russian invasion of Ukraine, Middle Eastern flows have increasingly replaced lost Russian supply, particularly in refined products such as jet fuel. While Europe still refines around 50% of its jet fuel domestically, this share has been gradually declining amid a long-term downscaling regional trend driven by refinery closures and capacity rationalization. The remaining 50% is met through imports, with Gulf suppliers playing a central role (around 70%), particularly Kuwait as one of the largest single exporters, while the UK remains among the most import-dependent markets due to its structural refining deficit.


Figure 3: Destination region of refined oil products transiting via the Strait of Hormuz, thousand barrels per day



Sources: BloombergNEF, Allianz Research

Europe’s kerosenemarket is structurally vulnerable, with most major economies running persistent deficits. The UK, Germany, France and Italy show the largest shortfalls (Figure 4), underscoring their reliance on external supply to meet aviation demand. Even traditionally well-refined economies like the Netherlands and Spain were only modestly in surplus last year, while several smaller markets remain balanced or slightly positive, indicating limited regional capacity to offset the gap. This imbalance effectively positions Europe as a net structural importer of kerosene. As a result, European aviation activity is indirectly exposed not only to global oil price dynamics but also to geopolitical and logistical risks along key supply routes, reinforcing the region’s dependence on external refining hubs for a fuel that is essential to long-haul connectivity.

Figure 4: Kerosene balances in 2025, by country in Europe, in thousand barrels per day




Sources: BloombergNEF, Allianz Research. Note: Negative numbers indicate supply shortfall, meaning consumption outweighs domestic refinery output of kerosene.

Meanwhile, the US is becoming a key marginal supplier to Europe’s jet-fuel market. Kerosene flows from the Middle East to Northwest Europe declined by -90% m/m in March while no shipments departed in April. Conversely, shipments from the US skyrocketed last month (+782% m/m), reflecting a sharp but likely temporary reorientation of regional supply flows. In the short term, replacing Middle Eastern jet fuel with supply from the US enhances Europe’s resilience, but this shift carries structural costs: longer transatlantic routes and increased transport expenses and emissions. Moreover, US crude yields less jet fuel per barrel – tightening refining economics – and dependence is not eliminated but redirected, concentrating risk in transatlantic supply chains (and increasing President Trump’s bargaining power), while heightening exposure to price volatility, making the arrangement a stabilizing stopgap rather than a durable solution.

More worryingly, even after accounting for increased inflows from the US, total kerosene imports into Northwest Europe continue to decline. When combining supplies from both the US and the Middle East, shipments so far in April are down by -82% m/m (see Figure 5-A), pointing to a material tightening of physical availability and raising the likelihood of an outright supply shortfall by late May if the trend persists. Two mitigating factors provide only temporary relief: first, Europe still produces roughly half of its kerosene demand domestically, anchored by its refining base and second, jet-fuel cargoes dispatched in April (albeit small) will continue to arrive with a lag into May, smoothing the abrupt impact. However, these buffers are finite, and with inventories already under pressure and logistics stretched, they primarily delay rather than eliminate the risk of fuel disruption.

Figure 5: A) Kerosene flows to Northwest Europe per departure month and by origin; B) Total kerosene flows from the Middle East and the US together to Northwest Europe (departure vs arrival month). In millions of barrels




Sources: Refinitiv Workspace, Allianz Research

Pricing power and capacity discipline are the first line of defense

To cope with the situation, airlines are increasingly converging on a common set of strategies centered on revenue enhancement. Airfares are among the most immediate pressure valves. Asia-Pacific carriers were the first to move decisively, announcing fare increases within days of the fuel-price spike, signaling early that higher jet-fuel costs would be passed through to passengers, either via direct fare hikes or the reintroduction of fuel surcharges. In magnitude, increases so far have generally landed in the high single-digit to mid-teens range (+5-15%) on international routes, with some long-haul segments seeing steeper adjustments depending on fuel intensity and route disruption. In the US, airlines adjusted prices later but more incrementally, delivering a +23% increase in domestic segments over the quarter and +60% on international flights (see Figure 6).

Figure 6: Average airfares in the US, in USD; A) domestic flights; B) international flights



Sources: KAYAK, Allianz Research

Beyond airfares, airlines are pulling additional revenue levers – expanding ancillary charges through higher baggage fees and explicit fuel surcharges – while, in parallel, capacity discipline is becoming the norm, particularly across Europe. Checked baggage fees have increased by USD5-10 per bag per segment, while ancillary revenues from preferred seating and extra-legroom options are up 10-20% (typically USD20-80 per segment, depending on route and demand). In parallel, several carriers – particularly in Europe and Asia – have reintroduced explicit fuel surcharges ranging from USD20-60 on short/medium-haul routes and USD80-150 on long-haul tickets. Should the jet-fuel shortage get worse, airlines could implement further fare increases in the range of +10–15%, with the upper end more likely on long-haul and fuel-intensive routes, complemented by higher fuel surcharges. To the detriment of travelers, airfares are unlikely to fully revert to pre-shock levels, regardless of the duration of the ongoing fuel disruption. Once higher fare levels are accepted by the market, they tend to persist and are rarely fully rolled back, even if underlying cost pressures subsequently ease. As a result, part of the recent fare increases is likely to remain embedded, supporting a persistently higher pricing floor across key routes.

At the same time, airlines are reinforcing pricing power through capacity discipline, particularly on lower-yield short-haul routes. This adjustment has been most pronounced among European carriers, where network trimming is increasingly used as a margin-protection lever. Structurally, airfares are driven by a relatively fixed cost composition: 60–75% operating costs, 20–40% taxes and airport/navigation fees and 0–10% profit margins, underscoring the industry's limited flexibility. European airlines in particular stand out for paying the highest airport charges globally, which make short-haul trips especially less profitable due to taxes and air navigation fees. Against this backdrop, several airlines in the region have begun rolling back planned capacity expansions for 2026, signaling a shift from growth to network rationalization and yield protection. With airport taxes largely fixed and difficult to pass through dynamically, carriers are increasingly reallocating capacity toward long-haul routes, where unit revenues are higher and fuel costs can be better absorbed, while reducing exposure to structurally weaker short-haul segments.


Despite visible route cuts, announced capacity reductions in Europe represent only 2–5% of total sector capacity, reflecting targeted network optimization rather than a systemic contraction. So far, the announced capacity adjustments in Europe remain relatively contained in aggregate terms, but are more meaningful at the margin. Overall, they represent roughly 2–5% of total sector capacity, reflecting targeted reductions rather than a broad-based contraction. The cuts are concentrated in lower-yield short-haul routes and secondary airports, where load factors are weaker and pricing power is limited. This pattern is particularly visible among low-cost carriers (LCC), which are actively reallocating aircraft toward higher-margin leisure routes and core city pairs. In contrast, long-haul and high-demand corridors remain largely unaffected in terms of capacity.

Low-cost carriers (LCCs) are indeed the most exposed segment in Europe: structurally thin margins and a short-haul-heavy network amplify vulnerability, facing substitution risk from rail operators. The low-cost carrier model in Europe is built on high-volume, short-haul, point-to-point operations; standardized fleets and maximized aircraft utilization, delivering typically 5–10% EBIT margins. This structure depends heavily on cost discipline, making fuel a disproportionately large share of total operating costs and a key source of margin volatility. As a reference, a 10% increase in jet-fuel prices typically reduces airline EBIT margins by roughly 0.3–0.8pp (depending on hedging coverage and the pricing power of each player). As a result, a doubling of jet-fuel prices would imply an EBIT margin compression of approximately 3–8pps at the sector level over a 12-month horizon. Indeed, some of the casualties announced so far in the sector have been low-cost carriers that were already under financial stress prior to the jet-fuel price spike and have so far relied on government support. Positively, evidence shows that in this particular sector, fuel-driven stress does not necessarily lead to failure, but rather to government-backed stabilization. Nevertheless, the customer base in this segment is highly price-sensitive, with fares often in the EUR40–100 range, limiting the ability to pass through higher costs without materially impacting demand. LCCs account for roughly 45% of intra-European flights and over 50% of short-haul capacity on many routes, where competition is most intense and pricing power is structurally constrained. Exposure is even more acute in markets such as Spain, Italy, France and Germany, where dense low-cost penetration overlaps with extensive high-speed rail corridors, increasing substitution risk on short distances.

Gulf hit severely and spillover to Asia and Africa

A different reality for Middle East tourism: the summer season is already being derailed. Beyond coping with jet-fuel scarcity, airlines in the Middle East also have to deal with big revenue losses and operational constraints due to the closure of almost all of the region's airspace. The conflict hit in the midst of a tourism boom: The UAE, Saudi Arabia and Qatar experienced a strong post-pandemic rebound (Figure 7), with the number of international visitors growing by +51% in 2025 vs 2019, thanks to simplified visa regimes and aggressive investment in aviation and tourism infrastructure across the region, while post-pandemic luxury tourism also turned the region into a high-growth destination. Before the start of the conflict, the region was set to witness a double digit increase in the number of international visitor arrivals (+13% y/y) this year. Now, it is likely to see a decline of -35-40% y/y if the conflict extends for another month, representing a loss of around USD70-75bn in tourism receipts over the year (Figure 8).

Figure 7: International tourist arrivals per year, million people



Source: UN Tourism, Allianz Research


Figure 8: Middle East tourism data



Source: UN Tourism, Oxford Economics, Allianz Research

Tourism-dependent economies are most at risk. The greatest direct vulnerability lies in economies where travel and visitor spending make up a meaningful share of national output. Lebanon, with tourism equivalent to 9.1% of GDP, is the most exposed (see Figure 9). Any prolonged decline in arrivals would likely intensify existing macroeconomic stress, reduce foreign-currency inflows and further weaken already fragile domestic demand. Bahrain (7.1%) and the UAE (6.2%) would also face material losses, particularly through reduced leisure travel, lower hotel occupancy and softer aviation and retail activity. In the UAE the effect would be partly cushioned by economic diversification, but tourism-linked sectors would still experience a noticeable slowdown. Jordan (5.9%) is similarly sensitive, as perceptions of regional instability often affect inbound travel regardless of whether the country is directly involved. This highlights a common feature of tourism shocks: reputational spillovers frequently matter as much as physical disruption. Saudi Arabia (4.4%) would likely see more moderate effects because of its larger domestic economy and expanding religious tourism base, though discretionary travel could soften. Qatar (2.9%) and Oman (2.4%) face lower direct exposure, but airlines, hospitality and transit traffic remain vulnerable. At the lower end, Egypt, Israel, and Kuwait appear less dependent on tourism as a share of GDP. However, even where direct exposure is smaller, foreign-exchange earnings can still be important. Countries with weaker current-account positions or limited reserve buffers may feel the impact quickly through exchange-rate pressure, tighter financing conditions and slower growth. In practice, this means smaller tourism-dependent economies are likely to absorb the sharpest immediate pain, while larger diversified economies would experience more manageable but still meaningful secondary effects.

Figure 9: Tourism direct share in GDP for Middle East countries (%)



Note: latest available year in parenthesis

Source: UN Tourism, Allianz Research

Overall, direct exposure to the conflict zone represents only 10% of global airline capacity. However, the Middle East is a critical global aviation hub, so any disruption has disproportionate spillover effects, impacting long-haul connectivity and tourism flows to regions that are largely dependent on air travel. With airports such as Dubai International Airport (DXB), Hamad International Airport and Abu Dhabi International Airport handling tens of millions of transit passengers annually and serving as critical connecting nodes between Europe and Asia, the ongoing airspace closure is also causing operational challenges for airlines, while impeding the connection to specific routes. DXB alone regularly ranks as the world's busiest international airport, with nearly 90mn passengers transported in 2025, underscoring the region's role as a structural bridge in global air traffic flows. As a result, any disruption in the region's airports has disproportionate ripple effects across long-haul global connectivity, particularly on Europe–Asia routes that rely heavily on Gulf-based transit hubs, and which have been growing significantly in the past year (Figure 10).


Figure 10: Average daily departure and arrival flights, Europe vs the rest of the world

Region	Average daily flights	% previous year		% 2019	
Intra-Europe	23,305	↑	+3%	↓	-1%
Europe ↔ Middle East	1,508	↑	+11%	↑	+6%
Europe ↔ North-Africa	1,375	↑	+11%	↑	+34%
Europe ↔ North Atlantic	1,369	↑	+3%	↑	+15%
Europe ↔ Asia/Pacific	969	↑	+13%	↑	+22%
Europe ↔ Other Europe	370	↑	+2%	↓	-64%
Europe ↔ Southern Africa	321	↑	+5%	↑	+3%
Europe ↔ Mid-Atlantic	176	↓	-1%	↑	+1%
Europe ↔ South-Atlantic	206	↑	+7%	↑	+11%
Non Intra-Europe	6,293	↑	+8%	↑	+2%

Sources: Eurocontrol, Allianz Research

Several key tourism destinations – such as the Seychelles, Maldives, Mauritius, Thailand, Indonesia, Australia, Singapore and the Philippines – are highly dependent on long-haul air connectivity routed through Middle Eastern hubs. The top three regional carriers play a pivotal role in linking these markets to Europe and the Americas, often providing the most efficient or even only viable one-stop connections. As a result, any disruption or capacity tightening in the Middle East aviation system has a disproportionate impact on these destinations, some which rely heavily on tourism income. The consequence is a form of structural vulnerability: tourism demand in these economies is highly elastic to air-connectivity conditions, meaning that even marginal increases in fuel costs or route rationalization can directly affect visitor flows. This dynamic is particularly relevant for island economies such as the Maldives, Seychelles and Mauritius, where aviation is not just a transport mode but the primary gateway for tourism demand (Figure 11).

Figure 11: Share (%) of international inbound passenger traffic carried by top three gulf airlines, by destination.




Sources: Various sources, Allianz Research

Expect some substitution in developed markets but confidence caps the upside

The European summer holidays are not fully at risk. Indeed, rail transport in Europe carries substantially more passengers annually than aviation, making it the dominant mode of mass mobility across the region. While air travel carries hundreds of millions of passengers per year, rail systems collectively handle volumes in the billions, particularly through dense short- and medium-distance networks (Figure 12). The evolution of rail demand in Europe over the past decade reinforces the potential for modal shift as a key solution for travelers. Between 2014 and 2024, railways transported 81.2bn passengers, compared with 8.9bn transported by air companies over the same period (10 times less). The rail network's high frequency, extensive geographic coverage and strong city-center connectivity allow it to absorb a meaningful share of displaced demand. In this context, rail does not merely complement aviation – it acts as a structural backstop for intra-European mobility, helping to sustain tourism flows even amid fluctuations in airline capacity. But ultimately, the question many households will face is not which mode of transport to choose, but rather how much they are willing to allocate to leisure spending this summer, given the prevailing economic conditions.


Figure 12: Evolution of demand in air transport and railways in Europe (2014 – 2024)



Source: Eurostat, Allianz Research

The substitution upside is bounded by consumer confidence, which has collapsed alongside the war in Iran. Low consumer confidence in the US and the Eurozone reduces the likelihood that households will fully substitute cancelled long-haul or Middle East travel with additional spending on domestic or nearby holidays. In theory, geopolitical disruption should redirect demand toward safer and closer destinations, benefiting local tourism markets. In practice, that mechanism weakens when households are already cautious about their finances. In the US, the University of Michigan consumer sentiment index fell to 53.0 in March from 56.4 in January, signaling deteriorating household confidence (see Figure 13). In the Eurozone, the Economic Sentiment Indicator declined to 96.6 in March from 99.2 in January, likewise pointing to softer consumer and business expectations. When confidence indicators fall to these levels, consumers tend to interpret uncertainty as a signal to preserve cash, postpone discretionary purchases and avoid committing to non-essential travel. Rather than replacing a cancelled overseas trip with a domestic break, many households simply travel less, shorten trips, downgrade accommodation or delay plans altogether. Tourism is a highly income-elastic category of spending, meaning it reacts quickly to shifts in sentiment, employment expectations and perceptions of future disposable income. Concerns over inflation, borrowing costs, slower growth and softer labor markets therefore weigh disproportionately on holiday demand. Even where some substitution occurs, domestic travel usually generates lower per-capita spending than long-haul international holidays involving flights, premium hotels and longer stays.


Figure 13: Economic sentiment indicator (ESI) and Consumer Sentiment indices



Source: Eurostat, Allianz Research

Equity markets are betting on West-Med hospitality with stock prices up more than 35% since the onset of the war in Iran. Regional airline baskets have repriced quite sharply (see Figure 14). Asian airlines have corrected the most (-18% since the war), with investors judging that bypass-routing and higher fares cannot offset the fuel cost. US airlines stand at -7%, with most US carriers running unhedged and fully exposed to spot fuel. The lower price decline is also related to pricing strategies as US fares can adjust faster. European airlines sit at -15% as investors see them squeezed between higher fuel cost and carbon costs. The substitution thesis is benefiting West-Med listed hotel firms, which are up +36% since 28 February (i.e. Spanish, Italian and Greek hospitality), in line with surging flight bookings to these destinations: Industry trackers cite +32% y/y for Spain and +20% for Italy, Greece and Portugal.

Figure 14: Selected equity indicators (Feb. 27 = 100)



Note: each index aggregates large, listed company by market cap
Source: Eurostat, Allianz Research

These assessments are, as always, subject to the disclaimer provided below.

FORWARD-LOOKING STATEMENTS

The statements contained herein may include prospects, statements of future expectations and other forward-looking statements that are based on management's current views and assumptions and involve known and unknown risks and uncertainties. Actual results, performance or events may differ materially from those expressed

or implied in such forward-looking statements.

Such deviations may arise due to, without limitation, (i) changes of the general economic conditions and competitive situation, particularly in the Allianz Group's core business and core markets, (ii) performance of financial markets (particularly market volatility, liquidity and credit events), (iii) frequency and severity of insured loss events, including from natural catastrophes, and the development of loss expenses, (iv) mortality and morbidity levels and trends,

(v) persistency levels, (vi) particularly in the banking business, the extent of credit defaults, (vii) interest rate levels, (viii) currency exchange rates including the EUR/USD exchange rate, (ix) changes in laws and regulations, including tax regulations, (x) the impact of acquisitions, including related integration issues, and reorganization measures,

and (xi) general competitive factors, in each case on a local, regional, national and/or global basis. Many of these factors may be more likely to occur, or more pronounced, as a result of terrorist activities and their consequences.

NO DUTY TO UPDATE

The company assumes no obligation to update any information or forward-looking statement contained herein, save for any information required to be disclosed by law.

Allianz Trade is the trademark used to designate a range of services provided by Euler Hermes.

Allianz Research | 23 April 2026

Energy shock and policy response: Once bitten, twice shy?

Ludovic Subran
Chief Investment Officer & Chief
Economist
ludovic.subran@allianz.com

Lluís Dalmau Taules
Economist Africa and Middle East
lluis.dalmau@allianz-trade.com

Bjoern Griesbach
Head of Macroeconomic and Capital
Markets Research
bjoern.griesbach@allianz.com

Maddalena Martini
Senior Economist Southern Europe
and Benelux
maddalena.martini@allianz.com

In Summary

- **The current energy shock is fundamentally a supply disruption, not just a price event, with Asia most at risk.** The blockage of the Strait of Hormuz has disrupted around 20mb/d of oil flows (roughly 1/5th of global supply), leaving a net shortfall of ~10mb/d despite rerouting, additional supply and strategic reserve releases. The resulting supply-demand gap has been closed through higher oil prices (Brent +30% since the war began), which have lowered consumption. Physical scarcity is already visible: jet fuel is tightening, diesel is being rationed and industrial users in price-controlled markets face outright shortages. Asia is most exposed (85% of Hormuz oil flows vs. less than 6% to Europe and 4% to the US), driving a persistent 6–7% price premium above Brent. A first policy response, mainly by Asian countries, has been demand rationing (shortened working weeks, energy usage and fuel restrictions) cutting global demand by ~1mb/d.
- **Fiscal support followed swiftly, primarily through fuel-tax relief and subsidies (~0.15% of GDP in developed markets and ~0.20% in emerging markets on average).** On price support Asia leads again, with large EMs (South Korea, India, Indonesia) deploying fiscal packages >1% of GDP vs. a global average of ~0.2%. Similarly, in developed markets, large EU energy importers average ~0.4% of GDP vs. ~0.15% overall. However, deeper pockets buy more: Germany's small 0.04% of GDP package translates into USD22 per capita while India's package of 1.5% of GDP, despite being 37 times larger in relative terms, only translates to a mere USD40 per capita. Should the conflict prove more prolonged (beyond May), a "second salvo" of fiscal support would raise spending to 0.6-0.8% of GDP, as governments would respond to higher energy price pass through with a 2-6 months lag. Nonetheless, no fiscal transfer can resolve the supply-demand mismatch, which will need to be cleared via demand destruction.
- **Compared to 2022, the shock differs in two key ways: fiscal support is materially smaller, but proportionally similar (around 50% in Europe as the size of the energy price shock is also lower) and price transmission has weakened due to structural shifts in the energy mix.** This reflects policy learning, the perception of a temporary shock and fiscal limits, with governments allowing greater pass-through rather than repeating large-scale shielding (e.g. France announcing later and more targeted measures, Belgium taking no measures). Meanwhile, in the European power market, higher wind and solar penetration has reduced the gas-to-electricity pass-through seen in 2022, but gas remains indispensable for grid balancing Spain stands out with electricity prices around 70% lower than in Italy, reflecting a higher renewables share and better storage capabilities.
- **For now, fiscal costs for most countries are negligible but EMs with high energy import dependence and a deteriorating current account are vulnerable (Türkiye, Egypt, Morocco, and Hungary).** Debt-servicing costs will rise by around 0.05% of GDP annually in advanced and emerging economies on average. entirely from a higher debt stock. Interest rates have risen only modestly (US/DE: +30bps, EM hard currency: +20bps, local currency: +30bps), with an even lower change in real rates. A gradual pass-through given slow debt rollover raises debt-servicing costs only marginally through this channel. On the flip side, higher inflation could even dampen the debt burden (inflation tax) and thus debt-service costs but the currently expected increase in inflation is far below 2022, thereby limiting this

effect. Some EMs nevertheless stand out: fiscal fragility and elevated debt service are most acute in Egypt, Hungary and Poland. Currency depreciation amplifies the damage in nations with a high share of hard currency debt: Türkiye, Argentina and Egypt. Nevertheless, not all emerging markets are on the losing side of the energy shock: Higher prices are a windfall for energy producers such as Nigeria, Brazil and Colombia.

From price signals to physical scarcity: a regime shift

The 2026 energy shock is not a price event with physical consequences, but a physical supply disruption with price consequences. This distinction is central for policy design, fiscal outcomes and market pricing. Unlike 2022, prices are not clearing the market: the shock is defined by missing volumes, not by temporarily expensive ones. Where supply is physically constrained, no price cap or political statement can fully offset a shock defined by missing volumes: the molecules are simply not there. This raises the risk that adjustment occurs through actual scarcity (notably in jet fuel and refined products), not just higher prices. Energy markets entered 2026 in a fragile but manageable equilibrium, and the initial repricing following the Iran escalation was driven largely by geopolitical risk premia: markets were pricing the probability of disruption, not disruption itself. That quickly gave way to reality.

From a supply gap of 20mb/d from the global supply – roughly one-fifth of world supply – to a shortfall of approximately 10 mb/d thanks to alternative supply sources and routes. Half of the missing gap is being rerouted via alternative pipelines, from the East-West pipeline in Saudi Arabia to the Emirate East coast pipeline. On top of that extra supply from Russia – whose oil has been temporarily excluded from the sanctions list – and the US – whose producers increased production by around 200,000 b/d from pre-war projections – is also substituting part of that supply, resulting in 10mb/d gap. The International Energy Agency (IEA) coordinated the release of around 400mn barrels from strategic reserves in the weeks following the start of the war – the largest joint drawdown ever announced — aimed at stabilizing prices and signaling policy control. Markets were largely unconvinced: prices retraced only marginally, underscoring that reserves can smooth flows temporarily but cannot offset a sustained physical disruption. This option is now fully exhausted.

The effective closure of the Strait of Hormuz marked the regime shift. Around 20% of globally traded oil and 25% of LNG transit the Strait, with no short-run substitute route of comparable capacity – making the closure not just a price shock but a physical supply event with no historical precedent of comparable scale. Since the start of military engagement, oil prices have risen by more than 50% and natural gas prices by over 65%, with storage levels falling towards their 2022 lows. But the aggregate numbers obscure where the pressure is most concentrated: Asia, as the region most dependent on Gulf supply, is bearing a disproportionate share of the cost. DME Oman futures have been trading at a premium of approximately USD10-11 per barrel above Brent since early April – a spread that under normal conditions runs in the opposite direction — capturing with unusual precision where the physical shortage is concentrated and who is paying for it.

Asia sits at the epicenter of the energy crisis, given its large dependency on Gulf imports, and has been the fastest to activate demand-side responses – with demand destruction, rather than fiscal absorption, emerging as the primary adjustment channel. The policy actions are estimated to have already curbed oil demand by around 1 mb/d in March and April, with global demand growth for 2026 revised down to 640,000 b/d – 210,000 b/d below the pre-war estimate. Around 56% of emerging market economies in the region have adopted this approach, reflecting tighter fiscal space and more binding financing constraints that make explicit subsidy expansion politically and fiscally untenable. The IEA's emergency demand toolkit – covering work-from-home mandates, reduced highway speed limits, public transport incentives and number-plate rotation schemes – has been the reference framework for most responses, targeting road transport, which accounts for around 45% of global oil demand. The Philippines declared a national energy emergency; Thailand proposed a three-stage fuel rationing contingency plan; Vietnam activated fuel contingency plans; and Indonesia began rationing energy and urging conservation. Pakistan introduced a four-day working week for public offices, with 50% of staff working from home; Sri Lanka reintroduced the weekly fuel ration it had last used during its 2022 debt crisis, alongside a four-day government working week. A secondary consequence has been an acceleration of renewable energy commitments: South Korea's president described the crisis as a strategic opportunity to transition to renewables, while the Philippines and Vietnam have recorded sharp increases in EV sales as consumers seek to reduce fuel exposure.

Markets are closing the remainder through price-driven demand destruction, especially in Asia, where the price premium is 6-7% above the Brent benchmark since the war started.

The policy response under a physical energy shock

The policy response to the Hormuz disruption has been unusually fast. Within days, governments moved to limit the pass through of higher energy prices to end consumers, announcing price caps and subsidy extensions, often complemented by emergency reserve releases – in many cases before physical shortages had fully materialized (see Table 1). The speed reflects a hard learned lesson from 2022: delayed intervention risks de-anchoring inflation expectations and ultimately leads to higher ex-post fiscal costs.

Yet the adjustment differs sharply across countries. Economies with deeper fiscal buffers and market access are once again absorbing a significant share of the shock onto public balance sheets, effectively smoothing the impact on household incomes and corporate margins. In contrast, tighter financing constraints are forcing a larger share of the shock to pass through to activity, combining selective fiscal support with demand compression. In these economies, higher energy prices translate more directly into lower consumption, reduced industrial activity and slower growth, rather than higher fiscal outlays.

Table 1: Fiscal response divergence between DMs and EMs: prices vs demand


	Country	Price adjustments	Demand adjustments	Supply adjustments (non-IEA release)	Supply adjustments – IEA release	Other measures
Developed Markets	United States	No	No	No	Yes	No
	United Kingdom	Yes	No	No	Yes	No
	Europe	No	No	No	Yes	No
	Germany	Yes	No	No	Yes	Yes
	France	Yes	No	No	Yes	No
	Italy	Yes	No	No	Yes	Yes
	Spain	Yes	No	No	Yes	No
	Netherlands	No	No	No	Yes	No
	Belgium	No	No	No	Yes	No
	Japan	Yes	No	No	Yes	No
	Australia	Yes	No	No	Yes	No
	Singapore	No	Yes	Yes	No	No
	AVERAGE	58%	8%	8%	92%	17%
Emerging Markets	China	No	No	Yes	No	No
	India	Yes	No	Yes	No	No
	Pakistan	Yes	Yes	No	No	Yes
	Thailand	Yes	Yes	No	No	No
	Philippines	No	Yes	Yes	No	Yes
	Sri Lanka	Yes	Yes	Yes	No	Yes
	Vietnam	Yes	Yes	Yes	No	No
	Brazil	Yes	No	Yes	No	No
	Turkiye	Partial	No	No	No	No
	South Africa	Yes	No	No	No	No
	Poland	No	No	No	Yes	No
	AVERAGE	89%	56%	67%	11%	33%

Sources: various Allianz Research

Advanced economies: price-based shielding, with heterogeneous scale and timing. Drawing on the 2022 playbook, the initial response in Europe is still overwhelmingly price-based rather than income-based – reflecting administrative ease and political incentives..

This logic is visible again in the current shock, albeit at a smaller scale and with sharper cross country divergence. Under a three-month 30% oil price shock, the direct market impact for large EU energy importers is estimated at around 0.4% of GDP on average, depending on energy dependency (Figure 1). To date, discretionary fiscal measures amounts to roughly 0.15% of GDP, leaving a material and uneven residual burden on households and firms.


Figure 1: Total cost of the war (% of GDP)



Sources: Eurostat, Allianz Research

National strategies nonetheless diverge markedly. Spain acted early and forcefully, rolling out more than 80 broad measures centered on price caps and subsidies, significantly cushioning households and firms but delaying demand adjustment and shifting a large share of the burden onto the public balance sheet. France, by contrast, intervened later and far more selectively showing some fiscal consolidation efforts, with a fiscal response close to one tenth of its 2022 effort, accepting greater price pass through up front. Other countries have so far refrained from announcing discretionary support altogether: Belgium, constrained by limited fiscal space, and the Netherlands, reflecting a more explicitly frugal policy stance, have largely allowed prices to adjust.

Figure 2: Fiscal positions have improved since 2021, but higher spending pressures and interest costs limit the room to act




Sources: national, LSEG Datastream, Allianz Research. Note: blue bubbles: announced measures (size=estimated fiscal cost, % of GDP). Yellow bubble: no fiscal support announced as of 22.04.2026. Non-filled bubbles are for better readability.

Fiscal responses across EMs have been uneven, averaging around 0.2% of GDP, but rising well above that level in more fiscally resilient economies. Singapore recently announced measures amounting to about USD780mn (0.12% of GDP), while larger Asian Ems – including South Korea, India and Indonesia – have deployed packages exceeding 1% of GDP. In this sense, the Iran war is acting as a stress test of fiscal credibility – one that the Covid-19 shock, cushioned by unprecedented global liquidity, did not fully impose.

Despite entering the current shock in better aggregate fiscal positions than during the pandemic, EM starting points remain highly heterogeneous. Unlike the near-universal fiscal deterioration of 2020, fiscal balances today diverge sharply: while Brazil and Vietnam have worsened since the pandemic, several Asian economies have consolidated materially. India has narrowed its deficit from -4% to -2% of GDP, Sri Lanka from -11% to -5%, Thailand from -7% to -2.6% and Malaysia from -6% to -3%. This dispersion in fiscal space is now the primary determinant of how much each government can absorb the energy shock without triggering financing stress or a disorderly

currency adjustment – and explains why, for many EMs, demand destruction remains the dominant adjustment mechanism.

Figure 3: Emerging markets' fiscal space determined the fiscal response to Iran war




Sources: national, LSEG Datastream, Allianz Research.

A “second salvo” coming? Should the conflict prove more prolonged than our baseline assumes, a second wave of fiscal pressure becomes plausible. Fiscal costs could rise non linearly, towards 0.6–0.8% of GDP, roughly double current package sizes — compounding debt dynamics at precisely the moment when the first round of support has already eroded buffers and reintroduced sustainability concerns. Energy price increases pass through to food, transport and services with a lag of several months — and when that broader inflation arrives at the household level, governments face renewed political pressure to respond. The countries most exposed to are those where fiscal space is already thin. In Emerging Markets, high social sensitivity to food and fuel prices puts Morocco, Egypt, Pakistan, Hungary and Poland in the front line. In Europe the picture is equally uncomfortable: France, Italy, Belgium and Poland are already subject to the European Commission's Excessive Deficit Procedure, meaning any additional fiscal response risks deepening consolidation challenges that were already difficult to meet.

Not 2022 reloaded: smaller fiscal shields and weaker price transmission

The Iran 2026 shock is meaningfully smaller than 2022, particularly for gas – the more relevant benchmark for advanced economies, given their greater exposure to gas-indexed power and heating costs. Dutch TTF rose from EUR 31/MWh pre-conflict to a peak of around EUR 70/MWh in early March, before easing back to EUR 40–42/MWh as ceasefire talks progressed – representing roughly 10–12% of the extreme 2022 peak levels, when TTF surged over 1,000%. Brent told a similar but less dramatic story, reaching around 70–75% of its 2022 peak move. Yet the external cost is already visible: in the first 44 days of the Iran crisis alone, the EU spent approximately EUR22bn more on fossil fuel imports – an early reminder of the speed at which external balances can deteriorate, as illustrated by 2022, when the EU's total fossil fuel import bill more than doubled from EUR313bn to EUR693bn at its peak (Figure 4).

Figure 4: Oil prices surged sharply, while gas prices remain below their 2022 peaks.



Sources: LSEG Datastream, Allianz Research


Compared with the 2022 energy crisis, the current shock differs in several important and reinforcing ways. First, the fiscal response has been materially smaller. With the current shock estimated at less than half of the intensity of 2022, fiscal exposure is correspondingly lower. Discretionary measures announced so far across Europe amount to roughly half the scale deployed in 2022. A full 2022 equivalent response (around 1.2% of GDP over the first year) scaled to a 12-week horizon would imply net fiscal costs of about 0.3% of GDP. Crucially, most have been designed as explicitly temporary – typically limited to two to three months. This contrasts sharply with 2022, when interventions initially announced as temporary were repeatedly extended, substantially amplifying their cumulative fiscal cost. The lessons of 2022 are instructive. The EU's net budgetary cost over 2022–2024 reached 2.2% of GDP cumulatively. Around 60% was spent on broad price measures, with only about a quarter targeted at vulnerable households and firms. The takeaway was clear: price suppression can briefly dampen inflation, but it weakens incentives to conserve energy, props up demand and creates large fiscal liabilities that become politically difficult to unwind. The more restrained approach today reflects not only policy learning, but also a less forgiving macroeconomic environment. In 2022, post-pandemic reopening, strong nominal GDP growth and still-accommodative financial conditions helped absorb large fiscal packages. Today, despite improved government balances, near-stagnant real growth limits governments' ability to smooth shocks through public balance sheets without undermining debt sustainability. Fiscal intervention has therefore been more selective, shorter-lived, and – in many cases – explicitly conditional on market developments.

Second, price transmission has structurally weakened relative to 2022, reflecting changes in the European energy mix. Higher penetration of wind and solar generation has reduced the mechanical link between gas prices and electricity prices that defined the 2022 crisis. Back then, gas-fired generation frequently set the marginal price across most hours, transforming a gas supply shock into a broad inflation and fiscal shock. Today, low-marginal-cost renewables increasingly anchor electricity prices, with gas determining prices mainly during periods of tight supply. This has dampened the inflationary impulse and reduced the immediate need for large-scale fiscal intervention – even as gas prices remain elevated.

Together, the smaller scale, temporary design and weaker price transmission point to clear policy learning and a widespread perception that the shock may be transitory. Governments – much like markets – appear more willing to tolerate greater price pass-through rather than repeat the broad and persistent price suppression seen in 2022. The policy response thus remains closer to 2022 in approach than in scale, favoring short-duration, reversible interventions over open-ended commitments.

The heterogeneity in electricity market structures further reinforces this pattern. Spain now operates with a structurally low gas-to-electricity pass-through, estimated at around 0.2–0.4, reflecting high renewable penetration and growing storage capacity. Electricity prices have therefore remained relatively contained during the current shock, limiting inflation spillovers and reducing the need for fiscal intervention. Italy, by contrast, remains constrained by gas-marginal pricing, with gas setting electricity prices most of the time. This implies a much stronger transmission – around 0.7–0.9 – keeping power prices closely linked to gas despite ongoing renewable expansion and sustaining pressure for policy support.

Figure 5: Increased reliance on renewables helped to manage electricity prices (wholesale, EUR/MWh)



Sources: EMBER, Allianz Research

This structural divergence matters for fiscal policy. Where renewables dampen transmission, governments can accept greater pass-through without triggering an inflation–fiscal feedback loop. Where gas remains marginal, price pressures are faster and more persistent, raising political incentives to intervene even when growth is weak and fiscal space constrained. Importantly, the current configuration embeds a non-linear downside risk. If physical disruptions persist beyond the horizon implicitly assumed in today's policy design, and temporary measures are repeatedly rolled over or expanded, fiscal costs in Europe could still rise sharply. In an environment of weak growth and higher interest rates, prolonged price interventions would once again convert a supply shock into a growing contingent fiscal liability, this time without the mitigating effect of strong nominal growth that characterized 2022.

Debt sustainability dynamics reinforce this concern. In 2022, double-digit inflation lifted nominal growth well above financing costs, providing a buffer. Today, with inflation around 3% and near-zero real growth, nominal GDP growth is too weak to offset rising deficits under price caps. Debt dynamics have therefore become increasingly sensitive to the duration of fiscal support.

Winners and losers from the Iran shock

For now, fiscal costs remain manageable for most countries. On average, debt-servicing costs are set to rise by around 0.05% of GDP in both advanced and emerging economies, driven entirely by a higher debt stock rather than higher rates. However, energy import dependence combined with a deteriorating current account determines who faces the largest shock. Interest rates have risen only modestly — US and German yields up around 30bps, EM hard currency spreads +20bps, local currency +30bps — with real rate changes even smaller. The slow pace of debt rollover means higher borrowing costs pass through only gradually, limiting the immediate debt-service impact. On the flip side, higher inflation could partially erode the real debt burden through an inflation tax effect, though the currently expected inflation increase is far below 2022 levels, limiting this offset.

Debt service and fiscal fragility. The countries most acutely exposed are those carrying both high near-term debt-service obligations and significant structural dependence on energy imports. Egypt sits in the most dangerous quadrant: debt service above 10% of GDP, heavy reliance on imported oil and gas, an IMF program already under strain and Suez Canal revenues simultaneously compressed by the conflict. Pakistan, with debt service at 5.4% of GDP and near-total dependence on Gulf crude, faces import cost inflation and financing stress simultaneously. Hungary and Poland carry elevated fiscal deficits with limited room to sustain price support measures, while Sri Lanka — only recently exited from sovereign default — and Bangladesh face the energy shock as a direct balance-of-payments stress, with thin reserve buffers and limited debt market depth meaning the import cost increase arrives before any fiscal response can be mounted. Peripheral Eurozone sovereign bonds have also experienced one of their worst months in a decade, with spreads widening most sharply in Italy, Greece and Portugal — markets are not primarily repricing current debt stocks but the open-ended flow commitment implied by sustained price interventions, and the longer caps and subsidies remain in place, the more the fiscal trajectory diverges from the consolidation paths embedded in current spread levels.


Currency depreciation and hard currency debt. Currency weakness amplifies the damage for sovereigns with a high share of hard currency debt, as depreciation mechanically increases the local currency cost of servicing existing obligations precisely when fiscal revenues are under strain. Türkiye is the most exposed large EM across all dimensions: importing 90% of its energy at an annual cost of USD 60-65bn, with a current account deficit of USD 25bn and a lira down 17% over the past year — each USD 10 oil price increase adds USD 4.5-5bn to its external deficit. Argentina and Egypt face the same amplification dynamic, both carrying meaningful foreign currency debt shares above 25-30% of total sovereign debt. Tunisia has recorded among the sharpest currency depreciation since the start of the conflict, while the Thai baht has absorbed part of the shock with a 6% depreciation. The OECD's 2025 Global Debt Report confirms that among large emerging markets, only Argentina and Türkiye carry foreign currency debt shares exceeding 30% of total sovereign debt — making them uniquely vulnerable to the depreciation-debt spiral that the energy shock is now activating.

Monetary policy and real rates. One month and a half into the conflict, energy import dependency continues to weigh more heavily than financial constraints on central bank decisions. Singapore has tightened monetary policy for the first time in four years; the Bank of Japan is expected to follow with a 25bps hike in Q2, most likely June. The risk-off shift has pushed markets to reprice central bank paths globally — in Europe, expectations have shifted from cuts to hikes for the ECB, though we anticipate adjustments only for the duration of the conflict. Countries recovering from the 2022-23 inflation cycle will slow their easing, while others are pushing rate cuts to 2027. Egypt's CBE held rates unchanged at its 2 April meeting, treating the war as an external supply shock and allowing the pound to

depreciate as the primary adjustment mechanism — a stance that preserves monetary credibility for now but leaves the carry trade inherently fragile, as the central bank could face an impossible choice between further depreciation and an aggressive rate hike to retain investors. Chile, Indonesia and Thailand have kept rates on hold, pending clearer signals on the duration of the disruption.

Current account deterioration. The Iran war is exerting pressure on sovereign bond markets through two simultaneous channels: a global risk-off repricing that widens EM spreads, and a direct fiscal deterioration driven by higher energy import costs. Oil above USD 100 per barrel widens current account deficits, accelerates currency depreciation and feeds domestic inflation, eroding sovereign creditworthiness precisely when refinancing needs are greatest. India, at 3.9% debt service, has larger domestic energy buffers but its bond market remains exposed to the broader EM risk-off and dollar strengthening that higher oil prices typically generate. Vietnam faces meaningful energy import dependence alongside limited fiscal space, while Kenya — a large energy importer carrying a high fiscal deficit after years of difficult IMF negotiations — has thin reserves providing little cushion against a sustained import cost shock. Morocco and Tunisia occupy a similar intersection of moderate-to-high debt service and meaningful energy import dependency, leaving their sovereign spreads vulnerable to both the global risk premium repricing and the direct pass-through of higher import costs.

Figure 6: Most exposed: where high debt burdens (% of GDP) meet energy import dependence



Sources: UN COMTRADE, World Bank and Allianz Research

Winners. Not all emerging markets are on the losing side. Oil-exporting economies stand to benefit materially from sustained prices above USD 100 per barrel, with the windfall flowing through both fiscal revenues and current account positions. Nigeria is the most prominent African beneficiary — every USD 10 increase in oil prices adds approximately USD 3-4bn to government revenues — though its structural paradox partially erodes the net gain, as a major crude exporter that nonetheless imports virtually all of its refined petroleum products. Brazil, as significant oil producer, will see their fiscal and external positions improve as Brent rises, while Colombia and Ecuador capture a similar tailwind from hydrocarbon exports. Chile is the notable exception within the region: its early and deep investment in renewable energy has structurally reduced exposure to oil price volatility, insulating the economy from the full pass-through hitting more fossil-fuel-dependent peers. Latin America stands out as the emerging market region least exposed to the shock overall — where for most economies the Iran war arrives not as a fiscal and external stress but as a terms-of-trade tailwind.

These assessments are, as always, subject to the disclaimer provided below.

FORWARD-LOOKING STATEMENTS

The statements contained herein may include prospects, statements of future expectations and other forward-looking statements that are based on management's current views and assumptions and involve known and unknown risks and uncertainties. Actual results, performance or events may differ materially from those expressed

or implied in such forward-looking statements.

Such deviations may arise due to, without limitation, (i) changes of the general economic conditions and competitive situation, particularly in the Allianz Group's core business and core markets, (ii) performance of financial markets (particularly market volatility, liquidity and credit events), (iii) frequency and severity of insured loss events, including from natural catastrophes, and the development of loss expenses, (iv) mortality and morbidity levels and trends,

(v) persistency levels, (vi) particularly in the banking business, the extent of credit defaults, (vii) interest rate levels, (viii) currency exchange rates including the EUR/USD exchange rate, (ix) changes in laws and regulations, including tax regulations, (x) the impact of acquisitions, including related integration issues, and reorganization measures,

and (xi) general competitive factors, in each case on a local, regional, national and/or global basis. Many of these factors may be more likely to occur, or more pronounced, as a result of terrorist activities and their consequences.

NO DUTY TO UPDATE

The company assumes no obligation to update any information or forward-looking statement contained herein, save for any information required to be disclosed by law.

Allianz Trade is the trademark used to designate a range of services provided by Euler Hermes.